

Python Zero to Godhood

The Definitive Guide from Python 1.0 to Python 3.14

Master the Runtime.

Shreejit Verma

June 10, 2026

Contents

1	Preface	1
2	Python Zero to Godhood: Complete Evolution and Comprehensive Feature Guide	2
2.1	Preface	2
2.1.1	About the Author	2
2.1.2	Book Purpose & Scope	2
2.2	Table of Contents	2
2.2.1	Volume I: Classic Python & Core Engine (Python 1.0 to 2.7)	2
2.2.2	Volume II: The Python 3 Schism & Core Enhancements (Python 3.0 to 3.2)	2
2.2.3	Volume III: Generators, Iterators, and Async Inception (Python 3.3 to 3.5)	3
2.2.4	Volume IV: Expressiveness & Developer Ergonomics (Python 3.6 to 3.7)	3
2.2.5	Volume V: Structural Shifts & Pattern Matching (Python 3.8 to 3.10)	3
2.2.6	Volume VI: Performance Leap & Runtime Mechanics (Python 3.11 to 3.12)	3
2.2.7	Volume VII: The GIL-less Future & JIT Compilers (Python 3.13 to 3.14+)	3
2.2.8	Volume VIII: Runtime Internals & C Extensions	3
2.2.9	Volume IX: High Performance & Low Latency Concurrency	3
I	Volume I: Classic Python & Core Engine	4
3	Python 1.0 to 1.6: Inception & the LL(1) Executable Pipeline	5
3.0.1	1.1 Chronological Inception & Core Design Philosophies	5
3.0.2	1.2 CPython Parser Evolution: LL(1) Left-Recursion Math vs. PEG	5
3.0.3	1.3 Compilation, AST, and the Symbol Table Pass	6
3.0.4	1.4 PyCodeObject & Bytecode Anatomy	7
3.0.5	1.5 The Virtual Machine Execution Loop & Frame Objects	8
4	Python 1.x: The PyObject Model & Reference Counting Core	11
4.0.1	2.1 The Unified PyObject & PyVarObject Structs	11
4.0.2	2.2 The Type-Object Slot System	12
4.0.3	2.3 Reference Counting Lifecycle	13
4.0.4	2.4 Built-in Caching and Interning Optimizations	13
5	Python 2.0 to 2.1: Comprehensions, Nested Scopes, & Cycle-Detecting GC	15
5.0.1	3.1 Python 2.0 List Comprehensions & Scope Leakage	15
5.0.2	3.2 PEP 227 Nested Scopes and LEGB Name Resolution	16
5.0.3	3.3 Closure Mechanics and Bytecode Compilation	16
5.0.4	3.4 Python 2.0 Cycle-Detecting Generational Garbage Collector	17

6	Python 2.2 to 2.3: Type-Class Unification, Descriptors, & C3 MRO	19
6.0.1	4.1 Type-Class Unification: The Inception of New-Style Classes (PEP 252 & PEP 253)	19
6.0.2	4.2 The Descriptor Protocol & Attribute Lookup Chain	20
6.0.3	4.4 Instance Lifecycle: Allocation vs. Initialization	25
7	Python 2.4 to 2.7: Decorators, Context Managers, & the 2.x Twilight	26
7.0.1	5.1 Decorators: Syntactic Sugar and Compiler Mechanics (PEP 318 & PEP 3129)	26
7.0.2	5.2 Context Managers and the <code>with</code> Statement (PEP 343)	27
7.0.3	5.3 Generator Enhancements: Coroutines and Frame Suspension (PEP 342)	28
7.0.4	5.4 Under-the-Hood CPython Data Structure Layouts	29
8	Python 2.x: Low-Level File I/O & Exceptions Unwinding Blocks	31
8.0.1	6.1 Low-Level File Management and Standard Streams wrapping	31
8.0.2	6.2 Exception Mechanics and the Thread State slots	32
8.0.3	6.3 The Try-Except-Finally Block Stack and Stack Restoration	33
II	Volume II: The Python 3 Schism & Core Enhancements	34
9	Python 3.0: The Unicode Paradigm Shift and Text vs. Bytes Separation	35
9.0.1	7.1 The Unicode Paradigm Shift and Text vs. Bytes Separation (PEP 358 & PEP 3112)	35
9.0.2	7.2 The <code>print()</code> Function Redesign	37
9.0.3	7.3 Division Unification (PEP 238)	38
9.0.4	7.4 Lazy Iterators & Dictionary View Objects	39
10	Python 3.1 to 3.2: Standard Library Consolidation and Threading Pools	40
10.0.1	8.1 Antoine Pitrou's New GIL (Python 3.2)	40
10.0.2	8.2 Thread Pools and Process Pools (<code>concurrent.futures</code> / PEP 3148)	41
10.0.3	8.3 Standard Library Consolidation	42
III	Volume III: Generators, Iterators, and Async Inception	44
11	Python 3.3: Yield From Generators and Implicit Namespace Packages	45
11.0.1	9.1 PEP 380: Generator Delegation via <code>yield from</code>	45
11.0.2	9.2 PEP 420: Implicit Namespace Packages	46
11.0.3	9.3 PEP 393: Flexible String Representation	47
11.0.4	9.4 <code>memoryview</code> and the Buffer Protocol (<code>Py_buffer</code>)	48
12	Python 3.4: Asyncio Inception, Pathlib, and Enum Architectures	49
12.0.1	10.1 Asyncio Inception: Generators and Event Loops (PEP 3156)	49
12.0.2	10.2 <code>pathlib</code> : Object-Oriented Filesystem Paths (PEP 428)	50
12.0.3	10.3 Enum: Metaclass and Singletons (PEP 435)	51
13	Python 3.5: Native Async/Await Coroutines and Matrix Operations	52
13.0.1	11.1 PEP 492: Native Coroutines (<code>async/await</code> Internals)	52
13.0.2	11.2 Matrix Multiplication Operator (<code>@</code> / PEP 465)	54
13.0.3	11.3 Extended Unpacking Generalizations (PEP 448)	56

IV	Volume IV: Expressiveness & Developer Ergonomics	58
14	Python 3.6: F-Strings Formatting, Variable Annotations, and Compact Dicts	59
14.0.1	12.1 PEP 498: Formatted String Literals (f-strings)	59
14.0.2	12.2 PEP 526: Syntax for Variable Annotations	60
14.0.3	12.3 CPython Compact Dictionary Design	62
15	Python 3.7: Dataclasses, Context Variables, and Dict Ordering Guarantees	64
15.0.1	13.1 PEP 557: Dataclasses under the hood	64
15.0.2	13.2 PEP 567: Context Variables (<code>contextvars</code>)	65
15.0.3	13.3 Dictionary Insertion-Order Guarantee	65
15.0.4	13.4 Properties & Cached Properties (Descriptor Protocol)	66
15.0.5	13.5 Instance Slots Optimization (<code>__slots__</code>)	67
V	Volume V: Structural Shifts & Pattern Matching	69
16	Python 3.8: Walrus Operator (<code>:=</code>) and Positional-Only Parameters (<code>/</code>)	70
16.0.1	14.1 The Assignment Expression (<code>:=</code> / PEP 572)	70
16.0.2	14.2 Positional-Only Parameters (<code>/</code> / PEP 570)	72
17	Python 3.9 to 3.10: PEG Parser, Dict Merge (<code> </code>), and Pattern Matching	74
17.0.1	15.1 CPython Parser Evolution: LL(1) to PEG (PEP 617)	74
17.0.2	15.2 PEP 584: Dictionary Union Operators (<code> </code> and <code> =</code>)	75
17.0.3	15.3 PEP 634: Structural Pattern Matching (Decision Trees)	75
18	Python 3.8 to 3.10: Type Hinting Protocols and Structural Subtyping	78
18.0.1	16.1 Nominal vs. Structural Subtyping	78
18.0.2	16.2 Runtime Protocols & <code>@runtime_checkable</code>	79
18.0.3	16.3 Static Type Checking and Variance	80
VI	Volume VI: Performance Leap & Runtime Mechanics	82
19	Python 3.11: Faster CPython Specializing Interpreter and Adaptive Bytecode	83
19.0.1	17.1 PEP 659 Specialized Adaptive Interpreter Architecture	83
19.0.2	17.2 Specialized Bytecode Opcodes and Inline Cache Memory Layout	84
19.0.3	17.3 Polymorphic Deoptimization Costs	85
20	Python 3.12: Native Generics (PEP 695), Type statement, and Subinterpreters	86
20.0.1	18.1 PEP 695 Type Parameter Syntax	86
20.0.2	18.2 Annotation Scopes and Lazy Evaluation	87
20.0.3	18.3 PEP 684: Per-Interpreter GIL and Subinterpreters	87
21	Python 3.11 to 3.12: Exception Groups (<code>except*</code>) and Traceback trees	90
21.0.1	19.1 Exception Groups (<code>ExceptionGroup</code> & <code>BaseExceptionGroup</code> / PEP 654)	90
21.0.2	19.2 The <code>except*</code> Clause and Exception Tree Filtering	91
21.0.3	19.3 Exception Tree Filtering Algorithm	92
21.0.4	19.4 Traceback Representation and <code>add_note()</code> Internals	93

VII	Volume VII: The GIL-less Future & JIT Compilers	96
22	Python 3.13: Free-Threaded Build & GIL Removal Internals	97
22.0.1	20.1 The Free-Threaded CPython Paradigm Shift	97
22.0.2	20.2 Biased Reference Counting (BRC) and Object Headers	97
22.0.3	20.3 Thread-Safe Allocation via mimalloc	99
22.0.4	20.4 Lock-free Collections & Thread Safety	100
22.0.5	20.5 Generational Garbage Collection (GC) in a GIL-less World	101
23	Python 3.13: Copy-and-Patch JIT Compiler Architecture	103
23.0.1	21.1 The Copy-and-Patch JIT Paradigm (PEP 744)	103
23.0.2	21.2 Stencil Relocation and C-Level Internals	103
23.0.3	21.3 Tier 2 Micro-Ops (uops) and the Optimization Pipeline	104
23.0.4	21.4 Relocation Hole Patching and Memory Security	105
24	Python 3.12 to 3.13: Subinterpreters & Per-Interpreter GIL Parallelism	107
24.0.1	22.1 Isolation of Interpreter State (PEP 684)	107
24.0.2	22.2 The C-API Subinterpreter Initialization Configuration	108
24.0.3	22.3 Cross-Interpreter Communication Channels	109
VIII	Volume VIII: Runtime Internals & C Extensions	112
25	CPython Memory Allocator (PyMalloc) & Generational Garbage Collection	113
25.0.1	23.1 CPython's Memory Allocation Engine (PyMalloc)	113
25.0.2	23.2 Generational Garbage Collection and Cycle Detection	114
25.0.3	23.3 The Cycle Detection Algorithm	115
26	C Extensions & Python C-API Interoperability	117
26.0.1	24.1 Writing a Pure C Extension Module	117
26.0.2	24.2 Reference Counting and Ownership Rules in C-API	118
26.0.3	24.3 C-API Stable ABI (PEP 384)	119
27	Metaclasses, Descriptor Protocol, and type Slots	121
27.0.1	25.1 The Metaclass Class Creation Pipeline	121
27.0.2	25.2 The Descriptor Protocol and Attribute Lookup Precedence	122
27.0.3	25.3 CPython Type Slots	123
IX	Volume IX: High Performance & Low Latency Concurrency	125
28	LOW-LEVEL MEMORY OPTIMIZATION TECHNIQUES	126
28.0.1	26.1 <code>__slots__</code> Internals and Memory Mapping	126
28.0.2	26.2 Buffer Protocol and <code>memoryview</code> Zero-Copy Slicing	127
28.0.3	26.3 Compact Serialization: array vs. struct	128
29	CPU & I/O BOUND SYSTEM CONCURRENCY	129
29.0.1	27.1 Concurrency Paradigms: Threads vs. Processes vs. Coroutines (Asyncio)	129
29.0.2	27.2 Asyncio Event Loops and <code>uvloop</code> Internals	129
29.0.3	27.3 Multiprocessing Shared Memory IPC	130

30 NUMERICAL HIGH-PERFORMANCE DATA STRUCTURES	132
30.0.1 28.1 NumPy Contiguous Layouts and Strides	132
30.0.2 28.2 PyArrow Columnar Format and Zero-Copy Sharing	133
31 PROFILING, BENCHMARKING, AND DIAGNOSTICS	135
31.0.1 29.1 Deterministic Profiling (<code>cProfile</code>) vs. Sampling Profiling (<code>py-spy</code>) .	135
31.0.2 29.2 Memory Leak Diagnostics via <code>tracemalloc</code>	136
31.0.3 29.3 Crash Debugging via <code>faulthandler</code>	137
32 CAPSTONE PROJECT: HIGH-FREQUENCY ORDER BOOK & TRADING ENGINE	139
32.0.1 30.1 High-Frequency Trading Engine Architecture	139
32.0.2 30.2 Optimized Order Book Implementation	139
32.0.3 30.3 High-Throughput Asyncio Server	142

Chapter 1

Preface

Chapter 2

Python Zero to Godhood: Complete Evolution and Comprehensive Feature Guide

Author: Shreejit Verma

2.1 Preface

2.1.1 About the Author

Shreejit Verma is a systems architect, quantitative software engineer, and high-performance computing practitioner. This guide represents a masterclass in CPython internals, language syntax evolution, runtime mechanics, and hardware-sympathetic programming.

2.1.2 Book Purpose & Scope

This book provides a transition pathway from standard Python development to **low-level CPython mastery and high-performance computing**. It spans the entire chronological architecture of the language from Python 1.0 to Python 3.14 revealing how bytecode interpreters, memory systems, GIL execution models, and compiler optimizations interact.

2.2 Table of Contents

2.2.1 Volume I: Classic Python & Core Engine (Python 1.0 to 2.7)

- Chapter 1: [Python 1.0 to 1.6: Inception & the LL\(1\) Executable Pipeline](#)
- Chapter 2: [Python 1.x: The PyObject Model & Reference Counting Core](#)
- Chapter 3: [Python 2.0 to 2.1: Comprehensions, Nested Scopes, & Cycle-Detecting GC](#)
- Chapter 4: [Python 2.2 to 2.3: Type-Class Unification, Descriptors, & C3 MRO](#)
- Chapter 5: [Python 2.4 to 2.7: Decorators, Context Managers, & the 2.x Twilight](#)
- Chapter 6: [Python 2.x: Low-Level File I/O & Exceptions Unwinding Blocks](#)

2.2.2 Volume II: The Python 3 Schism & Core Enhancements (Python 3.0 to 3.2)

- Chapter 7: [Python 3.0: The Unicode Paradigm Shift and Text vs. Bytes Separation](#)
- Chapter 8: [Python 3.1 to 3.2: Standard Library Consolidation and Threading Pools](#)

2.2.3 Volume III: Generators, Iterators, and Async Inception (Python 3.3 to 3.5)

- Chapter 9: [Python 3.3: Yield From Generators and Implicit Namespace Packages](#)
- Chapter 10: [Python 3.4: Asyncio Inception, Pathlib, and Enum Architectures](#)
- Chapter 11: [Python 3.5: Native Async/Await Coroutines and Matrix Operations](#)

2.2.4 Volume IV: Expressiveness & Developer Ergonomics (Python 3.6 to 3.7)

- Chapter 12: [Python 3.6: F-Strings Formatting, Variable Annotations, and Compact Dicts](#)
- Chapter 13: [Python 3.7: Dataclasses, Context Variables, and Dict Ordering Guarantees](#)

2.2.5 Volume V: Structural Shifts & Pattern Matching (Python 3.8 to 3.10)

- Chapter 14: [Python 3.8: Walrus Operator \(:=\) and Positional-Only Parameters \(/\)](#)
- Chapter 15: [Python 3.9 to 3.10: PEG Parser, Dict Merge \(|\), and Pattern Matching](#)
- Chapter 16: [Python 3.8 to 3.10: Type Hinting Protocols and Structural Subtyping](#)

2.2.6 Volume VI: Performance Leap & Runtime Mechanics (Python 3.11 to 3.12)

- Chapter 17: [Python 3.11: Faster CPython Specializing Interpreter and Adaptive Bytecode](#)
- Chapter 18: [Python 3.12: Native Generics \(PEP 695\), Type statement, and Subinterpreters](#)
- Chapter 19: [Python 3.11 to 3.12: Exception Groups \(except*\) and Traceback trees](#)

2.2.7 Volume VII: The GIL-less Future & JIT Compilers (Python 3.13 to 3.14+)

- Chapter 20: [Python 3.13: Free-Threaded Build & GIL Removal Internals](#)
- Chapter 21: [Python 3.13: Copy-and-Patch JIT Compiler Architecture](#)
- Chapter 22: [Python 3.12 to 3.13: Subinterpreters & Per-Interpreter GIL Parallelism](#)

2.2.8 Volume VIII: Runtime Internals & C Extensions

- Chapter 23: [CPython Memory Allocator \(PyMalloc\) & Generational Garbage Collection](#)
- Chapter 24: [C Extensions & Python C-API Interoperability](#)
- Chapter 25: [Metaclasses, Descriptor Protocol, and type Slots](#)

2.2.9 Volume IX: High Performance & Low Latency Concurrency

- Chapter 26: [Low-Level Memory Optimization \(Slots, memoryviews, and Weak References\)](#)
 - Chapter 27: [Concurrency Architectures: Threading vs. Multiprocessing vs. Asyncio](#)
 - Chapter 28: [Numerical and Scientific Data \(NumPy SIMD, PyArrow Zero-Copy\)](#)
 - Chapter 29: [Profiling, Benchmarking, and System Diagnostics](#)
 - Chapter 30: [Capstone Project: High-Frequency Order Book and Trading Engine](#)
-

Part I

Volume I: Classic Python & Core Engine

Chapter 3

Python 1.0 to 1.6: Inception & the LL(1) Executable Pipeline

3.0.1 1.1 Chronological Inception & Core Design Philosophies

Python was conceived in December 1989 by Guido van Rossum at CWI (Centrum Wiskunde & Informatica) in the Netherlands. It was created as a successor to the **ABC language**, a programming language designed for beginners that was elegant but suffered from fatal architectural flaws: 1. **Monolithic Runtime**: ABC was difficult to extend. It did not support dynamic library loading or simple integration with C. 2. **Strict Filesystem Database**: ABC stored variable states globally in a monolithic filesystem database, which severely limited performance and concurrent operations. 3. **Closed Ecosystem**: ABC was a closed system, making it impossible to interface directly with low-level OS structures or hardware calls.

Van Rossum designed Python as an open-source, extensible scripting language that addressed these issues: * **Whitespace Significance**: Unlike languages that use curly braces (`{}`), CPython's tokenizer tracks indentations (`INDENT`) and dedentations (`DEDENT`) as formal syntax control elements. This ensures visual block structure matches execution scopes. * **Unified Object Model**: Every data structure is a heap-allocated object (`PyObject`), allowing unified interfaces and straightforward extensibility via C extensions. * **Names and Bindings**: Python variables are not declared memory slots containing values; they are references (pointers) to heap-allocated objects. Thus, `x = 5` binds the label `x` in the namespace dictionary to the `PyObject` representing the integer 5.

3.0.2 1.2 CPython Parser Evolution: LL(1) Left-Recursion Math vs. PEG

Until Python 3.9, CPython used a custom **LL(1) parser** (Left-to-right, Leftmost derivation with 1 token lookahead) to construct its concrete parse trees.

1. Left-Recursion in LL(1)

An LL(1) parser is a top-down parser. It cannot parse grammar rules that exhibit **left-recursion**. A grammar rule is left-recursive if the leftmost symbol on the right-hand side of a production is the non-terminal itself:

$$A \rightarrow A\alpha \mid \beta$$

When a top-down parser attempts to expand A using this rule, it loops infinitely without consuming any input:

$$A \rightarrow A\alpha \rightarrow A\alpha\alpha \rightarrow A\alpha\alpha\alpha \rightarrow \dots$$

2. Grammar Rewriting Workarounds

To prevent infinite recursion, developers had to rewrite left-recursive rules into right-recursive forms:

$$A \rightarrow \beta A'$$

$$A' \rightarrow \alpha A' \mid \epsilon$$

While mathematically equivalent, this workaround made the compiler's grammar definition complex, hard to maintain, and less readable.

3. Parsing Expression Grammar (PEG) Parser (PEP 617)

Python 3.9 replaced the LL(1) parser with a **PEG parser**. * **Packrat Parsing (Memoization)**: A PEG parser handles left-recursion by tracking and caching parsed rules at each input index. If the parser backtracks or encounters left-recursion, it references the cache to break infinite loops. * **Infinite Lookahead**: Unlike LL(1)'s single-token limit, the PEG parser can look ahead arbitrarily far, allowing more expressive syntax. * **Direct AST Generation**: The PEG parser generates the Abstract Syntax Tree (AST) directly from the tokenizer stream, eliminating the overhead of building an intermediate Concrete Parse Tree.

```

1 LL(1) Parser Pipeline (Classic):
2 [Source] -> [Tokenizer] -> [Concrete Parse Tree] -> [AST Generator] -> [AST]
3
4 PEG Parser Pipeline (Modern 3.9+):
5 [Source] -> [Tokenizer] -> [PEG Parser (Memoized)] -> [AST]
```

3.0.3 1.3 Compilation, AST, and the Symbol Table Pass

After generating the AST, CPython compiles the code. This begins with a static analysis pass to build the symbol table (`Python/symtable.c`).

1. Variable Scope Classification

CPython classifies variable names into four distinct scopes: 1. **Local**: Bound within the current function scope. 2. **Global Explicit**: Declared via `global x`. 3. **Global Implicit**: Referenced but not bound within the current function (resolved at runtime in module globals). 4. **Enclosing (Closure/Free)**: Variables defined in an outer function scope and accessed by a nested inner function.

2. Closure Cells (`PyCellObject`)

When a local variable is accessed by a nested function, it cannot be stored in the standard fast-locals array of the outer function because its lifetime must extend beyond the outer function's execution. CPython handles this by wrapping the variable in a `PyCellObject`.

```

1 typedef struct {
2     PyObject_HEAD
3     PyObject *ob_ref;          /* Pointer to the enclosed PyObject */
4 } PyCellObject;
```

The cell object resides on the heap, allowing both the outer and inner functions to share access to the same object reference, even after the outer function's stack frame is popped.

3. The Dynamic Namespace Trap

Python allows dynamic namespace modifications using features like `exec()`, `eval()`, or `from module import *`. When the compiler detects these dynamic features inside a function, it cannot determine variable scopes at compile-time. As a result, it disables fast local lookups: * **Standard Local Lookup**: Uses the `LOAD_FAST` bytecode, which retrieves references from an array index in $\mathcal{O}(1)$ time. * **Dynamic Fallback**: Reverts to the slow `LOAD_NAME` bytecode, which searches the local, enclosing, global, and builtin dictionaries sequentially at runtime.

4. Symbol Table Inspection

We can inspect this static compilation analysis using the built-in `symtable` module:

```

1 import symtable
2
3 code_text = """
4 def outer_func(x):
5     z = 10
6     def inner_func():
7         nonlocal z
8         return x + z
9     return inner_func
10 """
11
12 table = symtable.symtable(code_text, filename("<string>", compile_type="exec")
13 outer_scope = table.lookup("outer_func").get_namespace()
14
15 print("Outer variables:", outer_scope.get_identifiers())
16 print("Is 'z' a cell variable in outer_func?", outer_scope.lookup("z").is_cell
17       ())
18 print("Is 'x' a cell variable in outer_func?", outer_scope.lookup("x").is_cell
19       ())
20
21 inner_scope = outer_scope.lookup("inner_func").get_namespace()
22 print("Inner variables:", inner_scope.get_identifiers())
23 print("Is 'z' a free (closure) variable in inner_func?", inner_scope.lookup("z"
24       ).is_free())

```

3.0.4 1.4 PyCodeObject & Bytecode Anatomy

The compiler processes the AST and symbol table to generate a `PyCodeObject` struct, which serves as the immutable blueprint for execution.

1. C-Level Structure of PyCodeObject

Defined in `Include/cpython/code.h`, this structure holds the compiled bytecode, constants pool, and variable name catalogs:

```

1 struct PyCodeObject {
2     PyObject_HEAD
3     int co_argcount;           /* Number of positional arguments */
4     int co_posonlyargcount;    /* Positional-only arguments (Python 3.8+) */
5     int co_kwonlyargcount;    /* Keyword-only arguments */
6     int co_nlocals;           /* Number of local variables */
7     int co_stacksize;         /* Maximum value stack depth needed by VM */
8     int co_flags;             /* Compiler configuration flags (e.g.
9     CO_GENERATOR) */

```

```

10 PyObject *co_code;          /* Instruction sequence (bytes or 16-bit code
    units) */
11 PyObject *co_consts;        /* Immutable literal constants (tuple of
    PyObject) */
12 PyObject *co_names;         /* Non-local/Global names referenced (tuple of
    strings) */
13 PyObject *co_varnames;      /* Local parameter and variable names (tuple)
    */
14 PyObject *co_freevars;      /* Free variables captured in closures (tuple)
    */
15 PyObject *co_cellvars;      /* Local variables enclosed by nested scopes (
    tuple) */
16 };
17 typedef struct PyCodeObject PyCodeObject;

```

2. Instruction Decoding and EXTENDED_ARG

- **Bytecode Formats:**

- *Pre-3.11*: Instructions used a 2-byte format (1 byte for the opcode, 1 byte for the argument).
- *Python 3.11+*: Instructions use 16-bit code units (2 bytes), combining opcode, arguments, and inline caches.

- **The EXTENDED_ARG Opcode:** The argument field of a standard instruction is limited to 8 bits (representing indices 0 to 255). If an index exceeds 255 (e.g., loading the 300th constant in `co_consts`), the compiler prefixes the instruction with an EXTENDED_ARG opcode.
 - The EXTENDED_ARG instruction loads the high-order bits of the argument value into an internal register.
 - The subsequent instruction shifts this value and adds its own argument payload, allowing the VM to resolve indices up to 32 bits wide.

Let's write a script to inspect these low-level compiled fields inside a live Python runtime:

```

1 def example_fn(a, b):
2     local_val = 42
3     return a + b + local_val
4
5 code_obj = example_fn.__code__
6
7 print("Local Variable Names (co_varnames):", code_obj.co_varnames)
8 print("Constant Pool (co_consts):", code_obj.co_consts)
9 print("Instruction Byte Sequence (co_code):", list(code_obj.co_code))
10 print("Scoping/Compiler Flags (co_flags):", bin(code_obj.co_flags))

```

3.0.5 1.5 The Virtual Machine Execution Loop & Frame Objects

When a function is called, the CPython virtual machine allocates a new stack frame representation, the `PyFrameObject`, on the heap.

1. C-Level Structure of `PyFrameObject`

Defined in `Include/internal/pycore_frame.h`, the frame object tracks runtime execution state:

```

1 struct _frame {
2     PyObject_HEAD
3     struct _frame *f_back;      /* Link to previous frame (caller's frame) */
4     PyCodeObject *f_code;       /* Code object executed in this frame */
5     PyObject *f_builtins;       /* Builtins symbol lookup dictionary */
6     PyObject *f_globals;        /* Global namespace dictionary */
7     PyObject *f_locals;         /* Local namespace dictionary */
8     PyObject **f_valuestack;    /* Points to the base of the evaluation stack
9     */
10    int f_lasti;                 /* Last instruction index executed */
11    /* ... additional traceback and debugging fields ... */
12 };
13 typedef struct _frame PyFrameObject;

```

2. CPython VM Evaluation Loop (_PyEval_EvalFrameDefault)

The core execution engine is defined in `Python/ceval.c` inside the function `_PyEval_EvalFrameDefault()`. This function contains a monolithic evaluation loop:

```

1 PyObject* _PyEval_EvalFrameDefault(PyThreadState *tstate, PyFrameObject *f, int
2     throwflag) {
3     // Local optimization pointers
4     PyObject **stack_pointer = f->f_valuestack;
5     const _Py_CODEUNIT *next_instr = (_Py_CODEUNIT *)PyBytes_AS_STRING(f->
6     f_code->co_code);
7
8     for (;;) {
9         _Py_CODEUNIT word = *next_instr++;
10        int opcode = _Py_OPCODE(word);
11        int oparg = _Py_OPARG(word);
12
13        switch (opcode) {
14            case LOAD_FAST: {
15                PyObject *value = GETLOCAL(oparg);
16                Py_INCREF(value);
17                PUSH(value);
18                DISPATCH();
19            }
20            case BINARY_OP: {
21                PyObject *right = POP();
22                PyObject *left = POP();
23                PyObject *res = PyNumber_Add(left, right);
24                Py_DECREF(left);
25                Py_DECREF(right);
26                PUSH(res);
27                DISPATCH();
28            }
29            case RETURN_VALUE: {
30                PyObject *retval = POP();
31                return retval;
32            }
33            /* ... additional opcodes ... */
34        }
35    }
36 }

```

3. Value Stack Execution Walkthrough

CPython is a stack-based virtual machine. Operations pop arguments from the evaluation stack and push results back. Here is a trace of evaluating the expression `a + b`:

1	Stack State Trace for BINARY_OP (Add):			
2				
3	1. LOAD_FAST 0 (a)	2. LOAD_FAST 1 (b)	3. BINARY_OP	4.
	RETURN_VALUE			
4	+-----+	+-----+	+-----+	
	+-----+			
5	val_a	val_b	val_a+b	(
	Empty)			
6	+-----+	+-----+	+-----+	
	+-----+			
7	(Empty) -->	val_a -->	(Empty) -->	(
	Empty)			
8	+-----+	+-----+	+-----+	
	+-----+			

The program counter (`next_instr`) increments, fetching code units, popping parameters off the stack, and executing the corresponding C functions.

Chapter 4

Python 1.x: The PyObject Model & Reference Counting Core

4.0.1 2.1 The Unified PyObject & PyVarObject Structs

In CPython, every python value is a heap-allocated C struct. There are no primitive variables on the stack. The fundamental base structure is defined in `Include/object.h`:

```
1 /* CPython source definition from Include/object.h */
2 #ifndef Py_TRACE_REFS
3 /* Doubly-linked list pointers for active heap tracing in debug builds */
4 #define _PyObject_HEAD_EXTRA \
5     struct _object *_ob_next; \
6     struct _object *_ob_prev;
7 #else
8 #define _PyObject_HEAD_EXTRA
9 #endif
10
11 typedef struct _object {
12     _PyObject_HEAD_EXTRA           /* Doubly-linked list pointers */
13     Py_ssize_t ob_refcnt;         /* Reference count counter */
14     struct _typeobject *ob_type; /* Pointer to the type definition
15     object */
16 } PyObject;
17
18 typedef struct {
19     PyObject ob_base;             /* Base type fields */
20     Py_ssize_t ob_size;          /* Size of the variable portion (e.g.
21     sequence length) */
22 } PyVarObject;
```

1. Comparison to C++ Memory Layout

Unlike C++, where objects are laid out contiguously on the stack or heap based on their static type declarations and resolved via virtual function tables (`vtable`), CPython objects are entirely dynamic: 1. **Uniform Pointer Width**: Every variable in Python is a pointer (`PyObject*`), consuming exactly 8 bytes (on 64-bit platforms). 2. **No VTable Indirection**: Polymorphism is resolved dynamically by dereferencing the type pointer `ob_type` at runtime to lookup operations.

```
1 Python Reference Pointer:
2 [Name: x] ---> [ Heap Allocation: PyObject ]
3
4 | _PyObject_HEAD_EXTRA (Tracking) |
5 | ob_refcnt: 1 |
6 | ob_type: ----> [ PyTypeObject (int) ] |
7 | Integer Data Value Payload |
```

8 +-----+

4.0.2 2.2 The Type-Object Slot System

The type of an object is defined by the `PyTypeObject` struct (which is itself a subclass of `PyObject`).

1. Type Struct Slots layout

The type struct defines a series of “slots” (pointers to C functions) that determine how the object behaves when subjected to operators.

```

1  /* Conceptual CPython definition from Include/cpython/object.h */
2  struct _typeobject {
3      PyObject_VAR_HEAD
4      const char *tp_name;           /* For printing, in format <module>.<
   name> */
5      Py_ssize_t tp_basicsize, tp_itemsize; /* For allocation */
6
7      /* Type Slots (Static Function Pointers) */
8      destructor tp_dealloc;         /* Deallocation handler */
9      reprfunc tp_repr;             /* Representation builder */
10     getattrfunc tp_getattr;        /* Attribute lookup hook */
11     setattrfunc tp_setattr;        /* Attribute assignment hook */
12     hashfunc tp_hash;              /* Hashing calculation function
   pointer */
13
14     /* Protocol Table Slots */
15     PyNumberMethods *tp_as_number; /* Math function pointers (e.g. nb_add
   ) */
16     PySequenceMethods *tp_as_sequence; /* Indexing function pointers (e.g.
   sq_item) */
17     PyMappingMethods *tp_as_mapping; /* Map lookup function pointers (e.g.
   mp_subscript) */
18 };
19 typedef struct _typeobject PyTypeObject;

```

2. Dunder Methods to Slot Mapping

When defining a class in Python, the interpreter automatically maps custom dunder methods to these internal C slots: * `__init__` maps to `tp_init`. * `__new__` maps to `tp_new`. * `__hash__` maps to `tp_hash`. * `__add__` maps to `tp_as_number->nb_add`. * `__getitem__` maps to `tp_as_mapping->mp_subscript` OR `tp_as_sequence->sq_item`.

```

1  # Customizing slots via class definitions
2  class CustomObject:
3      def __init__(self, val):
4          self.val = val
5
6      def __hash__(self):
7          # Maps to the tp_hash C slot
8          return hash(self.val)
9
10 obj = CustomObject("test")
11 print("Hash via type slot:", hash(obj))

```

4.0.3 2.3 Reference Counting Lifecycle

CPython manages memory directly via reference counting. Reference counts are modified using C macros.

1. C-Level Reference Modification Macros

Defined in `object.h`, the macros expand as:

```

1 #define Py_INCREF(op) ( \
2     _Py_INC_REFTOTAL _Py_REF_DEBUG_EXTRA \
3     (op)->ob_refcnt++)
4
5 #define Py_DECREF(op) do { \
6     if (_Py_DEC_REFTOTAL _Py_REF_DEBUG_EXTRA \
7         --(op)->ob_refcnt == 0) \
8         _Py_Dealloc((PyObject *) (op)); \
9 } while (0)

```

To prevent segmentation faults when referencing a null pointer, CPython provides the NULL-safe macros `Py_XINCREF(op)` and `Py_XDECREF(op)`, which check if `(op != NULL)` before modifying the counts.

2. The Deallocation Pipeline

When `Py_DECREF(obj)` drops the reference count to zero: 1. **Invoke Deallocator:** The VM invokes `obj->ob_type->tp_dealloc(obj)`. 2. **Recursive Decrements:** The deallocator function decrements references of any nested objects inside this object. For example, if a list is deallocated, it decrements the reference counts of all items it contains. 3. **Release Memory:** The deallocator releases the raw memory back to `PyMalloc` or the system allocator.

3. Leak Tracking in C Extensions

In custom C extensions, failing to decrement references leads to memory leaks. In debug builds, compiling CPython with `Py_TRACE_REFS` tracks all active references in a global doubly-linked list. We can trace reference counts using the `sys` module:

```

1 import sys
2
3 # Reference counts tracking
4 my_list = [100, 200]
5 print("Starting references count:", sys.getrefcount(my_list) - 1)
6
7 ref_a = my_list
8 ref_b = my_list
9 print("References after bindings:", sys.getrefcount(my_list) - 1)
10
11 del ref_a
12 print("References after deleting one binding:", sys.getrefcount(my_list) - 1)

```

4.0.4 2.4 Built-in Caching and Interning Optimizations

To avoid allocating memory for common objects, CPython uses global caches:

1. Small Integer Cache

CPython pre-allocates an array of integer objects from -5 to 256 (NSMALLNEGINTS and NSMALLPOSINTS). * **Mechanism:** Any assignment in this range returns a pointer to the pre-existing static object, bypassing heap allocations. * **C-level array:** Inside objects/longobject.c, this is defined as:

```
1 static PyLongObject small_ints[NSMALLNEGINTS + NSMALLPOSINTS];
```

2. String Interning

String literals resembling identifiers are automatically **interned**. * **Mechanism:** Interned strings are stored in a global dictionary inside unicodeobject.c. If a string is already interned, any new instantiation returns the existing string pointer. * **Performance:** This allows matching strings via fast pointer comparison ($O(1)$ address checks) instead of byte-by-byte comparison ($O(N)$ string scans).

```
1 /* Under the hood string comparison */
2 if (str1 == str2) {
3     return 1; /* Match! (O(1) address check) */
4 }
```

3. Empty Collections and Singletons

To save memory: * **Empty Tuple Cache:** CPython allocates exactly one global empty tuple object (&Py_EmptyTuple). Every call to tuple() returns a pointer to this singleton. * **Built-in Singletons:** None, True, False, Ellipsis, and NotImplemented are statically allocated global structures.

```
1 # Small Integer Caching Check
2 int_a = 256
3 int_b = 256
4 print("Are 256 pointers identical?", int_a is int_b) # True (cached)
5
6 int_c = 257
7 int_d = 257
8 print("Are 257 pointers identical?", int_c is int_d) # False (dynamically
9     allocated)
10
11 # String Interning Check
12 str_a = "godhood_string"
13 str_b = "godhood_string"
14 print("Are literal strings interned?", str_a is str_b) # True
15
16 # Empty tuple check
17 tup_a = ()
18 tup_b = tuple()
19 print("Are empty tuples cached singletons?", tup_a is tup_b) # True
```

Chapter 5

Python 2.0 to 2.1: Comprehensions, Nested Scopes, & Cycle-Detecting GC

5.0.1 3.1 Python 2.0 List Comprehensions & Scope Leakage

Introduced in Python 2.0, list comprehensions provided a concise syntax to create lists from sequences. However, their early implementation in the compiler had a major design flaw: they leaked loop variables into the surrounding scope.

1. Inline Loop Compilation

In Python 2.x, a list comprehension did not create a new stack frame or lexical scope. Instead, it was compiled as an inline loop inside the host function. Consider the following Python 2.x list comprehension:

```
1 # Python 2.x Behavior
2 x = 999
3 data = [x for x in range(5)]
4 print(x) # Outputs 4! The variable x was overwritten.
```

2. Bytecode Analysis

If we trace the bytecode of this list comprehension under Python 2.x, we see the compiler emits direct modifications to the local frame variables array:

1	2	0 BUILD_LIST	0	
2		3 LOAD_GLOBAL	0 (range)	
3		6 LOAD_CONST	1 (5)	
4		9 CALL_FUNCTION	1	
5		12 GET_ITER		
6	>>	13 FOR_ITER	12 (to 28)	
7		16 STORE_FAST	0 (x)	# Direct write to host
	local variable x!			
8		19 LOAD_FAST	0 (x)	
9		22 LIST_APPEND	2	
10		25 JUMP_ABSOLUTE	13	
11	>>	28 STORE_FAST	1 (data)	

Because it uses `STORE_FAST 0 (x)`, the loop index directly overwrites the host function's local variable `x`.

3. Modern Python 3 Isolation

To resolve this leakage, Python 3.0 redefined list, set, and dictionary comprehensions to execute inside a compiler-generated nested function scope. This creates a separate `PyFrameObject` during evaluation, protecting the host function’s namespace from leakage.

5.0.2 3.2 PEP 227 Nested Scopes and LEGB Name Resolution

Before Python 2.1, CPython used a simple three-tier lookup model: **LGB (Local, Global, Built-in)**. * **The Limitation:** If a function was nested inside another function, the inner function could not access local variables defined in the outer function.

```

1 # Pre-Python 2.1 Behavior
2 def outer():
3     x = 10
4     def inner():
5         return x # Raised NameError! x was not in local, global, or builtins.

```

To access outer variables, developers had to pass them explicitly as default arguments (e.g., `def inner(x=x):`).

1. The LEGB Resolution Engine (PEP 227)

Introduced as an optional feature in Python 2.1 (enabled via `from __future__ import nested_scopes`) and made standard in Python 2.2, PEP 227 added the **Enclosing** scope to name resolution, establishing the modern **LEGB** lookup hierarchy:

```

1 [Name Resolution Lookup Path]
2 |
3 1. Local Scope (Array offset / LOAD_FAST)
4 |
5 2. Enclosing Scopes (Cell reference / LOAD_DEREF)
6 |
7 3. Global Scope (Module dict / LOAD_GLOBAL)
8 |
9 4. Built-in Scope (Builtins dict / LOAD_GLOBAL)

```

1. **Local (L):** Fast array offset reads inside the active `PyFrameObject`.
2. **Enclosing (E):** Searches cells (`PyCellObject`) containing enclosed variable pointers.
3. **Global (G):** Searches the module globals dictionary (`__dict__`).
4. **Built-in (B):** Searches the standard built-in namespace dictionary (`__builtins__`).

5.0.3 3.3 Closure Mechanics and Bytecode Compilation

When a function accesses variables from an enclosing scope, CPython compiles the enclosing variables as **free variables** and stores them on the heap inside `PyCellObject` structs.

1. Bytecode Compilation Trace

Consider the following nested scope implementation:

```

1 def parent(x):
2     def child():
3         return x
4     return child

```

The CPython compiler processes this nested structure using three specialized bytecode instructions: `LOAD_CLOSURE`, `MAKE_FUNCTION` (with flags), and `LOAD_DEREF`.

Let's trace the bytecode for both the parent and child functions:

Bytecode for parent:

1	2	0	<code>LOAD_CLOSURE</code>	0 (x)	<code>/* Load cell reference for cell variable x */</code>
2	2		<code>BUILD_TUPLE</code>	1	<code>/* Package cell reference inside a tuple */</code>
3	4		<code>LOAD_CONST</code>	1	<code>(<code object child>)</code>
4	6		<code>LOAD_CONST</code>	2	<code>('parent.<locals>.child')</code>
5	8		<code>MAKE_FUNCTION</code>	8	<code>/* Flag 8 tells the VM to bind closure cells tuple */</code>
6	10		<code>STORE_FAST</code>	1 (child)	
7					
8	4	12	<code>LOAD_FAST</code>	1 (child)	
9		14	<code>RETURN_VALUE</code>		

Bytecode for child:

1	3	0	<code>LOAD_DEREF</code>	0 (x)	<code>/* Dereference closure cell and load value */</code>
2		2	<code>RETURN_VALUE</code>		

2. Instruction Walkthrough

1. `LOAD_CLOSURE`: Pushes the cell reference representing the cell variable `x` onto the stack.
2. `MAKE_FUNCTION 8`: Creates a new function object (`PyFunctionObject`). The flag 8 (in modern CPython, this uses `MAKE_FUNCTION` with cell tuples on the stack) instructs the VM to pop the cells tuple and assign it to the function's `__closure__` attribute.
3. `LOAD_DEREF`: Executed inside `child`. The VM bypasses local dictionary or local array lookups. It accesses the closure tuple at index 0, dereferences the `PyCellObject`, and pushes its `ob_ref` payload onto the evaluation stack.

```
1 import inspect
2
3 closure_fn = parent(42)
4 # Inspect closure cells at runtime
5 print("Closure cells:", closure_fn.__closure__)
6 print("Cell value:", closure_fn.__closure__[0].cell_contents)
7 print("Parent cell vars:", parent.__code__.co_cellvars)
8 print("Child free vars:", closure_fn.__code__.co_freevars)
```

5.0.4 3.4 Python 2.0 Cycle-Detecting Generational Garbage Collector

Reference counting is deterministic and fast, but it cannot reclaim self-referencing cycles:

```
1 # Reference Cycle Example
2 a = []
3 b = []
4 a.append(b)
5 b.append(a)
6 del a, b # reference counts of both objects remain 1, leaking memory!
```

To solve this, Python 2.0 introduced a generational cycle-detecting Garbage Collector (GC) to periodically identify and sweep cycles.

1. Tracked Objects and GC Headers

Only container objects (lists, dictionaries, tuples, custom classes) can participate in reference cycles. Simple types like integers or strings cannot. CPython prefixes all tracked container allocations with a `PyGC_Head` header, linking them into a doubly-linked list.

2. Generational Strategy

The GC groups objects into three generations (Gen 0, Gen 1, Gen 2). Newly allocated objects enter Gen 0. If they survive a garbage collection pass, they are promoted to Gen 1, and eventually to Gen 2. This strategy is based on the **weak generational hypothesis**: most objects die young.

3. Cycle Isolation Algorithm

1. **Copy Reference Counts:** The GC copy-assigns the reference count of every object in the collection generation to its `gc_refs` field in the `PyGC_Head` header.
 2. **Trace and Subtract:** The GC traverses the reference pointers of every object in the list. If a target object is also in the list, the GC decrements the target's `gc_refs` value.
 3. **Analyze Remainders:**
 - If an object's `gc_refs` drops to zero, it is unreachable from outside the generation list (candidate for garbage).
 - If `gc_refs > 0`, the object is reachable from outside. The object and all objects it references are marked as reachable and promoted.
 4. **Deallocate:** The GC breaks cycles for remaining unreachable objects and frees their memory.
-

Chapter 6

Python 2.2 to 2.3: Type-Class Unification, Descriptors, & C3 MRO

6.0.1 4.1 Type-Class Unification: The Inception of New-Style Classes (PEP 252 & PEP 253)

Before Python 2.2, user-defined classes and built-in types existed as entirely separate entities under the hood. This architectural separation introduced severe language inconsistencies.

1. The Classic Class Inconsistency

In the classic class model (pre-Python 2.2):

- * **Classic Classes:** Created using the `class MyClass` : syntax without inheriting from any class. All user-defined classes were represented by a single C type, `PyClass_Type`, defined by the `PyClassObject` struct.
- * **Classic Instances:** Every instance of any user-defined class was represented by a single C type, `PyInstance_Type`, mapped to a `PyInstanceObject` struct.
- * **The Dichotomy:** The `type()` of any classic class instance always returned `<type 'instance'>`, rather than the class itself. The actual class of the instance was accessible only via the `__class__` attribute: “python # Python 1.5 - 2.1 Classic Class Behavior

```
class Classic: pass
```

```
c = Classic() print(type(c)) # <type 'instance'> print(c.__class__) # <class 'main.Classic'>
```

“* **Built-in Types**”: In contrast, built-in types (e.g., `int`, `list`, `dict`, `str`) were instances of `PyType_Type`, represented in C by a static `PyTypeObject` struct. * **The Subclassing Barrier**”: Because of this dichotomy, user-defined classes could not inherit from built-in types. Subclassing `list` or `dict` to extend their behavior was impossible, because the classic class object engine could not parse or manage the memory layouts of built-in C types.

2. The Solution: New-Style Classes and object

Introduced in Python 2.2 by PEP 252 and PEP 253, **new-style classes** unified types and classes.

- * **The Unified Base:** A new-style class is defined by inheriting (directly or indirectly) from the built-in `object` type (the root of the unified type tree).
- * **Type Unification:** Under the unified model, user-defined classes are themselves type objects (instances of `type`), just like built-in types. Checking `type(c)` now returns the class itself. “python # Python 2.2+ New-Style Class Behavior

```
class NewStyle(object): pass
```

```
n = NewStyle() print(type(n)) # <class 'main.NewStyle'> ““
```

3. Low-Level C Representation: `PyTypeObject`

Every class (built-in or new-style user-defined) is represented at the C-level as an instance of `PyTypeObject`. Let's inspect the core fields in CPython's `Include/object.h`:

```

1 typedef struct _typeobject {
2     PyObject_VAR_HEAD
3     const char *tp_name;           /* For printing, in format "<module>.<
4     Py_ssize_t tp_basicsize, tp_itemsize; /* For allocation sizes */
5
6     /* Methods to implement standard operations */
7     destructor tp_dealloc;
8     printfunc tp_print;
9     getattrofunc tp_getattr;
10    setattrofunc tp_setattr;
11
12    /* Attribute lookup slot */
13    getattrofunc tp_getattro;       /* Pointing to PyObject_GenericGetAttr
14    setattrofunc tp_setattro;       /* Pointing to PyObject_GenericSetAttr
15
16    /* Protocol slot mappings */
17    PyNumberMethods *tp_as_number;
18    PySequenceMethods *tp_as_sequence;
19    PyMappingMethods *tp_as_mapping;
20
21    /* Inheritance and lookup structures */
22    PyObject *tp_dict;              /* Namespace dictionary */
23    PyObject *tp_bases;             /* Tuple of base classes */
24    PyObject *tp_mro;               /* Method Resolution Order tuple */
25
26    /* Allocation / Initialization slots */
27    newfunc tp_new;                 /* __new__ allocation entry */
28    initproc tp_init;               /* __init__ initialization entry */
29    allocfunc tp_alloc;             /* Low-level memory allocator */
30
31    /* Flags and inheritance info */
32    unsigned long tp_flags;
33    struct _typeobject *tp_base;    /* Direct base pointer */
34 } PyTypeObject;

```

4. Slot Wrapper Descriptors

Because C slots (like `tp_init`) expect standard C function signatures (e.g., `int (*tp_init)(PyObject *, PyObject *, PyObject *)`), but user-defined classes write Python methods (`def __init__(self, ...)`), CPython utilizes **slot wrappers** to bridge the boundary: * **Python to C (Slot Fillers)**: When a new-style class defines a Python method like `__init__`, the type compiler dynamically wraps this method in a C function (such as `slot_tp_init`) and assigns it to the type object's `tp_init` slot. When `tp_init` is called, `slot_tp_init` executes the Python function. * **C to Python (Wrapper Descriptors)**: To expose built-in C slots (like object's default initialization code) to Python, CPython wraps them in descriptor objects (e.g. `__init__` is exposed as `<slot wrapper '__init__' of 'object' objects>`).

6.0.2 4.2 The Descriptor Protocol & Attribute Lookup Chain

The descriptor protocol is the underlying mechanism that enables properties, methods, class-methods, and staticmethods in Python.

1. The Descriptor Protocol Definition

A descriptor is an object that implements at least one of the three descriptor protocol methods:

```

1 def __get__(self, instance, owner=None):
2     """Invoked when getting the attribute."""
3     pass
4
5 def __set__(self, instance, value):
6     """Invoked when setting the attribute."""
7     pass
8
9 def __delete__(self, instance):
10    """Invoked when deleting the attribute."""
11    pass

```

- **Data Descriptor:** Implements both `__get__` AND `__set__` (and/or `__delete__`).
- **Non-Data Descriptor:** Implements only `__get__` (e.g. functions, classmethods, staticmethods).

2. The Attribute Lookup Algorithm (PyObject_GenericGetAttr)

When an attribute is accessed (`obj.name`) on a new-style instance `obj` of class `Class`, CPython invokes the `tp_getattro` slot of the type object. By default, this points to `PyObject_GenericGetAttr` in `Objects/object.c`.

The lookup follows this strict resolution flow: 1. **MRO Search:** CPython calls `_PyType_Lookup(Class, name)`. This searches the class's namespace dictionary `tp_dict` and the namespace dictionaries of all parent classes in the Method Resolution Order (`tp_mro`). Let the resolved object be `descr`. 2. **Data Descriptor Check:** If `descr` is found, and its type implements the `__set__` (or `__delete__`) slot (`descr->ob_type->tp_descr_set` is not NULL): * Call `descr->ob_type->tp_descr_get(descr, obj, Class)` and return the result immediately. 3. **Instance Dict Lookup:** Check the instance dictionary of `obj` (the `__dict__` table at `obj->ob_dict`). If `name` exists in the instance dictionary, return its associated value. 4. **Non-Data Descriptor Check:** If `descr` is found: * If `descr` has a `__get__` slot (`descr->ob_type->tp_descr_get` is not NULL): * Call `descr->ob_type->tp_descr_get(descr, obj, Class)` and return the result. * If `descr` does not have a `__get__` slot (e.g. a plain class attribute like a string or int): * Return `descr` directly. 5. **Fallback to `__getattr__`:** If the attribute is still not resolved, CPython raises an `AttributeError`, which triggers the invocation of the fallback `__getattr__(self, name)` method if defined on the class.

```

1  [Start Lookup: obj.name]
2      |
3      v
4  [Search MRO for "name"]
5      |
6  +-----+-----+
7  |         |         |
8  (Not Found)       (Found)
9  |         |         |
10 |         |         v
11 |         |         [Is it a Data Descriptor?]
12 |         |         (has __get__ and __set__)
13 |         |         |
14 |         |         +-----+-----+
15 |         |         | Yes           | No
16 |         |         |         |         |
17 |         |         |         |         |
18 |         |         |         |         |
19 |         |         |         |         |
20 +-----+-----> [Check obj.__dict__]

```

```

21      |
22      +-----+-----+
23      | Found          | Not Found
24      v              |
25      [Return dict val] v
26      [Is it a Non-Data Descriptor?]
27      (has __get__ but no __set__)
28      |
29      +-----+-----+
30      | Yes          | No
31      v              v
32      [Call __get__]  [Is it class attribute?]
33      [Return result] |
34      +-----+-----+
35      | Yes          | No
36      v              v
37      [Return value]  [Raise AttributeError]
38                      |
39                      [Call __getattr__]

```

3. C-Level descriptor specifications: PyGetSetDef

In C extensions, properties are frequently defined using the `PyGetSetDef` struct in the class's `tp_getset` slot:

```

1 typedef struct PyGetSetDef {
2     const char *name;
3     getter get;          /* C function: PyObject *(*getter)(PyObject *, void
4                          *) */
5     setter set;          /* C function: int (*setter)(PyObject *, PyObject
6                          *, void *) */
7     const char *doc;
8     void *closure;       /* Context pointer passed to getter/setter */
9 } PyGetSetDef;

```

4. Practical Implementation of Core Decorators

Here is how decorators leverage descriptors to modify binding behavior: * **Bound Methods:** Python functions are non-data descriptors. When accessed via `obj.method`, `FunctionType.__get__(func, obj, Class)` is called. It returns a `PyMethod` object wrapping the function and the instance:

`obj.method` $\equiv$ `Class.method.__get__(obj, Class)`

* **@classmethod:** Implements `__get__(self, instance, owner)`. When accessed, it ignores the `instance` argument and returns a bound method wrapping the function and the `owner` (the class object itself). * **@staticmethod:** Implements `__get__(self, instance, owner)`. It returns the underlying raw function object directly, bypassing any method binding. * **@property:** A data descriptor that wraps getter, setter, and deleter functions: “python class CustomProperty(object): def **init**(self, fget, fset=None): self.fget = fget self.fset = fset

```

1 def __get__(self, instance, owner):
2     if instance is None:
3         return self
4     return self.fget(instance)
5
6 def __set__(self, instance, value):
7     if self.fset is None:
8         raise AttributeError("can't set attribute")
9     self.fset(instance, value)

```

```

1 ---
2
3 ### 4.3 Method Resolution Order (MRO) & C3 Linearization
4
5 Method Resolution Order determines how Python traverses the inheritance tree
  during attribute search.
6
7 ##### 1. The Classic DFLR Algorithm & The Diamond Problem
8 Classic classes (pre-Python 2.2) resolved attributes using a **Depth-First,
  Left-to-Right (DFLR)** tree traversal.
9 In a diamond inheritance hierarchy:

```

```

1  A
2  / \

```

B C / D

```

1 The declaration 'class D(B, C)' inherits from 'B' and 'C', both of which
  inherit from 'A'.
2 * DFLR path: '[D, B, A, object, C, A, object]'.
3 * Removing duplicates keeping only the **first** occurrence gives: '[D, B, A,
  object, C]'.
4 * **The Failure**: If 'A' defines a method 'method()' and 'C' overrides it,
  calling 'D().method()' resolves to 'A.method()' rather than 'C.method()'.
  This violates the principle that specialized subclasses ('C') should
  override generic ancestors ('A').
5
6 ##### 2. Python 2.2 MRO: The Last Occurrence Rule & Monotonicity Failure
7 To fix this, Python 2.2 introduced a new-style MRO calculation:
8 1. Perform DFLR traversal of the class and all its ancestors.
9 2. Remove all duplicates except the **last** occurrence.
10
11 For the diamond hierarchy 'D(B, C)':
12 * Traversal: '[D, B, A, object, C, A, object]'.
13 * Keeping only the **last** occurrence: '[D, B, C, A, object]'.
14 While this solved the diamond problem, the algorithm was **non-monotonic**.
15 * **Monotonicity**: If class  $XX$  precedes  $YY$  in the MRO of class  $PP$ , then  $XX$ 
  must precede  $YY$  in the MRO of any subclass  $SS$  derived from  $PP$ .
16
17 ##### Samuele Pedroni's Monotonicity Violation Example:
18 Consider this class configuration in Python 2.2:
19 '''python
20 class A(object): pass
21 class B(object): pass
22 class C(object): pass
23 class D(object): pass
24 class E(object): pass
25
26 class K1(A, B, C): pass
27 class K2(D, B, E): pass
28 class K3(D, A): pass
29
30 class Z(K1, K2, K3): pass

```

Let's calculate the Python 2.2 MRO for $K1$ and $K2$: * $K1$ raw: $[K1, A, object, B, object, C, object]$. Keeping last: $[K1, A, B, C, object]$. (Here, A precedes B). * $K2$ raw: $[K2, D, object, B, object, E, object]$. Keeping last: $[K2, D, B, E, object]$. (Here, B precedes E). * $K3$ raw: $[K3, D, object, A, object]$. Keeping last: $[K3, D, A, object]$.

Now let's resolve $Z(K1, K2, K3)$ under Python 2.2: * Raw DFLR list: $[Z, K1, A, object, B, object, C, object, K2, D, object, B, object, E, object, K3, D, object, A, object]$. * Keeping only the last occurrence: $[Z, K1, C, K2, K3, D, B, E, A, object]$. * **The Monotonicity Fail-**

ure: * In the parent class κ_1 's MRO, A preceded B . * In the child class z 's MRO, B precedes A ($\dots D, B, E, A \dots$). * The child class has reversed the relative order of A and B established in the parent, violating monotonicity.

3. Python 2.3+ MRO: The C3 Linearization Algorithm

To guarantee monotonicity and local precedence ordering, Python 2.3 adopted **C3 Linearization**.

Mathematical Formulation: Let $L(C)$ be the linearization (MRO) of class C . For a class C inheriting from direct parents $B_1, B_2, \dots, B_N$:

$$L(C) = [C] + \text{merge}(L(B_1), L(B_2), \dots, L(B_N), [B_1, B_2, \dots, B_N])$$

Where $L(\text{object}) = [\text{object}]$.

The Merge Operation:

1. Examine the head (index 0) of the first list inside the merge block: $H = L(B_1)[0]$.
2. If H does not appear in the **tail** (index 1 to the end) of any other list in the merge block, it is a **good head**.
 - Append H to the linearization of C .
 - Remove H from all lists in the merge block.
 - Repeat the merge step.
3. If H appears in the tail of any other list, it is not a good head. Move to the next list in the merge block and check its head.
4. If no candidate head can be selected across all lists, the merge is impossible. Python raises a `TypeError`.

Step-by-Step Mathematical Calculation of $Z(\kappa_1, \kappa_2, \kappa_3)$: Let's resolve the MRO of class z from the Pedroni example using C3 Linearization.

We have the parent linearizations:

$$L(K1) = [K1, A, B, C, \text{object}]$$

$$L(K2) = [K2, D, B, E, \text{object}]$$

$$L(K3) = [K3, D, A, \text{object}]$$

Now calculate $L(Z)$:

$$L(Z) = [Z] + \text{merge}(L(K1), L(K2), L(K3), [K1, K2, K3])$$

$$L(Z) = [Z] + \text{merge}([K1, A, B, C, \text{obj}], [K2, D, B, E, \text{obj}], [K3, D, A, \text{obj}], [K1, K2, K3])$$

Step 1: Check head of first list: κ_1 . * Does κ_1 appear in the tail of $[K2, D, B, E, \text{obj}]$, $[K3, D, A, \text{obj}]$, or $[K1, K2, K3]$? No. * Extract κ_1 :

$$L(Z) = [Z, K1] + \text{merge}([A, B, C, \text{obj}], [K2, D, B, E, \text{obj}], [K3, D, A, \text{obj}], [K2, K3])$$

Step 2: Check head of first list: A . * Does A appear in the tail of other lists? Yes, it appears in the tail of $[K3, D, A, \text{obj}]$. Skip A . * Move to the next list head: κ_2 . * Does κ_2 appear in the tail of other lists? No (it only appears at the head of the last list $[K2, K3]$). * Extract κ_2 :

$$L(Z) = [Z, K1, K2] + \text{merge}([A, B, C, \text{obj}], [D, B, E, \text{obj}], [K3, D, A, \text{obj}], [K3])$$

Step 3: Check head of first list: **A**. * Does **A** appear in the tail of other lists? Yes, in the tail of $[K3, D, A, \text{obj}]$. Skip. * Check head of second list: **D**. * Does **D** appear in the tail of other lists? Yes, in the tail of $[K3, D, A, \text{obj}]$. Skip. * Check head of third list: **K3**. * Does **K3** appear in the tail of other lists? No. * Extract **K3**:

$$L(Z) = [Z, K1, K2, K3] + \text{merge}([A, B, C, \text{obj}], [D, B, E, \text{obj}], [D, A, \text{obj}])$$

Step 4: Check head of first list: **A**. * Does **A** appear in the tail of other lists? Yes, in the tail of $[D, A, \text{obj}]$. Skip. * Check head of second list: **D**. * Does **D** appear in the tail of other lists? No (it is at the head of $[D, A, \text{obj}]$). * Extract **D**:

$$L(Z) = [Z, K1, K2, K3, D] + \text{merge}([A, B, C, \text{obj}], [B, E, \text{obj}], [A, \text{obj}])$$

Step 5: Check head of first list: **A**. * Does **A** appear in the tail of other lists? No (it is at the head of $[A, \text{obj}]$). * Extract **A**:

$$L(Z) = [Z, K1, K2, K3, D, A] + \text{merge}([B, C, \text{obj}], [B, E, \text{obj}], [\text{obj}])$$

Step 6: Check head of first list: **B**. * Does **B** appear in the tail of other lists? No. * Extract **B**:

$$L(Z) = [Z, K1, K2, K3, D, A, B] + \text{merge}([C, \text{obj}], [E, \text{obj}], [\text{obj}])$$

Step 7: Check head of first list: **C**. * Does **C** appear in the tail of other lists? No. * Extract **C**:

$$L(Z) = [Z, K1, K2, K3, D, A, B, C] + \text{merge}([\text{obj}], [E, \text{obj}], [\text{obj}])$$

Step 8: Check head of first list: **obj**. * Does **obj** appear in the tail of other lists? Yes, in the tail of $[E, \text{obj}]$. Skip. * Check head of second list: **E**. * Does **E** appear in the tail of other lists? No. * Extract **E**:

$$L(Z) = [Z, K1, K2, K3, D, A, B, C, E] + \text{merge}([\text{obj}], [\text{obj}], [\text{obj}])$$

Step 9: Extract **obj**:

$$L(Z) = [Z, K1, K2, K3, D, A, B, C, E, \text{object}]$$

This calculation resolves the Method Resolution Order cleanly, preserving both local precedence and global monotonicity across all classes.

6.0.3 4.4 Instance Lifecycle: Allocation vs. Initialization

CPython splits object creation into two phases, controlled by different type slots: 1. **Allocation** (`__new__`): * Maps to the `tp_new` slot in `PyTypeObject`. * This is a static method responsible for allocating memory on the heap (using `PyType_GenericNew()`) and initializing the object's `PyObject` headers (`ob_refcnt` and `ob_type`). * It must return a new instance of the class (or subclass). 2. **Initialization** (`__init__`): * Maps to the `tp_init` slot in `PyTypeObject`. * This is an instance method called immediately after `__new__` returns a valid instance. It populates the instance dictionary `__dict__` with fields.

Chapter 7

Python 2.4 to 2.7: Decorators, Context Managers, & the 2.x Twilight

7.0.1 5.1 Decorators: Syntactic Sugar and Compiler Mechanics (PEP 318 & PEP 3129)

Prior to Python 2.4, wrapping a function with utility behavior required declaring the function first, then immediately reassigning it in the host namespace:

```
1 def query():
2     pass
3 query = transaction(log(query))
```

This approach separated the metadata from the function declaration. PEP 318 introduced the `@decorator` syntax to resolve this.

1. Compiler Translation & Bytecode Analysis

The `@` symbol is syntactic sugar resolved at compile-time. When the compiler encounters decorators, it loads the decorators onto the evaluation stack *before* compiling the target function, then applies them in bottom-up (reverse-textual) order.

Consider this nested decorator structure:

```
1 @dec1
2 @dec2
3 def target():
4     pass
```

The compiler translates this structure into the following bytecode operations:

1	1	0	LOAD_NAME	0 (dec1)	/* Push dec1 onto stack */
	*/				
2		3	LOAD_NAME	1 (dec2)	/* Push dec2 onto stack */
	*/				
3		6	LOAD_CONST	0 (<code object target>)	
4		9	LOAD_CONST	1 ('target')	
5		12	MAKE_FUNCTION	0	/* Create function object target */
6		15	CALL_FUNCTION	1	/* Call dec2(target) */
7		18	CALL_FUNCTION	1	/* Call dec1(dec2(target)) */
8		21	STORE_NAME	2 (target)	/* Bind result to name target */

2. Decorator Execution Flow

1. The compiler pushes the decorator references `dec1` and `dec2` onto the evaluation stack.
2. `MAKE_FUNCTION` instantiates a `PyFunctionObject` from the compiled <code object `target`> and pushes it onto the stack.
3. The VM executes `CALL_FUNCTION 1`, popping `target` as the argument and `dec2` as the callable, pushing the returned wrapped function onto the stack.
4. The VM executes the second `CALL_FUNCTION 1`, popping the wrapped function and `dec1` as the callable, pushing the final wrapped function onto the stack.
5. `STORE_NAME` binds the resulting callable to the name `target` in the local namespace.

This bottom-up composition is mathematically equivalent to:

$$\text{target} = \text{dec1}(\text{dec2}(\text{target}))$$

3. Class Decorators (PEP 3129)

Introduced in Python 2.6, class decorators apply the same syntactic translation to class creation. After the class body is evaluated and the class type object is constructed, the compiler passes the type object to the decorator:

```
1 @class_decorator
2 class Model(object):
3     pass
```

This translates directly to:

$$\text{Model} = \text{class_decorator}(\text{Model})$$

7.0.2 5.2 Context Managers and the with Statement (PEP 343)

Introduced in Python 2.5, the `with` statement encapsulates clean acquisition and release patterns for resources.

1. The Context Manager Protocol

An object is a context manager if it implements: * `__enter__(self)`: Acquires the resource and returns the object to be bound to the `as target` (if present). * `__exit__(self, exc_type, exc_val, exc_tb)`: Invoked when leaving the block. * If an exception was raised, `exc_type`, `exc_val`, and `exc_tb` contain the exception details. If `__exit__` returns a truthy value, the exception is silenced. * If no exception occurred, all three arguments are passed as `None`.

2. Compiler Translation to Try-Finally

The compiler translates a `with` statement block:

```
1 with expression as target:
2     suite
```

Into a low-level equivalent logic block:

```
1 mgr = expression
2 exit_method = type(mgr).__exit__
3 value = type(mgr).__enter__(mgr)
4 exc = True
5 try:
```

```

6         try:
7             target = value # Bind target if "as" clause exists
8             suite
9         except:
10            exc = False
11            if not exit_method(mgr, *sys.exc_info()):
12                raise
13    finally:
14        if exc:
15            exit_method(mgr, None, None, None)

```

3. Bytecode Implementation and the Block Stack

To guarantee that `__exit__` is called even when exceptions propagate, the VM uses its internal **block stack** inside the execution frame.

When the VM executes a `with` block:

1. **SETUP_WITH**: Evaluates the context manager expression, calls `__enter__`, pushes the `__exit__` method onto the evaluation stack, pushes a `finally` block onto the frame's block stack, and pushes the `__enter__` return value onto the evaluation stack.
2. **WITH_CLEANUP** (or modern equivalent block unwinding): When leaving the `with` block (normally or via an exception), the VM pops the block off the block stack.
 - * If an exception propagates, the VM leaves the exception (`type`, `value`, `traceback`) on the evaluation stack and calls `__exit__(type, value, traceback)`.
 - * If `__exit__` returns `True`, the VM clears the exception state, preventing propagation. If `False`, the VM re-throws the exception.

7.0.3 5.3 Generator Enhancements: Coroutines and Frame Suspension (PEP 342)

Introduced in Python 2.5, PEP 342 expanded generators from simple, passive data producers into active coroutines by allowing bidirectional communication.

1. Bidirectional Methods

- `.send(value)`: Resumes the generator and passes `value` back. Inside the generator, the active yield expression evaluates to this `value`. Calling `.send(None)` is equivalent to `next()`.
- `.throw(type, value=None, traceback=None)`: Raises the specified exception inside the generator frame at the suspended yield point.
- `.close()`: Raises a `GeneratorExit` exception at the suspended yield point. The generator must clean up resources and exit. If it yields another value instead of returning or exiting, the VM raises a `RuntimeError`.

2. Frame Suspension Mechanics (PyFrameObject)

When a generator encounters a `yield` statement, the execution loop in `_PyEval_EvalFrameDefault()` suspends execution:

```

1 [Active Thread State] -> [Running Generator Frame] -> Yield Encountered
2 |
3 v
4 1. Save instruction pointer: frame->f_lasti = active_instruction_offset
5 2. Save evaluation stack pointer depth: frame->f_stackdepth = current_depth
6 3. Set generator state: generator->gi_state = GEN_SUSPENDED
7 4. Detach frame: thread_state->frame = frame->f_back
8 |
9 v
10 [Return Control to Caller]

```

1. **f_lasti Preservation:** The VM saves the offset of the next instruction in `frame->f_lasti`.
2. **Stack Conservation:** The evaluation stack pointer is frozen at its current depth.
3. **State Transition:** The generator's state is set to `GEN_SUSPENDED`.
4. **Frame Detachment:** The frame is unlinked from the active thread state execution chain, returning control back to the caller while keeping the frame alive on the heap.

7.0.4 5.4 Under-the-Hood CPython Data Structure Layouts

1. Python List Layout (PyListObject)

A Python list is a contiguous dynamic array of `PyObject*` pointers. Because list sizes change dynamically, CPython overallocates memory blocks to achieve $O(1)$ amortized append performance.

Let's inspect the `PyListObject` struct defined in `Include/listobject.h`:

```
1 typedef struct {
2     PyObject_VAR_HEAD
3     PyObject **ob_item;           /* Vector of pointers to list items */
4     Py_ssize_t allocated;        /* Number of slots allocated in ob_item memory */
5 } PyListObject;
```

When a list grows beyond its current capacity, CPython resizes the underlying array using the formula:

$$\text{allocated} = \text{newsize} + (\text{newsize} \gg 3) + (\text{newsize} < 9 ? 3 : 6)$$

This formula balances memory overhead and reallocation speed. By adding $\approx 12.5\%$ extra slots as the list grows, it minimizes heap reallocations and memory copying costs.

Here is the resulting capacity growth trace: | Item Count (*newsize*) | Allocated Capacity |
 Overallocation Factor | — | — | — | 0 | 0 | - | 1 | 4 | 400% | 5 | 8 | 160% | 9 | 16 | 177% |
 17 | 25 | 147% | 1000 | 1129 | 112.9% |

2. Classic Sparse Dictionary Layout (Pre-Python 3.6)

Before Python 3.6 (PEP 468), dictionaries were sparse hash tables consisting of an array of `PyDictEntry` structs.

The entry struct `PyDictEntry` was defined in `Include/dictobject.h` as:

```
1 typedef struct {
2     Py_ssize_t me_hash;           /* Cached hash value of me_key */
3     PyObject *me_key;            /* Pointer to the key PyObject */
4     PyObject *me_value;          /* Pointer to the value PyObject */
5 } PyDictEntry;
```

The dictionary header `PyDictObject` was defined as:

```
1 struct _dictobject {
2     PyObject_HEAD
3     Py_ssize_t ma_fill;          /* Active entries + dummy entries */
4     Py_ssize_t ma_used;          /* Active entries only */
5     Py_ssize_t ma_mask;          /* size of table - 1 */
6     PyDictEntry *ma_table;       /* Pointer to the sparse entry table array */
7 };
```

This layout was memory-inefficient: * Every slot in the table array consumed 24 bytes (on 64-bit systems), regardless of whether it contained an entry or was empty/deleted (dummy). * To keep search collisions low, the table size was always a power of 2 and kept at least 1/3 empty. * Memory was highly sparse, leading to poor cache locality.

```

1 Classic Sparse Dictionary Memory Layout:
2 [ Index 0: <me_hash, key_ptr, value_ptr> (24 bytes) ]
3 [ Index 1: <0, NULL, NULL> (24 bytes - null padding) ]
4 [ Index 2: <me_hash, key_ptr, value_ptr> (24 bytes) ]
5 [ Index 3: <0, NULL, NULL> (24 bytes - null padding) ]

```

3. Sets Optimization (PySetObject)

Python sets are implemented as open-addressed hash tables similar to dictionaries, but without values. Lookups skip value retrieval entirely, evaluating only keys and hashes in tight C loops. A set's entries are instances of `PySetEntry`:

```

1 typedef struct {
2     PyObject *key;
3     long hash;           /* Cached hash value of key */
4 } PySetEntry;

```

4. Tuples Optimization (PyTupleObject)

Tuples are immutable sequences stored as a single contiguous memory block containing the type header, the item count, and an array of `PyObject*` pointers:

```

1 typedef struct {
2     PyObject_VAR_HEAD
3     PyObject *ob_item[1]; /* Inline array of pointers (allocated dynamically) */
4 } PyTupleObject;

```

- **Tuple Free Lists:** To avoid the overhead of the system allocator, CPython maintains an array of free lists for small tuples up to size 20. When a tuple under size 20 is deallocated, its memory block is cached in the corresponding size slot of the free list for immediate reuse during subsequent tuple allocations.

Chapter 8

Python 2.x: Low-Level File I/O & Exceptions Unwinding Blocks

8.0.1 6.1 Low-Level File Management and Standard Streams wrapping

Under Python 2.x, the built-in `file` type and `open()` function were thin wrappers around standard C stdio library streams (`FILE *`).

1. The CPython 2.x PyFileObject Layout

Let's inspect the `PyFileObject` C structure defined in CPython 2.x's `Include/fileobject.h`:

```
1 typedef struct {
2     PyObject_HEAD
3     FILE *f_fp;           /* Pointer to standard C stdio stream */
4     PyObject *f_name;     /* File name string object */
5     PyObject *f_mode;     /* File mode string object */
6     int (*f_close)(FILE *); /* Close wrapper function */
7     int f_softspace;      /* Print statement space tracking state */
8     int f_binary;        /* Binary mode flag indicator */
9     PyObject *f_encoding; /* File encoding string */
10    PyObject *f_errors;   /* Error handling mode */
11    PyObject *f_newlines; /* Decoded newlines object */
12 } PyFileObject;
```

2. The `f_softspace` Printing Flag

In Python 2.x, the `print` statement was keyword-based rather than a function. The runtime tracked layout separation using the `f_softspace` slot inside the standard output stream: * When a `print` statement completed printing an item, it evaluated the next token. If a trailing comma was present (e.g., `print x,`), the VM set `f_softspace = 1` and omitted the newline character. * During the subsequent `print` operation, CPython checked `f_softspace`. If it was 1, it wrote a space to the file stream before printing the next element and reset `f_softspace = 0`.

3. Redundant Locking Bottlenecks under the GIL

Because Python 2.x wrapped standard C `FILE *` streams, I/O operations were subject to the underlying C library's internal buffering and synchronization locking. * Standard C libraries acquire internal mutex locks per read/write call to ensure multi-threaded safety. * In CPython, since the Global Interpreter Lock (GIL) already guarantees single-threaded interpreter execution, these nested C stdio locks were redundant, introducing unnecessary CPU context-switching overhead and locking bottlenecks during high-frequency concurrent file operations.

4. The Layered I/O System (PEP 3116)

To eliminate stdio bottlenecks, Python 3.0 replaced the old stream-based system with a layered I/O architecture backported to Python 2.6+. This system interacts directly with OS file descriptors via system calls (e.g., `read(2)`, `write(2)`), bypassing C's stdio library:

1. **Raw Layer** (`RawIOBase` / `FileIO`): A thin wrapper over the POSIX OS file descriptor, executing raw system-level reads and writes.
2. **Buffered Layer** (`BufferedIOBase` / `BufferedReader` / `BufferedWriter`): Maintains an internal memory block (typically 8KB) to minimize context switches between user space and kernel space.
3. **Text Layer** (`TextIOBase` / `TextIOWrapper`): Handles encoding/decoding and system-specific universal newline translations.

8.0.2 6.2 Exception Mechanics and the Thread State slots

CPython manages exceptions at the thread level, storing active and caught exception states inside the execution thread's state structure.

1. Exception Slots in PyThreadState

CPython maintains exception metadata on a per-thread basis within the `PyThreadState` struct (`Include/pystate.h`):

```

1 typedef struct _ts {
2     struct _ts *next;
3     PyInterpreterState *interp;
4     struct _frame *frame;          /* Active execution frame */
5
6     /* Exception currently propagating */
7     PyObject *curexc_type;
8     PyObject *curexc_value;
9     PyObject *curexc_traceback;
10
11     /* Exception currently caught and handled */
12     PyObject *exc_type;
13     PyObject *exc_value;
14     PyObject *exc_traceback;
15 } PyThreadState;

```

- **curexc_*** (**Current Exception**): The active exception currently propagating. When an instruction fails (e.g., division by zero), CPython sets these pointers and returns `NULL` up the C execution chain.
- **exc_*** (**Caught Exception**): The exception currently being handled inside an active except block. This keeps the exception accessible via `sys.exc_info()` during nesting, while preventing active propagation from overwriting the caught state.

2. The String Exception Era

In early Python versions (pre-2.6), exceptions were not required to be class instances; you could raise plain string literals:

```

1 # Pre-Python 2.6 Behavior
2 raise "DatabaseConnectionFailed"

```

During lookup, the VM evaluated string exceptions by identity (`is` check) rather than class inheritance. This made hierarchy categorization and error composition difficult. Python 2.6 deprecated and removed string exceptions, requiring all exceptions to inherit from the built-in base type `BaseException`.

8.0.3 6.3 The Try-Except-Finally Block Stack and Stack Restoration

CPython handles structured control flow (loops, exception blocks) using a static **block stack** located inside each execution frame (`PyFrameObject`).

1. The PyTryBlock Structure

The block stack consists of up to 20 (`CO_MAXBLOCKS`) `PyTryBlock` structs:

```
1 typedef struct {
2     int b_type;           /* Block type: SETUP_LOOP, SETUP_EXCEPT,
3     SETUP_FINALLY */
4     int b_handler;        /* Target instruction offset of the handler */
5     int b_level;          /* Stack pointer depth at block entry */
6 } PyTryBlock;
```

2. VM Exception Unwinding Flow

When an instruction returns `NULL`, the VM loop (`_PyEval_EvalFrameDefault`) enters unwinding mode:

```
1 [Exception Raised: C API returns NULL]
2     |
3     v
4 [Is frame->f_blockstack empty?]
5     |           |
6     Yes        No
7     |           |
8     v           v
9 [Pop Frame]    [Pop Try Block]
10 [Traceback]   |
11             [Is block SETUP_EXCEPT or SETUP_FINALLY?]
12             |           |
13             No          Yes
14             |           |
15             v           v
16             [Continue loop]  1. Restore stack pointer: stack_pointer =
17                                frame->b_level
18                                2. Push traceback, value, type onto stack
19                                3. Set PC: frame->f_lasti = block->b_handler
20                                4. Resume bytecode execution in handler
```

1. **Stack Cleanup via `b_level`:** The VM pops the current try block. It immediately restores the frame's evaluation stack pointer back to the offset saved in `b_level`. This discards any unused variables or intermediate states created inside the `try` block, preventing memory leaks.
2. **Context Setup:** The VM pushes the exception's `traceback`, `value`, and `type` (in Python 2.x) onto the evaluation stack.
3. **Frame Traversal fallback:** If the frame's block stack is exhausted without finding a handler, CPython pops the frame, instantiates a `PyTracebackObject` mapping the exception to the current line number, and continues unwinding in the caller's frame (`frame->f_back`).

Part II

Volume II: The Python 3 Schism & Core Enhancements

Chapter 9

Python 3.0: The Unicode Paradigm Shift and Text vs. Bytes Separation

9.0.1 7.1 The Unicode Paradigm Shift and Text vs. Bytes Separation (PEP 358 & PEP 3112)

Python 3.0 introduced a structural boundary between textual characters and binary data. In Python 2.x, `str` served double-duty as both raw bytes and text, causing silent, non-deterministic bugs when ASCII decoding failed implicitly during concatenation. Python 3.0 solved this by establishing a strict boundary between `str` (text) and `bytes`/`bytearray` (binary).

1. The Python 2.x Concatenation Trap and Coercion Ban

In Python 2.x, mixed operations between `str` (raw 8-bit characters) and `unicode` (arbitrary code points) were implicitly resolved. When executing `"rsum" + u" (French)"`, CPython attempted to promote the `str` by decoding it via the default codec (usually ASCII):

```
1 coerced = PyUnicode_FromEncodedObject(str_obj, "ascii", "strict");
```

If the raw string contained non-ASCII bytes (e.g., `\xe9` for Latin-1, or `\xc3\xa9` for UTF-8), this implicit decoding step raised a `UnicodeDecodeError` at runtime. Python 3.0 eliminated this by raising an unconditional `TypeError` on any implicit mixed-type operation:

```
1 # Python 3.0 runtime coercion ban
2 >>> b"data" + "string"
3 Traceback (most recent call last):
4   File "<stdin>", line 1, in <module>
5 TypeError: can't concat bytes to str
```

At the C-API level, binary and text domains are kept completely disjoint. `PyUnicode_Compare` returns a failure state if one operand is a bytes object, and `PyBytes_Concat` does not accept unicode inputs.

2. C-Level Representations of Binary Objects: `PyBytesObject` & `PyByteArrayObject`

CPython represents the immutable `bytes` type at the C-level as `PyBytesObject`, defined in `Include/bytesobject.h`:

```
1 typedef struct {
2     PyObject_VAR_HEAD
3     Py_hash_t ob_shash;           /* Cached hash value of bytes; -1 if uncomputed
4     /*
5     char ob_sval[1];             /* Contiguous byte vector; dynamically sized to
6     [ob_size + 1] */
7 } PyBytesObject;
```

Here, `ob_sval` is a character array that is allocated inline with the rest of the object structure, terminating with a trailing `\0` to remain compatible with standard C library functions (e.g., `strlen` and `strcpy`).

The mutable sibling, `bytearray`, is represented by `PyByteArrayObject` defined in `Include/bytearrayobject.h`:

```
1 typedef struct {
2     PyObject_VAR_HEAD
3     Py_ssize_t ob_alloc;           /* Number of bytes allocated in the buffer */
4     char *ob_bytes;               /* Pointer to the start of the allocated memory
5     block */
6     char *ob_start;               /* Pointer to the start of the logical byte
7     sequence */
8     Py_ssize_t ob_exports;        /* Reference count of active buffer exports (
9     locking) */
10 } PyByteArrayObject;
```

`ob_bytes` and `ob_start` are split to allow efficient head-deletion/prepends. When a user calls `del array[0]`, `ob_start` is simply incremented and `ob_size` decremented, avoiding an $O(n)$ memory shift. Memory buffer growth is managed via an overallocation growth factor:

$$\text{ob_alloc} = \text{new_size} + (\text{new_size} \gg 3) + (\text{new_size} < 9 ? 3 : 6)$$

This guarantees that append operations have an amortized $O(1)$ time complexity. `ob_exports` tracks references from active `memoryview` objects; if `ob_exports > 0`, any operation that attempts to resize or reallocate the underlying `ob_bytes` buffer will raise a `BufferError`.

3. Low-Level C Representation: Pre-PEP 393 `PyUnicodeObject`

Prior to the string optimization introduced in Python 3.3 (PEP 393), CPython represented Unicode strings using a uniform array of `Py_UNICODE` units. `PyUnicodeObject` was defined in `Include/unicodeobject.h` as:

```
1 typedef struct {
2     PyObject_HEAD
3     Py_ssize_t length;            /* Number of code points */
4     Py_UNICODE *str;             /* Pointer to the character array */
5     Py_hash_t hash;              /* Cached hash value; -1 if uncomputed */
6     PyObject *defenc;            /* Cached UTF-8 encoded bytes representation */
7 } PyUnicodeObject;
```

The representation type `Py_UNICODE` was defined at compile-time as: *** UCS-2 (Narrow Build)**: Compiled with `wchar_t` as a 16-bit type (size 2 bytes). Characters beyond `U+FFFF` (e.g., emojis) had to be represented as surrogate pairs, causing index and len calculations to mismatch. *** UCS-4 (Wide Build)**: Compiled with `wchar_t` as a 32-bit type (size 4 bytes). Every Unicode code point mapped to exactly one array index, but ASCII-only strings consumed 4 times the memory they required.

4. PEP 383: The surrogateescape Error Handler Mechanics

To allow POSIX filesystems (which use arbitrary null-terminated byte sequences for paths) to round-trip pathnames containing invalid UTF-8 bytes without throwing exceptions, PEP 383 introduced the `surrogateescape` error handler. During decoding, any invalid byte (which does not form a valid UTF-8 sequence) is mapped to a high-surrogate code point in the range `U+DC80` to `U+DCFF` via:

$$\text{code_point} = 0xDC00 + \text{byte_value}$$

During encoding, these specific surrogate code points are mapped back to their original raw bytes:

$$\text{byte_value} = \text{code_point} - 0xDC00$$

This allows arbitrary binary data to be round-tripped through Unicode `str` representations:

```
1 # Raw undecodable path round-trip simulation
2 raw_bytes = b"bad_\xff_path.txt"
3 decoded_str = raw_bytes.decode("utf-8", "surrogateescape")
4 # decoded_str contains code point U+DCFF where \xff was.
5 reencoded_bytes = decoded_str.encode("utf-8", "surrogateescape")
6 assert reencoded_bytes == raw_bytes
```

The CPython implementation in `Objects/unicodeobject.c` handles this matching logic:

```
1 /* Pseudocode of surrogateescape decoding branch */
2 if (status == INVALID_BYTE) {
3     *unicode_ptr++ = 0xDC00 + (unsigned char)input_byte;
4 }
```

9.0.2 7.2 The `print()` Function Redesign

1. Bytecode Comparison

In Python 2.x, `print` was a core language statement with specialized bytecodes. In Python 3.0, it was unified into a standard built-in function.

Python 2.7 compiler translation for `print "hello":`

```
1 0 LOAD_CONST          0 ('hello')
2 3 PRINT_ITEM
3 4 PRINT_NEWLINE
```

The VM execution loop (`ceval.c`) maps `PRINT_ITEM` directly to standard output stream writing operations, meaning the behavior was fixed at compilation.

Python 3.0 compiler translation for `print("hello"):`

```
1 0 LOAD_GLOBAL          0 (print)
2 3 LOAD_CONST           1 ('hello')
3 6 CALL_FUNCTION        1
```

The interpreter dynamically resolves `print` at runtime via standard namespace lookups. This enables runtime overriding:

```
1 import builtins
2 def custom_print(*args, **kwargs):
3     builtins.print("[LOG]", *args, **kwargs)
4 builtins.print = custom_print
```

2. C-Level stream writing & `TextIOWrapper`

The C-level entrypoint for the `print()` function is `builtin_print` inside `Python/builtinmodule.c`:

```
1 static PyObject *
2 builtin_print(PyObject *self, PyObject *args, PyObject *kwargs)
3 {
4     static char *kwlist[] = {"sep", "end", "file", "flush", 0};
5     PyObject *sep = Py_None, *end = Py_None, *file = Py_None;
6     int flush = 0;
```

```

7 // Parsing keywords using PyArg_ParseTupleAndKeywords...
8 ...
9 }

```

- **Stream Resolution:** If `file` is `Py_None` or omitted, the function queries `sys.stdout` via `PySys_GetObject("stdout")`.
- **String Conversions:** For each positional argument, `builtin_print` calls `PyObject_Str(arg)` to compute its string representation.
- **Buffer Flushing:** After calling `file.write()`, if `flush=True` is set, CPython attempts to invoke the `flush()` method of the file object. If `sys.stdout` wraps a standard output descriptor (`stdout`), this triggers `fflush(stdout)` in the underlying C library, executing a synchronous write syscall.

9.0.3 7.3 Division Unification (PEP 238)

1. The Division Operators and Bytecodes

Python 2.x's division operator `/` mapped to the `BINARY_DIVIDE` bytecode, which executed floor division if both operands were integers, and true division if either was a float. This led to silent precision loss bugs. Python 3.0 unified this by assigning `/` to true division and introducing `//` for floor division: * **True Division** (`/`) maps to the `BINARY_TRUE_DIVIDE` bytecode, which always delegates to the `nb_true_divide` slot. * **Floor Division** (`//`) maps to the `BINARY_FLOOR_DIVIDE` bytecode, delegating to the `nb_floor_divide` slot.

2. C-Level Integer Division Mechanics

CPython defines number methods in the `PyNumberMethods` struct inside `Include/object.h`:

```

1 typedef struct {
2     binaryfunc nb_add;
3     binaryfunc nb_subtract;
4     ...
5     binaryfunc nb_floor_divide;
6     binaryfunc nb_true_divide;
7 } PyNumberMethods;

```

For integer types, these slots point to C functions in `Objects/longobject.c`: * `nb_floor_divide` points to `long_div`. It computes the integer quotient and rounds towards negative infinity. * `nb_true_divide` points to `long_true_divide`.

The execution path of `long_true_divide(x, y)`: 1. It extracts the size of integers `x` and `y`. If they fit within C double precision (53 bits of mantissa), it converts them to double:

$$x_{\text{double}} = (\text{double})x; \quad y_{\text{double}} = (\text{double})y;$$

2. If the integer values exceed double limits, CPython runs an arbitrary-precision float conversion algorithm (`_PyLong_Format` or manual bit shift scaling) to extract the most significant 53 bits. 3. It performs the float division and instantiates a new `PyFloatObject` to hold the output:

$$\text{result} = x_{\text{double}} / y_{\text{double}}$$

4. If `y = 0`, it raises a `ZeroDivisionError` via `PyErr_SetString`.

9.0.4 7.4 Lazy Iterators & Dictionary View Objects

1. Lazy Conversions & rangeobject Memory Layout

Python 3.0 replaced list-producing functions with lazy iterators. The `range()` built-in returns a `rangeobject` defined in `Objects/rangeobject.c`:

```
1 typedef struct {
2     PyObject_HEAD
3     PyObject *start;
4     PyObject *stop;
5     PyObject *step;
6     PyObject *length;
7 } rangeobject;
```

Because the sequence elements are not pre-allocated, a `rangeobject` consumes $O(1)$ memory. To compute the value at a specific index i , the type's sequence method (`range_item`) performs a direct arithmetic step calculation:

$$\text{value} = \text{start} + i \times \text{step}$$

This math enables $O(1)$ index access and $O(1)$ containment checks for numeric targets:

$$\text{remainder} = (\text{target} - \text{start}) \pmod{\text{step}}$$

If `remainder == 0` and `start ≤ target < stop` (or reverse for negative steps), the value is in the range.

2. Dictionary View Objects

Methods like `keys()`, `values()`, and `items()` return dynamic views containing a direct reference back to the parent dictionary, instead of copying elements into a new list. The underlying struct is `PyDictViewObject`:

```
1 typedef struct {
2     PyObject_HEAD
3     PyDictObject *dv_dict;      /* Pointer to parent dictionary struct */
4 } PyDictViewObject;
```

Chapter 10

Python 3.1 to 3.2: Standard Library Consolidation and Threading Pools

10.0.1 8.1 Antoine Pitrou's New GIL (Python 3.2)

To resolve multi-core performance bottlenecks, Python 3.2 replaced the legacy ticker-based GIL with an interval-based GIL designed by Antoine Pitrou.

1. The Convoy Effect & GIL Battle under the Ticker-Based GIL

In Python 2.x and pre-3.2, thread switching was based on a bytecode execution ticker (`sys.checkinterval`, default 100 instructions). Once a thread T_1 executed 100 bytecodes, it released the GIL, signaled waiting threads, and immediately re-attempted to acquire it: 1. T_1 releases GIL $\rightarrow T_1$ immediately calls `acquire` again. 2. Because T_1 is already running on a CPU core with warm caches, it is highly likely to re-acquire the GIL before a sleeping thread T_2 on another core can wake up. 3. This resulted in the **convoy effect** (or GIL battle), causing thread starvation and wasting CPU cycles on rapid, unsuccessful mutex context-switching.

2. The Interval-Based GIL Architecture

The new GIL manages thread switching using a time interval (default 5000 microseconds / 5ms). The state of the GIL is maintained in a global runtime structure (`ceval_gil.h`):

```
1 typedef struct {
2     Py_Mutex_T mutex;           /* Mutex guarding the GIL status fields */
3     Py_COND_T cond;            /* Condition variable for waiting threads */
4     int locked;                 /* Flag indicating if the GIL is locked (1
5     or 0) */
6     unsigned long switch_interval; /* Maximum thread execution duration before
7     switch (in microseconds) */
8
9     PyThreadState* volatile tstate; /* Pointer to the thread state currently
10    holding the GIL */
11    int volatile gil_drop_request; /* Shared atomic flag indicating a thread
12    must drop the GIL */
13 } _gil_runtime_state;
```

3. Step-by-Step Thread State Transitions & Mutex Flow

When thread T_1 holds the GIL and thread T_2 requests execution:

Thread 1 (Active)	Thread 2 (Waiting)
-----	-----
Holds GIL (locked=1)	

4	Executes bytecodes...	Requests GIL via <code>take_gil()</code>
5	with 5ms timeout	Acquires mutex, waits on cond
6	released GIL)	Timeout expires! (T_1 hasn't
7		Sets <code>gil_drop_request = 1</code>
8	Checks <code>gil_drop_request</code> in eval loop	
9	Detects flag set to 1	
10	Releases GIL via <code>drop_gil()</code>	
11	Resets <code>locked=0</code> , <code>tstate=NULL</code>	
12	Signals cond, releases mutex	
13		Wakes up from cond, acquires
14	mutex	Sets <code>locked=1</code> , <code>tstate=T2</code>
15		Resets <code>gil_drop_request=0</code>
16		Starts executing bytecodes...

1. **Requesting the GIL:** Thread T_2 enters `take_gil()`, acquires the mutex, and waits on the condition variable `cond` with a timeout of `switch_interval` (5ms).
2. **Timeout & Requesting Release:** If the timeout expires and T_1 has not released the GIL (e.g., due to executing an I/O operation or blocking system call), T_2 sets the global atomic flag `gil_drop_request = 1`.
3. **Interpreter Eval Loop Check:** In CPython's evaluation loop (`_PyEval_EvalFrameDefault` in `Python/ceval.c`), the interpreter checks `gil_drop_request` between bytecode execution steps:

```

1 /* Check GIL drop request flag */
2 if (_Py_atomic_load_relaxed(&gil_runtime_state.gil_drop_request)) {
3     /* Give up the GIL and wait for reschedule */
4     PyThreadState *tstate = _PyThreadState_GET();
5     PyEval_SaveThread();    /* Drops GIL, signals cond */
6     PyEval_RestoreThread(); /* Re-acquires GIL, blocks if held */
7 }

```

4. **Acquiring the GIL:** Once T_1 calls `PyEval_SaveThread()`, it sets `locked = 0`, signals `cond` to wake up T_2 , and releases the mutex. T_2 wakes up, sets `locked = 1`, clears `gil_drop_request = 0`, and starts its execution phase.

10.0.2 8.2 Thread Pools and Process Pools (`concurrent.futures` / PEP 3148)

PEP 3148 unified concurrent task execution under a shared API using `ThreadPoolExecutor` and `ProcessPoolExecutor`.

1. ThreadPoolExecutor Mechanics

`ThreadPoolExecutor` manages a pool of worker threads using a task queue (`queue.SimpleQueue`): *

Task Submission: When `submit(fn, *args)` is called, CPython wraps the callable in a `Future` object and a `_WorkItem` struct, pushing it to the executor's task queue. *

Worker Execution: Worker threads execute a loop inside `_worker()`: `python def _worker(executor_reference, work_queue): while True: work_item = work_queue.get(block=True) if work_item is not None: work_item`

`.run()` * **GIL Constraints:** Because all worker threads reside in the same memory space, they share the GIL. They can execute I/O-bound operations (e.g., waiting on socket descriptors or file operations) concurrently by releasing the GIL at the C-API level during blocking calls, but cannot execute CPU-bound tasks in parallel.

2. ProcessPoolExecutor & IPC Serialization Bottlenecks

To bypass the GIL for CPU-bound operations, `ProcessPoolExecutor` spawns independent worker processes. Because processes do not share memory, CPython must serialize task arguments and deserialize results using the `pickle` protocol.

The math of Process Pool Execution Overhead:

$$\text{Total Execution Time} = t_{\text{serialization}} + t_{\text{IPC transfer}} + t_{\text{computation}} + t_{\text{IPC return}} + t_{\text{deserialization}}$$

$$\text{IPC Overhead} \propto \text{size of arguments} + \text{size of results}$$

1. **Serialization:** The parent process pickles the function and arguments into a byte stream.
2. **IPC Transfer:** The byte stream is written to an OS pipe or socket descriptor, crossing the user-kernel space boundary twice:

$$\text{Parent User Space} \xrightarrow{\text{write()}} \text{Kernel Buffer} \xrightarrow{\text{read()}} \text{Child User Space}$$

3. **Computation:** The child process runs the task within its own interpreter instance and GIL.
4. **Return:** The child pickles and writes the result back through a return pipe. For small, fast computations, the serialization and pipe I/O overhead can exceed the execution time of the computation itself.

3. Future State Tracking

A `Future` object tracks task completion state. The state transitions are guarded by a thread-local reentrant lock (`threading.RLock`):

```

1 # Future state definitions inside concurrent.futures.futures
2 PENDING = 'PENDING'
3 RUNNING = 'RUNNING'
4 CANCELLED = 'CANCELLED'
5 FINISHED = 'FINISHED'

```

- Calling `cancel()` transitions the state from `PENDING` to `CANCELLED`, waking up any threads waiting on `result()` via a condition variable.
- When the task completes, `set_result(val)` transitions the state to `FINISHED`, stores the return value in `self._result`, and invokes all callback functions registered via `add_done_callback(fn)`.

10.0.3 8.3 Standard Library Consolidation

1. OrderedDict (PEP 372)

Introduced in Python 3.1, `OrderedDict` preserves key insertion order. Since standard dictionaries were unordered at this time, `OrderedDict` maintained order by wrapping a standard dictionary with a private doubly-linked list. The keys are stored as nodes in the linked list:

```

1 # Doubly-linked list node format: [PREV, NEXT, KEY]
2 root = []
3 root[:] = [root, root, None] # Pointer loop representing empty list

```

- **Insertion:** When a new key is added, `OrderedDict` appends it to the end of the doubly-linked list by adjusting pointers:


```

1 last = root[0]
2 last[1] = root[0] = [last, root, key]
3 self.__map[key] = root[0]

```

- **Deletion:** Deleting a key updates the pointers of the neighboring nodes in $O(1)$ time:

```

1 link_prev, link_next, key = self.__map.pop(key)
2 link_prev[1] = link_next
3 link_next[0] = link_prev

```

2. Counter

A dictionary subclass designed for tallying hashable elements. It optimizes standard dictionary lookups by wrapping loops in C:

```

1 # Counter update logic
2 for elem in iterable:
3     self[elem] = self.get(elem, 0) + 1

```

At the C level, it accesses dictionary lookups via `PyDict_GetItemWithError` to bypass Python-level attribute lookup overhead, speeding up frequency counting. ##### 3. `argparse` Introduced in Python 3.2 to replace `optparse`, `argparse` uses a parser state machine to process command-line arguments. It builds a hierarchical action registry (`_actions` list) containing `Action` objects (e.g., `_StoreAction`, `_StoreConstAction`). During `parse_args()`, it iterates over argument strings, matches positional arguments and optional flags (using a regular expression mapping state machine), and applies type coercion hooks (e.g., `int`, `float`) before writing variables to a `Namespace` object.

Part III

Volume III: Generators, Iterators, and Async Inception

Chapter 11

Python 3.3: Yield From Generators and Implicit Namespace Packages

11.0.1 9.1 PEP 380: Generator Delegation via `yield from`

`yield from <iterable>` delegates generator operations directly to a sub-generator or iterable, acting as a transparent channel between the caller and the active sub-generator.

1. Bytecode Compilation and Stack Transitions

When the compiler encounters `yield from`, it emits two specialized bytecodes: `* GET_YIELD_FROM_ITER`: Pops the target iterable from the evaluation stack, checks if it is a generator or iterator, and prepares it for delegation. If it is a generator, it is pushed back directly; otherwise, CPython calls `PyObject_GetIter` to obtain an iterator. `* YIELD_FROM`: Establishes the active delegation channel in the main loop. The VM pops the sub-generator, pulls its yielded value, and pushes it back as the output of the active generator, suspending the frame.

The stack transitions during delegation:

```
1 Active Frame Stack (During yield from)
2 +-----+
3 |      Return Value      | <-- (After StopIteration catches)
4 +-----+
5 |  Sub-generator Ref     | <-- Managed dynamically by YIELD_FROM
6 +-----+
7 |      Receiver Value    | <-- Received from caller's send()
8 +-----+
```

2. C-Level Generator Mechanics: `PyGenObject`

CPython represents generators at the C-level as `PyGenObject` (`Include/genobject.h`):

```
1 typedef struct {
2     PyObject_HEAD
3     struct _frame *gi_frame;           /* Execution frame holding variables and
4     instruction pointer */
5     char gi_running;                   /* Active execution flag (1 if running, 0
6     if suspended) */
7     PyObject *gi_code;                 /* Code object associated with the
8     generator */
9     PyObject *gi_weakreflist;          /* List of weak references */
10    PyObject *gi_name;                  /* Generator name string */
11    PyObject *gi_qualname;              /* Qualified name string */
12    _PyErr_StackItem gi_exc_state;      /* Exception state saved across
13    suspensions */
14 } PyGenObject;
```

- **Suspension Phase:** When a generator yields, the VM sets `gi_running = 0`, saves the current instruction pointer (`f_lasti`) and stack depth within `gi_frame`, and returns control to the caller.
- **Resuming Phase:** When `send()` or `next()` is invoked, CPython sets `gi_running = 1`, restores the thread state's active frame to `gi_frame`, and resumes execution at the saved instruction pointer.

3. Exception & Value Routing Protocol

The CPython implementation inside `Python/ceval.c` maps `yield from` delegation rules: * `send()`

Routing: When the caller sends a value via `generator.send(val)`, the `YIELD_FROM` instruction catches the value, bypasses the delegator's frame, and passes it directly to the sub-generator: `cretval = _PyGen_Send(sub_gen, sent_val, &val); * throw()` **Propagation:** If the caller raises an exception using `generator.throw(type, val, tb)`, the VM passes the exception to the sub-generator's `throw()` method. If the sub-generator catches the exception and yields a new value, execution continues. If the sub-generator raises `StopIteration` or propagates a different exception, the delegating generator is resumed or unwound. * `close()` **Cleanup:** When `close()` is called on the delegating generator, the VM invokes the `close()` method of the sub-generator. If the sub-generator does not terminate, it raises a `RuntimeError`. * **Return Value Unpacking:** When the sub-generator completes by raising `StopIteration`, CPython extracts the `value` attribute from the exception object: `python # CPython StopIteration Unpacking Logic except StopIteration as e: result = e.value # Maps to return value of sub-generator` The value is pushed onto the stack, replacing the sub-generator reference, and execution of the delegating generator resumes.

11.0.2 9.2 PEP 420: Implicit Namespace Packages

PEP 420 introduced implicit namespace packages, allowing packages to span multiple directories on disk without an `__init__.py` file.

1. Import Search Algorithm in `importlib`

During `import foo`, the CPython import engine (`sys.meta_path` hooks) searches directory paths:

1. **Finders Traversal:** The import engine iterates over finders registered in `sys.meta_path` (primarily `PathFinder` which uses paths from `sys.path`).
2. **Standard Package Check:** For each search path in `sys.path`, `PathFinder` checks for a subdirectory `foo` containing `__init__.py`. If found, it returns a standard module spec.
3. **Namespace Path Accumulation:** If no standard package is found, but one or more subdirectories named `foo` exist across `sys.path`, `PathFinder` does not terminate with an error. Instead, it scans all search paths and accumulates all matching directory paths into a list.
4. **Spec Initialization:** It returns a `ModuleSpec` with: * `loader` set to `_NamespaceLoader`. * `submodule_search_locations` containing the accumulated directory list.
5. **Caching:** It registers the paths in `sys.path_importer_cache` to speed up subsequent submodule lookups.

2. Namespace Modules

The loader `_NamespaceLoader` instantiates a module whose `__file__` attribute is `None`, and whose `__path__` contains the list of accumulated directories:

```
1 # Namespace module path list
2 import company.core
3 print(company.core.__path__)
4 # Output: _NamespacePath(['/path1/company/core', '/path2/company/core'])
```

This allows separate wheels or libraries to distribute modules into the same namespace package dynamically.

11.0.3 9.3 PEP 393: Flexible String Representation

PEP 393 redesigned the internal representation of Unicode strings (`str`) to reduce memory usage. CPython now dynamically selects the narrowest character array encoding based on the maximum code point in the string.

1. C-Level Layout Headers

CPython defines three structs in `Include/unicodeobject.h` to represent strings: `* PyASCIIObject` (ASCII only, characters ≤ 127): `c typedef struct { PyObject_HEAD Py_ssize_t length; /* Number of code points */ Py_hash_t hash; /* Cached hash value; -1 if uncomputed */ struct { unsigned int interned:2; /* Interned state (e.g. SGI_INTERNED) */ unsigned int kind:3; /* character size kind (1, 2, or 4 bytes) */ unsigned int compact:1; /* compact layout flag */ unsigned int ascii:1; /* ASCII flag (1 if ASCII only) */ unsigned int ready:1; /* Ready state */ } state; wchar_t *wstr; /* Legacy wchar_t representation cache */ } PyASCIIObject;` The raw characters are stored in memory immediately following this header. `* PyCompactUnicodeObject` (Non-ASCII compact strings, characters  $\leq 65535$ ): `c typedef struct { PyASCIIObject _base; Py_ssize_t utf8_length; /* Length of UTF-8 representation */ char *utf8; /* Pointer to UTF-8 representation cache */ Py_ssize_t wstr_length; /* Length of wchar_t representation */ } PyCompactUnicodeObject;` `* PyUnicodeObject` (Non-compact legacy strings): Used primarily for backward compatibility with the C-API. It adds a pointer to the character data memory address (`data.any`).

2. String Kind Allocation and Promotion Rules

The character array representation is determined by the `kind` field: `* PyUnicode_1BYTE_KIND` (Latin-1, characters ≤ 255): 1 byte per character. `* PyUnicode_2BYTE_KIND` (UCS-2, characters ≤ 65535): 2 bytes per character. `* PyUnicode_4BYTE_KIND` (UCS-4, characters > 65535): 4 bytes per character.

If a string operation (e.g., concatenation or substitution) appends a character that exceeds the current string's maximum code point, CPython allocates a new string with the promoted `kind` and converts the existing characters:

```
1 # String Promotion Simulation
2 s = "abc"           # Kind: 1-byte (ASCII)
3 s += " "            # Promoted to Kind: 1-byte (Latin-1)
4 s += " "            # Promoted to Kind: 2-byte (UCS-2)
5 s += " "            # Promoted to Kind: 4-byte (UCS-4)
```

During promotion from 1-byte to 2-byte, CPython executes a C conversion loop:

```
1 /* C-level character expansion loop */
2 for (Py_ssize_t i = 0; i < length; i++) {
3     dest_2byte[i] = (Py_UCS2)source_1byte[i];
4 }
```

This design maintains $O(1)$ indexing for all strings while optimizing memory usage for ASCII and Latin-1 strings.

11.0.4 9.4 memoryview and the Buffer Protocol (Py_buffer)

The buffer protocol allows Python objects (e.g., bytes, bytearray, array.array) to expose their raw memory buffer directly to other objects without copying data.

1. The Py_buffer Struct

The interface is defined by the Py_buffer struct in Include/object.h:

```

1 typedef struct {
2     void *buf; /* Pointer to the start of the memory block */
3     PyObject *obj; /* Reference to the parent object providing the
4         buffer */
5     Py_ssize_t len; /* Total length of the buffer in bytes */
6     Py_ssize_t itemsize; /* Size of a single element in bytes */
7     int readonly; /* Read-only flag (1 if read-only, 0 if
8         writable) */
9     const char *format; /* Format string describing element type (
10         struct syntax) */
11     int ndim; /* Number of dimensions */
12     Py_ssize_t *shape; /* Array of sizes for each dimension */
13     Py_ssize_t *strides; /* Array of step strides in bytes for each
14         dimension */
15     Py_ssize_t *suboffsets; /* Suboffsets for nested arrays */
16     void *internal; /* Private storage for the buffer provider */
17 } Py_buffer;
    
```

2. Multidimensional Strides Mathematics

The offset in bytes of an element in a multi-dimensional buffer is computed using strides:

$$\text{Offset} = \text{start_offset} + \sum_{i=0}^{n-1} \text{index}_i \times \text{strides}_i$$

For a 2D matrix of shape [2, 3] containing 32-bit C-integers (itemsize = 4) stored in row-major order: * Shape array: [2, 3] * Strides array: [12, 4] (since each row is 3 × 4 = 12 bytes, and each column is 4 bytes). To access index [1, 2]:

$$\text{Offset} = (1 \times 12) + (2 \times 4) = 12 + 8 = 20 \text{ bytes}$$

CPython can transpose or slice buffers in $O(1)$ time by altering the shape and strides arrays without moving or copying the underlying data in memory.

3. Buffer Locking & Safety

To prevent memory corruption, objects must coordinate resizing with active buffers: 1. **Buffer Export:** When memoryview is created on a bytearray, it invokes the provider's bf_getbuffer slot, which populates the Py_buffer struct and increments the provider's export counter (ob_exports++).

2. **Resizing Ban:** While ob_exports > 0, any operation that attempts to resize or reallocate the memory buffer (e.g., bytearray.append() or bytearray.extend()) is blocked and raises a BufferError: python # Buffer protection verification data = bytearray(b"raw_bytes") view = memoryview(data) data.extend(b"_new") # Raises BufferError: Existing exports prevent resizing

3. **Buffer Release:** When the memoryview is garbage collected or explicitly closed, it calls PyBuffer_Release(), which invokes the provider's bf_releasebuffer slot to decrement the export counter (ob_exports--). Once ob_exports reaches 0, the buffer can be resized or freed.

Chapter 12

Python 3.4: Asyncio Inception, Pathlib, and Enum Architectures

12.0.1 10.1 Asyncio Inception: Generators and Event Loops (PEP 3156)

1. Generator-Based Coroutine Syntax

Prior to Python 3.5's native `async/await` syntax, asyncio coroutines were defined using generators decorated with `@asyncio.coroutine` and delegated execution using `yield from`:

```
1 import asyncio
2
3 @asyncio.coroutine
4 def fetch_data():
5     yield from asyncio.sleep(1)  /* Delegate to asyncio.sleep generator */
6     return "payload_data"
```

2. The Event Loop and Selector Multiplexing

Under the hood, `asyncio` does not use multi-threading to achieve concurrency. Instead, it runs an **Event Loop** on a single thread. The event loop manages a collection of tasks and schedules them using operating system-level I/O multiplexing. * **OS-Level Selector**: The event loop wraps Python's built-in `selectors` module, which maps to POSIX system calls: `select()`, `poll()`, Linux `epoll()`, or macOS `kqueue()`. * **Multiplexing Cycle**: 1. The event loop registers sockets or file descriptors with the OS selector, specifying the events to monitor (e.g., readability or writability) and registering associated callbacks. 2. The loop enters a blocking poll state, calling `selector.select()`. The OS blocks the thread until one or more registered file descriptors become ready. 3. When an event triggers, the OS returns the list of ready descriptors. The event loop iterates over this list and schedules the registered callbacks for immediate execution.

3. Future and Task Execution Cycle

The event loop schedules and executes coroutines using two core structures: * **Future**: Represents the eventual result of an asynchronous operation. It maintains an internal state (PENDING, FINISHED, CANCELLED) and stores a list of callbacks to invoke upon completion. * **Task**: A subclass of `Future` that wraps a coroutine. It acts as the bridge between the event loop and the coroutine's execution.

The Step-by-Step Task Loop:

1. **Task Initialization**: The event loop schedules the `Task`. It invokes the task's `_step` method, which calls `coroutine.send(None)` to start the coroutine.

2. **Suspension:** The coroutine runs until it encounters a blocking operation: `yield from future`. It yields control back to the task, returning the pending `Future` object.
3. **Callback Registration:** The `Task` receives this pending future. It calls `future.add_done_callback(task._step)` to register its own step method as a callback on the future, then yields control back to the event loop.
4. **Polling:** The event loop performs other work.
5. **Resuming:** When the I/O event completes, the selector changes the future's state to `FINISHED`. The future pops and schedules all registered callbacks. The task's `_step` method executes, calling `coroutine.send(result)` to resume execution inside the coroutine.

12.0.2 10.2 pathlib: Object-Oriented Filesystem Paths (PEP 428)

Before Python 3.4, filesystem paths were treated as raw strings, requiring developers to write platform-specific path manipulation code using the `os.path` module (e.g., `os.path.join()`, `os.path.split()`). PEP 428 introduced `pathlib` to represent paths as structured objects.

1. Class Hierarchy

`pathlib` splits path representations into a clear class hierarchy:

```

1          [ PurePath ] (Abstract Base / Path Math Only)
2          /         \
3      [ PurePosixPath ]   [ PureWindowsPath ]
4          |               |
5      [ PosixPath ]       [ WindowsPath ]
6          \               /
7          \-[ Path ]-/ (Instantiates PosixPath or WindowsPath
                        dynamically)

```

- **PurePath:** Provides pure path manipulations (string parsing, joining, metadata extraction) without executing OS-level system calls. It can be instantiated on any system (e.g., you can parse Windows paths on a Linux host using `PureWindowsPath`).
- **Path:** Inherits from `PurePath` and adds OS-level filesystem queries (like `.exists()`, `.glob()`, `.mkdir()`, and `.read_text()`).
- **Dynamic Instantiation:** When `Path()` is instantiated at runtime, CPython detects the host operating system and dynamically returns an instance of either `PosixPath` or `WindowsPath`.

2. Path Concatenation Operator Overloading

`pathlib` overloads the division operator (`/`) using the special method `__truediv__` to implement intuitive path joining:

```

1 from pathlib import Path
2 base_path = Path("/var")
3 log_path = base_path / "log" / "nginx.log"

```

Under the hood: 1. When `base_path / "log"` is evaluated, CPython calls `base_path.__truediv__("log")`. 2. The `__truediv__` method parses the argument, checks compatibility, joins the string segments using the host operating system's path separator (`/` or `\`), and returns a new `Path` object containing the joined path string, avoiding redundant string allocations.

12.0.3 10.3 Enum: Metaclass and Singletons (PEP 435)

PEP 435 introduced `Enum` to the standard library to support structured enumerations.

1. Metaclass Construction: `EnumMeta`

Enums are created using the custom metaclass `EnumMeta`. When the compiler processes an `Enum` class definition:

1. `EnumMeta.__prepare__()` returns a custom dictionary structure (`_EnumDict`) that tracks insertion order and raises errors if a duplicate member name is defined.
2. `EnumMeta.__new__()` parses the class namespace, separating attributes into `Enum` members and standard methods.
3. For each `Enum` member (e.g., `RED = 1`), `EnumMeta` instantiates a singleton instance of the `Enum` class, assigning the name (`RED`) and value (`1`) to the instance.
4. The metaclass updates the class dictionary, replacing the raw integer literals with the instantiated singleton object references.

2. Immutability and Protections

To ensure Enums act as stable constant mappings:

- * **Member Immutability:** `Enum` overrides `__setattr__` and `__delattr__` to block any attempt to modify or delete a member's value at runtime:

```
python class Color(Enum): RED = 1
    # This raises an AttributeError Color.RED.value = 2
    * **Iterability**: EnumMeta implements
    iter, allowing developers to iterate over enum members in declaration order.
    * **Lookup**: It implements
    getitem(lookup by name, e.g., Color['RED']) and call(lookup by value, e.g., Color(1)).
```

Chapter 13

Python 3.5: Native Async/Await Coroutines and Matrix Operations

13.0.1 11.1 PEP 492: Native Coroutines (async/await Internals)

Python 3.5 introduced native coroutines via PEP 492, establishing a clear separation between standard generators and asynchronous tasks to prevent logical developer errors.

1. The Generator-Coroutine Limitation

In Python 3.4, coroutines were decorated generators (`@asyncio.coroutine`). Under the hood, these coroutines were instances of `PyGenObject`. Because they shared the same C type as standard generators, they exposed the standard iterator protocol. Developers could mistakenly iterate over a coroutine using a `for` loop or call `next()` on it directly, causing unexpected runtime behavior. Additionally, this lack of syntactic separation meant that: * The interpreter could not perform static analysis to detect missing `await/yield from` calls on coroutine objects. * It was possible to call `next(coro())` directly, stepping the generator and bypassing the event loop's state machine. * The `yield from` syntax was overloaded, being used for both yielding from standard generators and awaiting asynchronous tasks, creating a high degree of cognitive load and making code refactoring prone to errors.

2. CPython Representation: PyCoroObject

Native coroutines declared using `async def` compile directly to a specialized C structure, `PyCoroObject` (Include/genobject.h):

```
1 typedef struct {
2     PyObject_HEAD
3     struct _frame *cr_frame;           /* Frame object executing coroutine code */
4     PyObject *cr_code;                 /* Code object compiled from async def */
5     PyObject *cr_name;                 /* Coroutine name */
6     PyObject *cr_qualname;             /* Qualified name */
7     PyObject *cr_origin;               /* Awaited task or future origin trace */
8     PyObject *cr_weakreflist;          /* Weak reference list pointer */
9     char cr_running;                   /* Flag indicating if coroutine is active */
10 } PyCoroObject;
```

Let's break down the role of each field in the CPython runtime: * `PyObject_HEAD`: The standard object header containing the reference counter (`ob_refcnt`) and the pointer to the type object (`ob_type`). In this case, `ob_type` points to the `PyCoro_Type` struct. * `cr_frame`: A pointer to the execution frame (`PyFrameObject`). It holds the local execution state, the bytecode pointer, the local evaluation stack, and local variable cells. * `cr_code`: Points to the `PyCodeObject` representing

the compiled bytecode. * `cr_name`: The coroutine's name (represented by a `PyUnicodeObject`). * `cr_qualname`: The qualified name (e.g. `ClassName.method_name`). * `cr_origin`: If debugging is enabled via `sys.set_coroutine_origin_tracking_depth()`, this pointer stores a traceback tuple showing where the coroutine was instantiated, enabling deep debugging of orphaned or un-awaited coroutines. * `cr_weakreflist`: Head of a doubly-linked list of weak reference objects pointing to this coroutine. * `cr_running`: A single-byte boolean flag. If `cr_running` is true, it indicates the coroutine frame is currently executing. Any attempt to resume the coroutine from another context will raise a `RuntimeError: coroutine already running`, protecting the runtime from re-entrant frame execution.

Because `PyCoroObject` is a distinct C type, it does not implement sequence or iterator slots (`tp_iter` is `NULL`), preventing standard iteration errors.

3. The `am_await` Slot Protocol

Instead of iteration slots, `PyCoroObject` implements the `tp_as_async` protocol table, specifically the `am_await` slot:

```
1 typedef struct {
2     unaryfunc am_await;           /* __await__ implementation pointer */
3     unaryfunc am_aiter;          /* __aiter__ implementation pointer */
4     unaryfunc am_anext;          /* __anext__ implementation pointer */
5 } PyAsyncMethods;
```

When an object is awaited via `await expr`, the CPython interpreter executes the following internal evaluation loop logic: 1. It checks if the object's type has `tp_as_async` populated and if `tp_as_async->am_await` is non-`NULL`. 2. If yes, it calls `am_await(expr)`. This slot function **MUST** return an iterator (an object that implements `tp_iternext`). 3. For native coroutines (`PyCoroObject`), their type has a custom `am_await` slot that wraps the coroutine in a `PyCoroWrapper` (or returns a wrapper that exposes standard iterator operations for the event loop to drive). 4. For user-defined classes, a custom `__await__` method is defined. At the C level, this maps to `am_await` resolving the python-level `__await__` method. It must return an iterator.

```
1 # User-level custom awaitable implementation
2 class DatabaseConnection:
3     def __await__(self):
4         # We yield control back to the event loop if the socket is not ready
5         while not self.is_connected():
6             yield None
7         return self._connection
```

4. Bytecode Compilation and Tracing: `GET_AWAITABLE`

When the compiler parses an `await` expression statement, it emits a specialized `GET_AWAITABLE` bytecode instruction: 1. `GET_AWAITABLE` pops the expression result off the evaluation stack. 2. The VM checks if the object implements the `am_await` slot (either natively or via class overrides). 3. If valid, the VM calls the slot to retrieve the awaitable iterator and pushes it onto the stack. If not valid, it raises a `TypeError`. 4. This is followed by a delegation loop (similar to `yield from`) that yields control back to the event loop if the awaitable is pending.

Let's trace the compiled bytecodes of an asynchronous execution path. Consider the following code:

```
1 async def get_val():
2     return 42
3
4 async def main():
5     val = await get_val()
```

The CPython 3.5 compiler generates the following disassemblies:

Disassembly of `get_val`:

```

1  2      0 LOAD_CONST                    1 (42)
2      3 RETURN_VALUE

```

Disassembly of `main`:

```

1  5      0 LOAD_GLOBAL                      0 (get_val)
2      3 CALL_FUNCTION                      0
3      6 GET_AWAITABLE                     0
4      9 LOAD_CONST                      0 (None)
5     12 YIELD_FROM                     0
6     15 STORE_FAST                      0 (val)
7     18 LOAD_CONST                      0 (None)
8     21 RETURN_VALUE

```

Step-by-Step VM Execution Trace of `main()`:

1. `LOAD_GLOBAL 0`: Resolves the identifier `get_val` from the global namespace dictionary and pushes the function object onto the evaluation stack.
2. `CALL_FUNCTION 0`: Invokes `get_val()`. Since it was defined with `async def`, the VM immediately creates a new `PyCoroObject` and pushes it onto the evaluation stack. Note that the body of `get_val` does not execute yet.
3. `GET_AWAITABLE`:
 - Pops the `PyCoroObject` off the stack.
 - Accesses `ob_type->tp_as_async->am_await`.
 - Executes the slot, which wraps the native coroutine or verifies it is ready to be driven.
 - Pushes the resulting awaitable iterator back onto the stack.
4. `LOAD_CONST 0`: Pushes `None` onto the stack as the priming value.
5. `YIELD_FROM`:
 - Performs a loop that mimics `yield from` behavior.
 - It pops the value (`None`) and sends it to the awaitable iterator by invoking its `tp_iternext` (or equivalent `send` slot).
 - The sub-frame `get_val` runs and reaches `RETURN_VALUE`. The runtime raises a `StopIteration` containing the return value 42.
 - `YIELD_FROM` catches `StopIteration`, extracts 42 from the exception's `value` attribute, and pushes it onto the evaluation stack.
6. `STORE_FAST 0`: Pops the value 42 off the stack and stores it in local variable `val`.

```

1 Evaluation Stack Transitions during 'await':
2
3 Step 1: [ get_val (fn) ] <-- LOAD_GLOBAL
4 Step 2: [ PyCoroObject ] <-- CALL_FUNCTION
5 Step 3: [ CoroWrapper ] <-- GET_AWAITABLE
6 Step 4: [ CoroWrapper ] <-- LOAD_CONST (None)
7         [ None ]
8 Step 5: [ 42 ] <-- YIELD_FROM (catches StopIteration(42))
9 Step 6: [ ] <-- STORE_FAST (val = 42)

```

13.0.2 11.2 Matrix Multiplication Operator (`@` / PEP 465)

Python 3.5 introduced the binary operator `@` (and its in-place version `@=`) to support clean mathematical syntax for matrix multiplication.

1. Low-Level C Slots Mapping

CPython maps the matrix multiplication operator to two new function pointer slots inside the `PyNumberMethods` struct (`Include/object.h`):

```
1 typedef struct {
2     /* ... */
3     binaryfunc nb_matrix_multiply;           /* Corresponds to __matmul__ */
4     binaryfunc nb_inplace_matrix_multiply;   /* Corresponds to __imatmul__ */
5 } PyNumberMethods;
```

When evaluating `A @ B`, CPython's binary operation dispatch system invokes `PyNumber_MatrixMultiply` (`A`, `B`):

- Left-to-Right Dispatch:** If the type of `A` defines `tp_as_number->nb_matrix_multiply`, it executes the slot function.
- Right-to-Left Fallback:** If `A`'s slot is `NULL`, or it returns `Py_NotImplemented`, CPython checks if the type of `B` defines `nb_matrix_multiply`.
- Subclass Precedence:** If `B` is a subclass of `A` and overrides `nb_matrix_multiply`, `B`'s slot is checked and called *before* `A`'s slot, allowing specialized subclass multiplication overrides.
- TypeError:** If neither slot returns a valid object (or both return `Py_NotImplemented`), the VM raises a `TypeError`: `unsupported operand type(s) for @`.

2. Python-Level Interface and Numeric Implementations

Developers can customize this operator behavior by implementing the Python magic methods `__matmul__`, `__rmatmul__`, and `__imatmul__`:

```
1 class CustomMatrix:
2     def __init__(self, grid):
3         self.grid = grid
4
5     def __matmul__(self, other):
6         if not isinstance(other, CustomMatrix):
7             return NotImplemented
8         # Compute the dot product matrix
9         rows_A, cols_A = len(self.grid), len(self.grid[0])
10        rows_B, cols_B = len(other.grid), len(other.grid[0])
11        assert cols_A == rows_B, "Dimension mismatch!"
12
13        result = [[0] * cols_B for _ in range(rows_A)]
14        for i in range(rows_A):
15            for j in range(cols_B):
16                result[i][j] = sum(self.grid[i][k] * other.grid[k][j] for k in
17range(cols_A))
18        return CustomMatrix(result)
```

3. High-Performance Implementations

This dedicated operator allowed numerical libraries (like NumPy) to bypass verbose nested function call chains. Consider this linear algebra comparison:

```
1 # Pre-Python 3.5 (Verbose and nesting-heavy)
2 result = A.dot(B).dot(C) + D.dot(E)
3
4 # Python 3.5+ (Clean, mathematical representation)
5 result = (A @ B @ C) + (D @ E)
```

NumPy defines C-level slots mapping for `nb_matrix_multiply` on its `ndarray` types, executing optimized BLAS or LAPACK matrix routines underneath. By pushing the matrix multiplication to these compiled C/Fortran libraries, they release the Global Interpreter Lock (GIL) during compilation of large arrays, permitting true multi-threaded parallel computation across CPU cores.

13.0.3 11.3 Extended Unpacking Generalizations (PEP 448)

PEP 448 expanded the capabilities of the unpacking operators `*` (iterable unpacking) and `**` (dictionary unpacking), allowing multiple unpacks inside collection literals and function calls.

1. Bytecode Compilation and Unpacking Instructions

Prior to Python 3.5, unpacking was limited to a single occurrence. To support multiple unpacking operations, the CPython compiler was updated to emit specialized group unpacking bytecodes: `* BUILD_LIST_UNPACK`: Pops multiple sequences off the stack, converts them to tuples/lists, and concatenates them into a single list. `* BUILD_TUPLE_UNPACK`: Merges multiple popped sequences into a tuple. `* BUILD_SET_UNPACK`: Merges sequences into a set, deduplicating elements. `* BUILD_MAP_UNPACK`: Pops multiple dictionaries/mappings from the stack and merges them into a single dictionary.

Note on Compilation Evolution: In later versions of CPython (e.g., Python 3.9+), these unpacking instructions were replaced by highly specialized, lower-overhead instructions such as `LIST_EXTEND`, `DICT_MERGE`, and `DICT_UPDATE` to avoid the temporary list creation overhead for every unpacked component.

2. Detailed Bytecode Traces

Case A: List Unpacking with Multiple Iterables Consider compiling this list literal containing mixed values and unpacking:

```
1 [1, *[2, 3], 4]
```

The Python 3.5 compiler generates the following bytecode:

1	1	0	LOAD_CONST	0	(1)	
2		3	BUILD_LIST	1		/* Push list [1] onto
	stack */					
3		6	LOAD_CONST	1	(2)	
4		9	LOAD_CONST	2	(3)	
5		12	BUILD_LIST	2		/* Push list [2, 3]
	onto stack */					
6		15	LOAD_CONST	3	(4)	
7		18	BUILD_LIST	1		/* Push list [4] onto
	stack */					
8		21	BUILD_LIST_UNPACK	3		/* Merge the 3 list
	elements on the stack */					

`BUILD_LIST_UNPACK 3` pops the three constructed list elements off the stack, performs dynamic sequence iteration to concatenate their values, and pushes the final unified list `[1, 2, 3, 4]` back onto the evaluation stack.

Case B: Dictionary Unpacking with Multiple Mappings Consider this dictionary literal merging two mappings:

```
1 y = {**d1, **d2, 'key': 42}
```

The Python 3.5 compiler generates the following bytecode:

1	1	0	LOAD_FAST	0	(d1)	/* Push dictionary d1
	*/					
2		3	LOAD_FAST	1	(d2)	/* Push dictionary d2
	*/					
3		6	LOAD_CONST	2	('key')	

4	9	LOAD_CONST	3	(42)	
5	12	BUILD_MAP	1		/* Push dictionary {'
		key': 42} */			
6	15	BUILD_MAP_UNPACK	3		/* Merge the 3 dict
		elements on the stack */			
7	18	STORE_FAST	2	(y)	

During execution: * `BUILD_MAP_UNPACK 3` evaluates the three mappings popped from the stack. * It iterates through their keys, updating the target dictionary. * In Python 3.5, key collisions do not raise an error for dictionary literals; rather, keys to the right overwrite keys to the left (e.g., `d2` overrides `d1`, and `'key'` overrides any prior matching key). * In contrast, if multiple unpacking overlaps occur during keyword argument calls (e.g. `func(**d1, **d2)`), duplicate keys trigger a runtime `TypeError` exception.

Part IV

Volume IV: Expressiveness & Developer Ergonomics

Chapter 14

Python 3.6: F-Strings Formatting, Variable Annotations, and Compact Dicts

14.0.1 12.1 PEP 498: Formatted String Literals (f-strings)

Python 3.6 introduced Formatted String Literals, commonly referred to as **f-strings**, to provide a clean, readable, and highly performant mechanism for string interpolation.

1. Legacy Formatting Overheads

Prior to Python 3.6, developers relied on `%` (printf-style) formatting or the `.format()` string method. Both approaches introduce significant runtime overhead:

- * **Percent (%) Formatting:** Uses ancient C-style printf routines. It is structurally rigid, struggles to handle complex container types gracefully, and requires allocating temporary tuples for positional parameters.
- * **Format Method (`.format()`):** Introduces two major layers of overhead:
 1. **Attribute Resolution:** Evaluating `"{}".format(x)` requires calling the `__getattr__` method on the string object to locate the `.format` method.
 2. **Function Call Overhead:** Invoking `.format()` pushes a new CPython call frame onto the stack, builds temporary positional arguments tuple `*args` and keyword arguments dictionary `**kwargs`, and parses the format string *at runtime* on every single invocation.

2. f-String Compilation Mechanics

f-strings bypass these overheads by converting interpolation directly into optimized bytecode operations at compile time. When the CPython parser encounters an f-string (e.g. `f"Name: {name}, Age: {age}"`), it parses the literal components and the embedded expressions into separate AST nodes:

1. **Static Literal Sections:** Treated as constant string literals.
2. **Embedded Expressions:** Extracted, parsed as separate AST sub-trees, and compiled into standard bytecode instruction sequences.

The compiler generates two primary bytecode instructions to handle f-strings:

- * **FORMAT_VALUE:** Pops a value off the stack, applies formatting parameters (str/repr/ascii conversion or formatting specifiers), and pushes the formatted string back onto the stack.
- * **BUILD_STRING:** Pops a specified number of strings off the stack, performs a high-performance C-level string concatenation (`PyUnicode_Concat` or pre-allocated buffer copy), and pushes the final concatenated string onto the stack.

CPython FORMAT_VALUE Flag Specification: The `FORMAT_VALUE` instruction takes an 8-bit integer argument representing rendering flags:

- * `0x00`: No conversion (calls `__format__` directly).

* 0x01: Call `str()` (corresponding to the `!s` flag). * 0x02: Call `repr()` (corresponding to the `!r` flag). * 0x04: Call `ascii()` (corresponding to the `!a` flag). * 0x08: Have a format specifier. If this bit is set, the instruction pops two values from the stack: the format specifier string (e.g., `".2f"`) followed by the value to format.

3. Bytecode Disassembly Analysis

Let's analyze the difference in compilation between `.format()` and an f-string:

```
1 import dis
2
3 def legacy_format(name, age):
4     return "Name: {}, Age: {}".format(name, age)
5
6 def fstring_format(name, age):
7     return f"Name: {name!r:10}, Age: {age}"
```

Disassembly of `legacy_format`:

1	2	0	LOAD_CONST	1	('Name: {}, Age: {}')
2		3	LOAD_ATTR	0	(format)
3		6	LOAD_FAST	0	(name)
4		9	LOAD_FAST	1	(age)
5		12	CALL_FUNCTION	2	
6		15	RETURN_VALUE		

Note: `LOAD_ATTR` and `CALL_FUNCTION` demonstrate the high overhead of method lookup and frame creation at runtime.

Disassembly of `fstring_format`:

1	5	0	LOAD_CONST	1	('Name: ')
2		3	LOAD_FAST	0	(name)
3		6	FORMAT_VALUE	2	/* repr (!r) conversion
	flag */				
4		9	LOAD_CONST	2	('10')
	string */				/* Format specifier
5		12	FORMAT_VALUE	10	/* Flag 0x02 (repr) + 0
	x08 (has spec) = 10 */				
6		15	LOAD_CONST	3	(' , Age: ')
7		18	LOAD_FAST	1	(age)
8		21	FORMAT_VALUE	0	/* Default formatting
	*/				
9		24	BUILD_STRING	5	/* Concatenate the 5
	stack items */				
10		27	RETURN_VALUE		

Because the VM calculates the lengths of all strings in the stack, `BUILD_STRING` 5 pre-allocates a single contiguous memory block in the C heap and copies the characters directly, completely avoiding intermediate string allocations and Python-level function calls.

14.0.2 12.2 PEP 526: Syntax for Variable Annotations

PEP 526 introduced variable type annotations to complement the function parameter annotations defined in PEP 3107.

1. Class and Module Level Annotations

At the module or class level, type annotations are evaluated at module import or class definition time. 1. The compiler evaluates the annotated type expression. 2. It creates or updates a dictionary named `__annotations__` in the module or class namespace. 3. The type annotation metadata is stored inside this dictionary: `{'variable_name': type_object}`.

Let's examine how the CPython compiler compiles a class with variable annotations:

```
1 class Profile:
2     name: str = "Anonymous"
3     age: int
```

Compiled Bytecode for Profile class namespace creation:

```
1 1          0 LOAD_NAME          0 (__name__)
2          3 STORE_NAME          1 (__module__)
3          6 LOAD_CONST          0 ('Profile')
4          9 STORE_NAME          2 (__qualname__)
5         12 SETUP_ANNOTATIONS          /* Initializes class
   __annotations__ */
6         15 LOAD_NAME          3 (str)
7         18 LOAD_NAME          4 (__annotations__)
8         21 LOAD_CONST          1 ('name')
9         24 STORE_SUBSCR          /* Store name: str in
   __annotations__ */
10        25 LOAD_CONST          2 ('Anonymous')
11        28 STORE_NAME          5 (name)
12        31 LOAD_NAME          6 (int)
13        34 LOAD_NAME          4 (__annotations__)
14        37 LOAD_CONST          3 ('age')
15        40 STORE_SUBSCR          /* Store age: int in
   __annotations__ */
16        41 LOAD_CONST          4 (None)
17        44 RETURN_VALUE
```

- **SETUP_ANNOTATIONS:** Emitted by the compiler to initialize the `__annotations__` dictionary in the active namespace if it does not already exist.
- **STORE_SUBSCR:** Updates `__annotations__` dynamically at runtime, showing that class-level annotations do carry a small runtime initialization overhead.

2. Function-Level Variable Annotations

In contrast to classes and modules, variable annotations defined inside function scopes are **completely ignored** at runtime:

```
1 def process():
2     x: int = 42
```

Disassembly of process:

```
1 2          0 LOAD_CONST          1 (42)
2          3 STORE_FAST          0 (x)
3          6 LOAD_CONST          0 (None)
4          9 RETURN_VALUE
```

Notice that there are no type checks, no `SETUP_ANNOTATIONS`, and no `__annotations__` dictionary overhead. The type annotation metadata `int` is entirely stripped by the compiler. This ensures that function execution paths (which are executed frequently) maintain maximum speed and zero memory allocation overhead.

3. Forward References and Runtime Inspection

Because class/module annotations are executed at definition time, declaring a type that has not yet been defined will raise a `NameError`:

```

1 class Node:
2     parent: Node # Raises NameError: name 'Node' is not defined

```

To bypass this, developers use string literals (forward references):

```

1 class Node:
2     parent: 'Node' # Compiles successfully as a string literal

```

At runtime, frameworks inspect these annotations using `inspect.get_type_hints()` or `typing.get_type_hints()`. These utilities automatically resolve string-based forward references by evaluating them within the global and local namespace of the target class or module.

14.0.3 12.3 CPython Compact Dictionary Design

Python 3.6 replaced CPython's legacy dictionary layout with a compact, ordered representation proposed by Raymond Hettinger.

1. Legacy Dictionary Layout (Pre-3.6)

Historically, a CPython dictionary was implemented as a single, large sparse hash table. Every slot in the table contained a 24-byte `PyDictKeyEntry` structure:

```

1 typedef struct {
2     Py_hash_t me_hash; /* 8 bytes */
3     PyObject *me_key; /* 8 bytes */
4     PyObject *me_value; /* 8 bytes */
5 } PyDictKeyEntry;

```

To maintain $O(1)$ lookups, the hash table was kept sparse with a maximum fill factor of $2/3$. For a dictionary containing 8 entries, CPython had to allocate a table of at least 16 slots. * **Memory Overhead:** Each empty slot in the table required 24 bytes of memory. If a dictionary was large, the amount of wasted memory in unallocated slots was massive:

$$\text{Wasted Memory} = \text{Empty Slots} \times 24 \text{ bytes}$$

Legacy Sparse Array Layout:

```

1 Indices/Entries Table (Sparse):
2 [ Slot 0: Hash | Key* | Val* ] (24 bytes)
3 [ Slot 1: NULL | NULL | NULL ] (Wasted 24 bytes)
4 [ Slot 2: Hash | Key* | Val* ] (24 bytes)
5 [ Slot 3: NULL | NULL | NULL ] (Wasted 24 bytes)
6 [ Slot 4: Hash | Key* | Val* ] (24 bytes)

```

2. The Compact Dictionary Layout

The new design splits the dictionary into two separate arrays: 1. **dk_indices (Sparse Index Array):** A small, sparse array containing integers (indices pointing into the dense array). Depending on the dictionary size, each index is stored as a 1-byte (`int8_t`), 2-byte (`int16_t`), or 4-byte (`int32_t`) integer. 2. **dk_entries (Dense Entry Array):** A dense, contiguous array containing `PyDictKeyEntry` structures (24 bytes each). Every slot in this array is fully packed in the order keys are inserted.

Compact Array Layout (CPython 3.6+):

```

1 dk_indices (Sparse Array of int8_t):
2 [ 0 | -1 | 1 | -1 | 2 ] (5 bytes, where -1 represents an empty slot)
3
4 dk_entries (Dense Array of PyDictKeyEntry):
5 [ Slot 0: Hash | Key* | Val* ] (24 bytes) - Inserts first
6 [ Slot 1: Hash | Key* | Val* ] (24 bytes) - Inserts second
7 [ Slot 2: Hash | Key* | Val* ] (24 bytes) - Inserts third

```

Dynamic Lookup Walkthrough: To look up a key (e.g. looking up `key` with hash value mapped to sparse index 2): 1. CPython computes the key’s hash and hashes it to a slot in `dk_indices`. 2. If `dk_indices[hash_slot]` is 1, it retrieves the entry at `dk_entries[1]`. 3. It compares the keys. If matched, it returns the value. If there is a collision, it continues probing within the sparse `dk_indices` array.

3. Core Structural Implications

- **Memory Savings:** Instead of wasting 24 bytes per empty slot, CPython now only wastes 1, 2, or 4 bytes per empty slot in `dk_indices`. This layout reduces dictionary memory footprints by **30% to 40%** in real-world workloads.
- **Preservation of Insertion Order:** Because entries are appended to `dk_entries` contiguously, the elements are stored in their exact insertion order. Iteration over the dictionary simply traverses the dense `dk_entries` array sequentially.
- **Faster Iteration:** Traversal does not require skipping empty slots, making dictionary iteration significantly faster due to CPU cache line friendliness.
- **Language Specification Guarantee:** While insertion order was introduced as an implementation detail in Python 3.6, it was officially codified as a language specification guarantee in Python 3.7.

Chapter 15

Python 3.7: Dataclasses, Context Variables, and Dict Ordering Guarantees

15.0.1 13.1 PEP 557: Dataclasses under the hood

Introduced in Python 3.7, the `@dataclass` decorator automates the generation of boilerplate methods (such as `__init__`, `__repr__`, and `__eq__`) by inspecting class type annotations.

1. Import-Time Code Generation Mechanics

Rather than intercepting attribute access or performing dynamic lookups at runtime, `@dataclass` is an import-time code generator: 1. **Annotation Reading:** When the class is imported, the decorator scans the class's `__annotations__` dictionary to identify the fields and their annotated types. 2. **String Code Assembly:** The decorator constructs a string representation of the Python source code for the requested magic methods. For example: `python def __init__(self, name: str, age: int): self.name = name self.age = age` 3. **Compilation and Execution:** It compiles this source string using the built-in `compile(source, "<string>", "exec")` function. It then executes the compiled code block via `exec()` within the class's namespace context to instantiate real function objects. 4. **Attachment:** The compiled function objects are attached directly to the class's dictionary (`__dict__`).

Because these methods are compiled into native CPython bytecode at import time, calling them at runtime incurs zero wrapper or interception overhead, executing at the exact same speed as manually written boilerplate methods.

2. Specialized Flags and post-init Hooks

- **Immutability (`frozen=True`):** If a class is declared with `@dataclass(frozen=True)`, the decorator dynamically generates custom `__setattr__` and `__delattr__` methods: `python def __setattr__(self, name, value): if type(self) is cls: raise FrozenInstanceError(f"cannot assign to field {name!r}") super().__setattr__(name, value)` This prevents direct attribute mutation. *Note:* Developers can bypass this protection by modifying attributes via `object.__setattr__(self, name, value)`.
- **The `__post_init__` Hook:** If the class defines a `__post_init__` method, the generated `__init__` function automatically appends a call to `self.__post_init__()` as its final instruction. This is useful for validating fields or computing dependent properties after initialization.

15.0.2 13.2 PEP 567: Context Variables (`contextvars`)

PEP 567 introduced context variables to provide a safe, high-performance mechanism for managing task-local state in asynchronous runtimes.

1. The Thread-Local Storage Flaw in Async Code

Historically, multi-threaded programs isolated state using thread-local storage (`threading.local`). However, this model breaks down in asynchronous runtimes (like `asyncio`):

- * In `asyncio`, many independent tasks run concurrently on a **single OS thread**, yielding control back and forth.
- * If task A sets a value in thread-local storage and awaits an I/O operation, control yields to task B on the same thread. Task B can read or overwrite task A's thread-local value, causing severe state leakage and race conditions.

2. Hash Array Mapped Trie (HAMT) Internals

Context Variables (`contextvars`) solve this by isolating state per-task using a C-level **Hash Array Mapped Trie (HAMT)** data structure:

- * **Immutability**: A HAMT is an immutable, persistent key-value mapping structure.
- * **Structural Sharing**: When a task modifies a context variable (`var.set(val)`), CPython does not modify the mapping in-place. Instead, it creates a new trie node representing the updated state. Unchanged branches are shared with the previous trie structure (structural sharing), minimizing memory allocation.
- * **Performance**: Due to its shallow 32-way branching factor, HAMT guarantees $O(\log_{32} N) \approx O(1)$ lookup, insertion, and deletion speeds.
- * **$O(1)$ Context Switching**: Because the context state is represented by an immutable HAMT node, saving or restoring a task's context during an `await` suspension is as simple as copying a pointer to the root trie node.

```

1 HAMT Structural Sharing on Update:
2     [ Root A ]                     [ Root B ] (New Context)
3     /       \                     /       \
4     [ Node1 ] [ Node2 ]           [ Node1 ] [ Node3 ] (New/Updated Entry)
5                               (Shared Branch)

```

3. Task Context Switching

In `asyncio`, each `Task` holds a reference to its own `contextvars.Context` object containing the HAMT structure.

1. When suspending at an `await` point, the task yields.
2. Before resuming a different task, the event loop calls `PyContext_Enter(new_context)` at the C level.
3. This restores the exact task-local state variables for the active task in $O(1)$ time, completely isolating concurrent execution paths.

15.0.3 13.3 Dictionary Insertion-Order Guarantee

In Python 3.6, compact dictionaries preserved insertion order as an implementation side-effect. In Python 3.7, this behavior was officially codified as a language specification guarantee.

1. Scope of the Guarantee

The order preservation applies to:

- * Dictionary iteration (e.g. `for key in my_dict`).
- * Views returned by `.keys()`, `.values()`, and `.items()`.
- * Keyword arguments passing (`**kwargs` preserves call-site order).
- * Class namespaces (attributes declared inside class definitions preserve their order inside `__dict__` and `__annotations__`).

2. Impact on the Ecosystem

- **JSON Serialization:** Serialization (`json.dumps`) becomes deterministic by default.
- **Deprecation of `collections.OrderedDict`:** While `OrderedDict` remains in the standard library (offering specialized operations like `.move_to_end()`), standard dictionaries are now preferred for general ordered mapping operations.

15.0.4 13.4 Properties & Cached Properties (Descriptor Protocol)

The descriptor protocol defines how Python handles attribute lookup on objects. Properties and cached properties leverage this protocol to customize attribute access.

1. Data vs. Non-Data Descriptors

A descriptor is an object that implements one or more of the methods: `__get__`, `__set__`, or `__delete__`. * **Data Descriptor:** Implements `__set__` or `__delete__`. * **Non-Data Descriptor:** Only implements `__get__`.

2. CPython Attribute Resolution Precedence

When resolving `instance.attribute`, the C-level function `PyObject_GenericGetAttr` searches namespaces in a strict order of precedence:

1. **Class Search:** Search the class MRO for a descriptor. If a descriptor is found and it is a **data descriptor**, call its `__get__` slot and return the result.
2. **Instance Dictionary:** If no data descriptor is found, check the instance's dictionary (`instance.__dict__`). If the attribute exists, return the value.
3. **Non-Data Descriptor:** If the attribute is not in `instance.__dict__` but a **non-data descriptor** was found on the class, call its `__get__` slot and return the result.
4. **Class Attribute:** Check if the attribute exists as a standard class attribute.
5. **AttributeError:** Raise `AttributeError`.

3. How `functools.cached_property` Exploits Precedence

The built-in `@property` is a **data descriptor** (it defines `__get__` and a default `__set__` that raises an error). Thus, it always overrides instance dictionary lookups.

In contrast, `functools.cached_property` is implemented as a **non-data descriptor** (it only implements `__get__`).

Conceptual Implementation of `cached_property`:

```

1 class cached_property:
2     def __init__(self, func):
3         self.func = func
4         self.__doc__ = func.__doc__
5
6     def __get__(self, instance, owner):
7         if instance is None:
8             return self
9         # Compute the value
10        value = self.func(instance)
11        # Write directly to the instance dict
12        instance.__dict__[self.func.__name__] = value
13        return value

```


Lookup Flow:

1. **First Lookup:** `instance.attribute` is queried. CPython MRO search finds `cached_property` (non-data descriptor). Since it is not in the instance `__dict__`, CPython executes `__get__`. The method computes the value, writes it into `instance.__dict__`, and returns it.
2. **Subsequent Lookups:** CPython MRO search finds `cached_property` (non-data descriptor). CPython then checks the instance `__dict__` and finds the cached value. Because instance dictionary lookups take precedence over non-data descriptors, the value is returned directly from `__dict__`, bypassing `__get__` completely.

15.0.5 13.5 Instance Slots Optimization (`__slots__`)

By default, every object instance allocates a dictionary (`__dict__`) to store attributes dynamically. This consumes significant memory.

1. Struct Layout Changes

Defining `__slots__ = ('x', 'y')` inside a class alters CPython's object layout: * The type constructor (`PyType_Ready`) suppresses the allocation of `__dict__` and `__weakref__` pointers in the instance's structure. * Instead, CPython reserves raw `PyObject*` pointer array slots directly inside the object's struct layout (immediately following `PyObject_HEAD`) for `x` and `y`.

2. Member Descriptors

For each slot defined in `__slots__`, CPython creates a `member_descriptor` object on the class: * The descriptor stores a hardcoded byte offset indicating where the variable's pointer resides relative to the head of the object structure in C memory. * When executing `obj.x`, the member descriptor accesses the pointer at `(char*)obj + offset` directly, replacing slow string hashing and dictionary lookups with fast C-level array indexing.

3. Trade-offs and Limitations

- **Memory Savings:** Eliminates the 100-150 byte overhead of the instance `__dict__` hash table, allowing developers to scale to millions of lightweight objects.
- **Attribute Blocking:** Prevents developers from dynamically adding arbitrary new attributes at runtime (raises `AttributeError`).
- **Weak References:** Since `__weakref__` is omitted, objects cannot be targets of weak references unless `'__weakref__'` is explicitly included in the `__slots__` tuple.
- **Subclassing Rules:** Subclasses do not inherit `__slots__`. If a subclass does not declare `__slots__`, it will allocate a standard `__dict__` and `__weakref__`, rendering the parent's memory savings moot for subclass instances.

```

1 import sys
2
3 class DictClass:
4     def __init__(self, x, y):
5         self.x = x
6         self.y = y
7
8 class SlotsClass:
9     __slots__ = ('x', 'y')
10    def __init__(self, x, y):
11        self.x = x

```

```
12         self.y = y
13
14 d = DictClass(1, 2)
15 s = SlotsClass(1, 2)
16
17 print("DictClass size:", sys.getsizeof(d) + sys.getsizeof(d.__dict__))
18 # SlotsClass has no __dict__, consuming significantly less memory
19 print("SlotsClass size:", sys.getsizeof(s))
```

Part V

Volume V: Structural Shifts & Pattern Matching

Chapter 16

Python 3.8: Walrus Operator (:=) and Positional-Only Parameters (/)

16.0.1 14.1 The Assignment Expression (:= / PEP 572)

Introduced in Python 3.8, the assignment expression operator (colloquially called the **walrus operator**) allows assigning values to variables inside expressions. This departs from Python's historical strict division between statements and expressions.

1. AST Representation

A standard assignment is a statement represented by the `Assign` node in the Abstract Syntax Tree (AST), which does not return a value. The walrus operator compiles to a `NamedExpr` node, which is an expression node returning the bound value.

	[Assign Statement]		[NamedExpr Expression]
1	/		/
2	\		\
3	Name(x) Constant(1)		Name(x) Constant(1)
4	(Returns void/no value)		(Returns 1 to parent context)

2. Bytecode & Stack Operations

Let's examine how CPython processes a normal assignment vs. an assignment expression. Consider the following source codes and their corresponding bytecodes:

```
1 # Standard Assignment
2 x = 1
```

```
1 2 LOAD_CONST          0 (1)
2 4 STORE_FAST          0 (x)
```

```
1 # Assignment Expression
2 (x := 1)
```

CPython 3.8 Bytecode:

```
1 2 LOAD_CONST          0 (1)
2 4 DUP_TOP
3 6 STORE_FAST          0 (x)
```

CPython 3.11+ Optimized Bytecode:

```

1 2 LOAD_CONST          0 (1)
2 4 COPY                1
3 6 STORE_FAST          0 (x)

```

The execution steps for the walrus operator stack operation are: 1. **LOAD_CONST**: Pushes the reference to the constant integer 1 onto the value stack. 2. **DUP_TOP / COPY 1**: Replicates the top element of the stack, resulting in two references to 1 on the stack. 3. **STORE_FAST**: Pops one of the references and binds it to the local variable name `x`. 4. **Stack Residual**: The remaining reference to 1 remains on the top of the stack, allowing parent expressions (such as `if` conditionals or loops) to consume it.

3. Scoping Rules and Symbol Table Traversal

To prevent variable leakage and namespace pollution, PEP 572 specifies complex scoping rules, especially within comprehensions: * **Comprehensions**: Python list, set, and dict comprehensions, as well as generator expressions, are executed in a nested function scope. * **Variable Binding (Hoisting)**: An assignment expression `x := ...` inside a comprehension binds the variable `x` in the *surrounding* (parent) scope, not the local comprehension scope. * **Symbol Table Resolution**: During compilation, the symbol table builder (`symtable.c`) traverses variables inside comprehensions. When it encounters a `NamedExpr`, it marks the target variable as a **free variable** in the comprehension scope and a **cell variable** in the parent scope. This hoists the reference out of the comprehension's inner namespace.

Let's examine this hoisting in action. Consider the following code:

```

1 def parent_func(data):
2     result = [y for x in data if (y := x * 2) > 2]
3     return result, y

```

Disassembly of parent_func:

```

1 2      0 LOAD_CLOSURE          0 (y)          /* Load the cell
   reference for y */
2      2 BUILD_TUPLE
3      4 LOAD_CONST              1
4      6 LOAD_CONST              1 (<code object <listcomp>...>)
5      8 MAKE_FUNCTION           2 ('parent_func.<locals>.<listcomp>')
   function closure */
6     10 LOAD_FAST               0 (data)
7     12 GET_ITER
8     14 CALL_FUNCTION            1
9     16 STORE_FAST              1 (result)
10
11 3     18 LOAD_FAST              1 (result)
12    20 LOAD_DEREF              0 (y)          /* Load y from closure
   cell */
13    22 BUILD_TUPLE              2
14    24 RETURN_VALUE

```

Disassembly of the nested <listcomp> code object:

```

1 2      0 BUILD_LIST              0
2      2 LOAD_FAST               0 (.0)        /* Load incoming
   iterator */
3      4 FOR_ITER                 26 (to 32)
4      6 STORE_FAST              1 (x)
5      8 LOAD_FAST               1 (x)
6     10 LOAD_CONST              1 (2)
7     12 BINARY_MULTIPLY
8     14 DUP_TOP

```

```

9      16 STORE_DEREF          0 (y)          /* Store into outer
    scope cell */
10      18 LOAD_CONST          2 (2)
11      20 COMPARE_OP          4 (>)
12      22 POP_JUMP_IF_FALSE   4
13      24 LOAD_DEREF          0 (y)          /* Load y from outer
    cell */
14      26 LIST_APPEND          2
15      28 JUMP_ABSOLUTE        4
16      32 RETURN_VALUE

```

By using `STORE_DEREF` and `LOAD_DEREF`, the nested list comprehension writes directly to the parent scope cell, allowing `y` to survive after the list comprehension has finished executing.

4. Compiler Restrictions

To prevent syntactically ambiguous or unstable code, the compiler enforces strict boundaries:

* **Iteration Variable Conflicts:** A target variable cannot share a name with a comprehension iteration variable. E.g. `[i for i in range(10) if (i := 2)]` raises `SyntaxError: assignment expression cannot rebind comprehension iteration variable`. * **Class Scope Prohibitions:** The walrus operator is explicitly blocked inside comprehensions in class bodies: `python class MyClass: data = [1, 2, 3] result = [y for x in data if (y := x * 2) > 2]` This raises `SyntaxError: assignment expression within a comprehension cannot be used in a class body`. *Why?* Class namespaces are evaluated as temporary dict namespaces, not standard function frames. They do not support closure cells (`PyCellObject`). Since the comprehension runs as a separate nested helper function, it cannot bind closures to class body locals, creating a scoping mismatch.

16.0.2 14.2 Positional-Only Parameters (/ / PEP 570)

Python 3.8 introduced the `/` marker to define positional-only parameters. Any parameters declared before the `/` marker cannot be passed as keyword arguments.

1. C-level Representation in PyCodeObject

The compiler encodes parameter boundaries directly into the function's bytecode descriptor struct. In `code.h`, the `PyCodeObject` structure contains specific integer fields to track positional boundaries:

```

1 typedef struct {
2     PyObject_HEAD
3     int co_argcount;          /* Number of positional and positional-only
    arguments */
4     int co_posonlyargcount;    /* Number of positional-only arguments */
5     int co_kwonlyargcount;    /* Number of keyword-only arguments */
6     /* ... additional fields ... */
7 } PyCodeObject;

```

For a function definition:

```

1 def func(a, b, /, c, d, *, e, f):
2     pass

```

The compiler maps the arguments as follows: * `a, b`: Positional-only. (`co_posonlyargcount` = 2) * `c, d`: Positional-or-keyword. (Total positional `co_argcount` = 4, representing positional-only + positional-or-keyword) * `e, f`: Keyword-only. (`co_kwonlyargcount` = 2)

2. Argument Parsing and Vectorcall Optimization (PEP 590)

When a function is called, CPython passes arguments using the **Vectorcall** calling protocol introduced in Python 3.8:

```
1 PyObject *vectorcall(PyObject *callable, PyObject *const *args, size_t nargsf,
    PyObject *kwnames);
```

Where: `*args` is a pointer to a contiguous array containing all passed argument values. `*nargsf` specifies the number of positional arguments. `*kwnames` is a tuple of strings containing the keyword argument names. The keyword values are stored in the `args` array starting at index `nargsf`.

The Vectorcall Matching Algorithm:

1. **Keyword Verification:** CPython parses the keyword strings in `kwnames`. If any string matches a parameter name located before index `co_posonlyargcount` (e.g. trying to pass `a=1`), the interpreter raises a `TypeError`.
2. **Bypassing Dictionary Checks:** Because positional-only parameters are guaranteed to be passed positionally, CPython completely bypasses string-matching hash loops for the first `co_posonlyargcount` arguments.
3. **Direct Array Offset Writing:** The VM copies references directly from the `args` array into the local frame variable cells by array index offset:

$$\text{LocalFrameOffset}[i] = \text{args}[i] \quad \text{for } 0 \leq i < \text{co_posonlyargcount}$$

This direct offset copy cuts down calling overhead by **10% to 15%**, which is highly beneficial for builtins and low-level utility functions that are called repeatedly.

Chapter 17

Python 3.9 to 3.10: PEG Parser, Dict Merge (`|`), and Pattern Matching

17.0.1 15.1 CPython Parser Evolution: LL(1) to PEG (PEP 617)

In Python 3.9, CPython replaced its concrete parser generator (`pgen`), which had been used since Python 1.0, with a modern **PEG (Parsing Expression Grammar)** parser.

1. Limitations of the Classic LL(1) Parser

The legacy compiler utilized an LL(1) (Left-to-right, Leftmost derivation with 1-token lookahead) parsing grammar:

- * **Lookahead Constraints:** With only one token of lookahead, the parser could not resolve ambiguities between grammar rules that looked similar initially.
- * **Left-Recursion Prohibitions:** LL(1) cannot parse left-recursive rules (e.g., $A \rightarrow A\alpha$), requiring complex grammar rewrites that made grammar maintenance difficult.
- * **Grammar Hacks:** To support complex syntax, developers had to implement custom AST-level hacks and validation passes after parsing, which bloated parser code.

2. The PEG Parser Architecture

PEG parsers resolve these limitations through two core features:

- * **Ordered Choice (`e1 / e2`):** Unlike CFGs (Context-Free Grammars) where choice is ambiguous, a PEG ordered choice evaluates `e1` first. If `e1` matches, the parser commits to it and completely ignores `e2`, resolving parsing ambiguities deterministically.
- * **Infinite Lookahead via Packrat Parsing:** PEG can look ahead arbitrarily far. To maintain linear time complexity $O(N)$ (where N is input length), the parser utilizes **Packrat Parsing** (memoization). It caches parser rules evaluation results at every character index offset in a memoization table (`ParserState` in `Parser/pegen.c`).

3. Enabling Syntax Improvements

The infinite lookahead and recursive capabilities of PEG enabled new syntax structures in Python 3.10, such as **parenthesized context managers**:

```
1 with (  
2     CtxManager1() as ctx1,  
3     CtxManager2() as ctx2,  
4 ):  
5     pass
```


Under the old LL(1) parser, parsing this construct was impossible because it could not distinguish between a parenthesized tuple and a grouped context manager block without infinite lookahead.

17.0.2 15.2 PEP 584: Dictionary Union Operators (`|` and `|=`)

Python 3.9 introduced the binary operators `|` (merge) and `|=` (update) directly on the `dict` class.

1. Legacy Workarounds

Prior to Python 3.9, merging dictionaries required verbose and slow operations: `*merged = {**d1, **d2}`: Highly performant but syntactically verbose and unreadable for larger expressions. `*merged = d1.copy(); merged.update(d2)`: Required multiple lines of statements, making it impossible to merge inline inside list comprehensions or lambda bodies.

2. Low-Level C-Level Slot Mappings

CPython implements `|` and `|=` by overloading the numeric bitwise OR slots inside `PyDict_Type` (`Objects/dictobject.c`):

```
1 /* Dict Type definition slot mappings */
2 PyTypeObject PyDict_Type = {
3     /* ... */
4     &dict_as_number,          /* tp_as_number */
5     /* ... */
6 };
7
8 static PyNumberMethods dict_as_number = {
9     .nb_or = (binaryfunc)dict_or,
10    .nb_inplace_or = (binaryfunc)dict_inplace_or,
11 };
```

When evaluating `A | B`: 1. CPython calls `dict_or(A, B)`. 2. The C function allocates a new dictionary object: `PyObject *new_dict = PyDict_Copy(A)`; 3. It then calls the dictionary update routine `PyDict_Update(new_dict, B)`; to merge the elements from the second mapping. 4. It returns the new dictionary. Right-hand elements overwrite duplicate keys present in the left-hand dictionary.

3. Type Constraints and Subclassing

- **Return Type:** `dict_or` always returns a standard `dict` instance, even if `A` or `B` is a subclass of `dict`, preserving typing boundaries.
- **Operand Support:** The left-hand operand `A` must be a `dict` (or subclass). The right-hand operand `B` can be any object that implements the mapping protocol (exposes a `.keys()` and key retrieval interface). If `B` is not a mapping, the C-level function returns `Py_NotImplemented`, triggering a runtime `TypeError`.

17.0.3 15.3 PEP 634: Structural Pattern Matching (Decision Trees)

Python 3.10 introduced Structural Pattern Matching (`match/case`).

1. Decision DAGs vs. Sequential Linear Scans

Unlike chains of `if/elif/else` statements, which execute linearly ($\mathcal{O}(N)$ lookup time), the CPython compiler compiles a `match` statement into a **Directed Acyclic Graph (DAG) decision tree**. The compiler groups cases matching the same pattern categories, checking target type boundaries and structural lengths once rather than executing redundant checks.

2. CPython Pattern Matching Bytecodes

CPython implements pattern matching using specialized stack-manipulation bytecodes: `* MATCH_SEQUENCE`: Checks if the object on the stack is a sequence (excluding `str`, `bytes`, and `bytearray`). `* MATCH_MAPPING`: Checks if the object is an instance of `collections.abc.Mapping`. `* MATCH_KEYS`: Pops a tuple of keys and a subject mapping. If all keys exist in the mapping, it pushes a tuple containing their values; otherwise, it pushes `None`. `* MATCH_CLASS`: Evaluates class-level matches.

3. Class Patterns and `MATCH_CLASS` Disassembly Trace

To match a class pattern:

```

1 class Point:
2     __match_args__ = ('x', 'y')
3     def __init__(self, x, y):
4         self.x = x
5         self.y = y
6
7 def check_point(pt):
8     match pt:
9         case Point(1, y):
10         return y

```

During execution, `case Point(1, y)` requires verifying that `pt` is an instance of `Point`, matching the first positional argument `x` to `1`, and binding the second argument `y` to local scope.

Disassembly of `check_point`:

1	8	0	LOAD_FAST	0	(pt)
2		2	LOAD_GLOBAL	0	(Point)
3		4	MATCH_CLASS	1	/* Match Point class
			with 1 positional arg */		
4		6	DUP_TOP		
5		8	LOAD_CONST	0	(None)
6		10	COMPARE_OP	9	(is not)
7		12	POP_JUMP_IF_FALSE	32	(to case mismatch)
8		14	UNPACK_SEQUENCE	2	/* Unpack matched x and
			y values */		
9		16	LOAD_CONST	1	(1)
10		18	COMPARE_OP	2	(==) /* Check if x == 1 */
11		20	POP_JUMP_IF_FALSE	28	(to case clean stack)
12		22	STORE_FAST	1	(y) /* Bind y to local
			scope */		
13		24	LOAD_FAST	1	(y)
14		26	RETURN_VALUE		
15					
16	>>	28	POP_TOP		/* Clean remaining
			unpacked values */		
17	>>	30	POP_TOP		
18	>>	32	POP_TOP		/* Fallback case
			mismatch target */		
19		34	LOAD_CONST	0	(None)
20		36	RETURN_VALUE		

C-Level Execution of MATCH_CLASS:

1. **Type Validation:** Checks `isinstance(pt, Point)`. If false, pushes `None` and exits.
 2. **Positional Resolution:** Reads the class's `__match_args__` tuple (`'x'`, `'y'`). The compiler indicated 1 positional argument, mapping the first argument to attribute `x`.
 3. **Keyword/Positional Extraction:** Extracts `pt.x` and `pt.y` via C-level attribute lookups.
 4. **Stack Result:** Pushes a tuple `(pt.x, pt.y)` onto the stack. If any attribute extraction fails, it pushes `None`.
 5. **Fail-Fast Stack Cleanup:** If `None` is pushed, `COMPARE_OP` determines a mismatch, jumps to the fallback, and pops the stack, leaving no bindings.
-

Chapter 18

Python 3.8 to 3.10: Type Hinting Protocols and Structural Subtyping

18.0.1 16.1 Nominal vs. Structural Subtyping

Python type hints support two distinct typing paradigms: Nominal Subtyping and Structural Subtyping (implemented via **Protocols** / PEP 544).

1. Nominal Subtyping

Nominal subtyping resolves type compatibility based on explicit inheritance hierarchies. A class `Dog` is considered a subtype of `Animal` if and only if it explicitly inherits from it:

```
1 class Animal:
2     def breathe(self) -> None:
3         pass
4
5 class Dog(Animal): # Explicitly nominal
6     pass
```

Under this model, even if a class defines all methods of `Animal`, static type checkers will reject it if there is no explicit base class relationship.

2. Structural Subtyping (Static Duck Typing)

Structural subtyping resolves type compatibility based on the structure (methods and attributes) of the class rather than its name. This is defined using `typing.Protocol`:

```
1 from typing import Protocol
2
3 class Renderable(Protocol):
4     def render(self) -> str:
5         ...
6
7 class Book:
8     # Book does NOT explicitly inherit from Renderable
9     def render(self) -> str:
10         return "Book Text"
11
12 def display(item: Renderable) -> None:
13     print(item.render())
14
15 # Static type checkers accept this call
16 display(Book())
```

Type checkers verify that the interface `Book` implements all attributes and methods declared in the protocol `Renderable` with matching signatures and types.

18.0.2 16.2 Runtime Protocols & `@runtime_checkable`

By default, protocols are static constructs erased during compilation. At runtime, evaluating `isinstance(Book(), Renderable)` raises a `TypeError`. However, decorating a protocol with `@runtime_checkable` enables standard runtime checks.

1. The `_ProtocolMeta` Metaclass Internals

When a class inherits from `typing.Protocol`, CPython uses the custom metaclass `typing._ProtocolMeta` (Lib/typing.py). 1. **Attribute Collection**: During class creation, `_ProtocolMeta` scans the namespace dict and parent MRO to collect all declared non-dunder attributes (methods, properties, and variable annotations). 2. **Caching**: It caches these names in a private set on the class structure named `_protocol_attrs`.

2. Metaclass Dunder Overrides

`@runtime_checkable` overrides the dunder methods of `_ProtocolMeta`: `* __instancecheck__`: Overrides `isinstance()`. `* __subclasscheck__`: Overrides `issubclass()`.

CPython `__subclasscheck__` Protocol Logic:

```

1 def __subclasscheck__(cls, subclass):
2     if not isinstance(subclass, type):
3         raise TypeError("issubclass() arg 1 must be a class")
4
5     # Fast path: nominal subclass checks
6     if super().__subclasscheck__(subclass):
7         return True
8
9     # Slow path: structural MRO checking
10    for attr in cls._protocol_attrs:
11        # Check if attribute exists in the subclass dictionary or any of its
12        # MRO parents
13        for entry in subclass.__mro__:
14            if attr in entry.__dict__:
15                break
16        else:
17            return False # Attribute missing in the entire MRO chain
18    return True

```

Critical Runtime Considerations:

[!WARNING] * **Omission of Type Safety**: Runtime `isinstance` checks only verify the **existence** of the attribute names. They do not inspect method signatures, parameter counts, or variable types. A class implementing `def render(self, a, b): pass` will successfully pass an `isinstance(obj, Renderable)` check, even though it violates the protocol signature statically. * **Performance Penalties**: Traversing the `__mro__` and inspecting `__dict__` for every protocol attribute introduces considerable execution overhead. Avoid using runtime protocol checks inside critical loops.

18.0.3 16.3 Static Type Checking and Variance

Static type checkers compile type annotations into a directed graph of subtypes to enforce type boundaries. The relationships between generic types depend on **Variance**.

1. Mathematical Definition of Variance

Let A and B be types where A is a subtype of B ($A \subseteq B$). Let F be a generic container type constructor (e.g., `List[T]`, `Iterable[T]`, or `Callable[[T], R]`). * **Covariance**: The container type preserves the subtype relationship:

$$A \subseteq B \implies F(A) \subseteq F(B)$$

In Python, this is defined via `T = TypeVar('T', covariant=True)`. *Example*: `Iterable[T]` is covariant. Because `Dog` $\subseteq$ `Animal`, `Iterable[Dog]` $\subseteq$ `Iterable[Animal]`. A list of dogs can be safely read as a list of animals. * **Contravariance**: The container type reverses the subtype relationship:

$$A \subseteq B \implies F(B) \subseteq F(A)$$

In Python, this is defined via `T = TypeVar('T', contravariant=True)`. *Example*: `Callable[[T], None]` arguments are contravariant. A function that accepts any `Animal` can be safely used where a function accepting a `Dog` is expected:

$$\text{Callable}[[\text{Animal}], \text{None}] \subseteq \text{Callable}[[\text{Dog}], \text{None}]$$

* **Invariance**: There is no relationship between container types:

$$F(A) \not\subseteq F(B) \quad \text{and} \quad F(B) \not\subseteq F(A)$$

In Python, standard generic classes are invariant by default. *Example*: `List[T]` is invariant because it allows both read (covariant) and write (contravariant) operations. Allowing a list of dogs to be treated as a list of animals could let someone insert a cat into the list, violating type safety.

2. Liskov Substitution Principle (LSP)

The variance of function arguments and return values is derived from the **Liskov Substitution Principle**: if S is a subtype of T , then objects of type T may be replaced with objects of type S without altering any of the desirable properties of the program.

Applying LSP to function subtypes leads to this mapping:

$$A \subseteq B \text{ and } C \subseteq D \implies (B \rightarrow C) \subseteq (A \rightarrow D)$$

This implies that for a function to substitute another: 1. **Arguments must be Contravariant**: It must accept a broader set of inputs. (Narrowing input scope is unsafe). 2. **Return values must be Covariant**: It must guarantee a narrower set of outputs. (Broadening output scope is unsafe).

```

1 Liskov Function Subtyping Mapping:
2 Input (Contravariant):   Animal (Parent)   ----->   Dog (Child)
3                               |
4                               v
5 Output (Covariant):      Dog (Child)       ----->   Animal (Parent)

```

3. Compile-Time Type Erasure

CPython compiles type annotations into bytecode, but completely ignores them during execution (unless checked by runtime libraries via `__annotations__`). Static type checking is entirely completed during pre-run compilation. Once bytecode is emitted, generic variables (e.g. `List[int]`) revert back to standard untyped collections (e.g. `List`) in memory, ensuring that typing extensions introduce zero runtime memory overhead.

Part VI

Volume VI: Performance Leap & Runtime Mechanics

Chapter 19

Python 3.11: Faster CPython Specializing Interpreter and Adaptive Bytecode

19.0.1 17.1 PEP 659 Specialized Adaptive Interpreter Architecture

Prior to Python 3.11, the bytecode interpreter loop executed instructions generically. For instance, a `LOAD_ATTR` instruction looked up attributes using hash table lookups every time. PEP 659 introduced a **specializing adaptive interpreter** to dramatically reduce lookup overhead.

1. Monomorphism vs. Polymorphism in Dynamic Execution

While Python is a dynamic language, execution paths in real-world application code are highly **monomorphic** (a specific call site or attribute access typically processes objects of the exact same type repeatedly). Generic bytecode instructions (like `LOAD_ATTR`) are designed to handle any type of object, performing expensive CPython namespace dictionary lookups and method searches on every execution.

2. Dynamic Instruction Patching

PEP 659 optimizes hot, monomorphic paths by dynamically modifying the instruction stream in memory at runtime:

1. **Generic State:** The compiler generates generic opcodes (e.g. `LOAD_ATTR`).
2. **Warm-up & Monitoring:** When executed, the generic opcode transitions into an **adaptive state** (e.g., `LOAD_ATTR_ADAPTIVE`). The adaptive instruction maintains an execution counter initialized to 8.
3. **Specialization:** On every execution, the interpreter monitors the operand types and decrements the counter. If the type is consistent when the counter reaches 0, the interpreter patches the bytecode in memory, replacing the adaptive opcode with a **specialized opcode** (e.g., `LOAD_ATTR_INSTANCE_VALUE`).
4. **Deoptimization:** If the operand type changes at a specialized call site, the opcode execution fails its fast path validation check. It branches to a deoptimization routine, restores the generic or adaptive state, and executes the slow lookup.

```
1 CPython 3.11 Instruction Specialization Lifecycle:
2
3   [ Generic Opcode ] (LOAD_ATTR)
4       |
5       | (Executed first time)
6       v
7   [ Adaptive Opcode ] (LOAD_ATTR_ADAPTIVE, Counter = 8)
8       |
9       | (Consistent types, Counter reaches 0)
```

```

10      v
11  [ Specialized Opcode ] (LOAD_ATTR_INSTANCE_VALUE)
12      /                  \
13  / (Type Match)         \ (Type Mismatch)
14  v                      v
15  [ Fast Direct Offset ] [ Deoptimize to Generic / Slow Lookup ]

```

19.0.2 17.2 Specialized Bytecode Opcodes and Inline Cache Memory Layout

CPython reserves extra memory slots directly adjacent to the instruction bytes inside the code object's instruction array. This contiguous memory block is called the **Inline Cache**.

1. Inline Cache Memory Alignment

Instructions are represented by `_Py_CODEUNIT` elements (16-bit blocks: 8-bit opcode and 8-bit argument). When a cacheable instruction is compiled, the compiler allocates empty `_Py_CODEUNIT` entries directly following the instruction to serve as the cache.

CPython Bytecode Stream Cache Layout:

```

1 Memory Address: | n | n+1 | n+2 | n+3 | n+4 | n+5 | n+6 |
2 Bytecode:      [ LOAD_ATTR ] [ Cache Slot 0 ] [ Cache Slot 1 ] [ Cache Slot 2 ]

```

When executing `LOAD_ATTR` at address `n`, the VM knows that the next 3 code units are reserved for the inline cache and skips them to decode the next instruction at address `n+4`.

2. C-level Cache Structures inside `pycore_code.h`

For attribute lookups (`LOAD_ATTR`), CPython maps the inline cache slots to the `_PyAttrCache` struct:

```

1 typedef struct {
2     uint16_t counter;           /* Adaptive execution counter */
3     uint32_t type_version;     /* Object type version pointer (tp_version_tag)
4     /*
5     uint16_t index;             /* Attribute array offset */
6 } _PyAttrCache;

```

For global variable lookups (`LOAD_GLOBAL`), the VM allocates the `_PyLoadGlobalCache` struct:

```

1 typedef struct {
2     uint16_t counter;           /* Adaptive counter */
3     uint32_t module_keys_version; /* Dictionary keys version of module globals
4     /*
5     uint32_t builtin_keys_version; /* Dictionary keys version of builtins dict
6     /*
7 } _PyLoadGlobalCache;

```

3. CPython Type Versioning (`tp_version_tag`)

To ensure cache validity: * Every class (`PyTypeObject`) in CPython has an internal `tp_version_tag` field (a 32-bit unsigned integer). * If a class is modified at runtime (e.g. adding a method or setting a class variable), CPython increments a global type version counter and assigns it to the class's `tp_version_tag`, immediately invalidating all inline caches containing the old version tag.

4. Execution Flow of Specialized Instructions

- **LOAD_ATTR_INSTANCE_VALUE:**
 1. Pops the target object off the stack.
 2. Compares the object's class `tp_version_tag` against the cached `type_version`.
 3. If they match, it skips dict lookups and reads the attribute directly from the object's instance values array (`obj->obj_values[index]`), pushing it onto the stack.
- **LOAD_GLOBAL_MODULE:**
 1. Compares the module dictionary's keys version against `module_keys_version`.
 2. If they match, it reads the value directly from the dictionary value array, completely bypassing hash collisions and key comparisons.

19.0.3 17.3 Polymorphic Deoptimization Costs

While specialization yields massive performance improvements, it is sensitive to code polymorphism: * **Deoptimization Paths:** If a call site receives objects of varying classes (e.g. calling a function with both `Dog` and `Cat` classes where both define `.x` but at different offsets), the instruction will constantly deoptimize. * **Thrashing:** The VM will oscillate between deoptimizing, executing slow-path lookups, warming up the adaptive counter, and specializing again (thrashing). This destroys performance because the interpreter incurs the combined costs of: 1. Cache tag validation mismatch checks. 2. Generic name dictionary hash lookups. 3. Adaptive state tracking and bytecode memory rewriting overhead.

Coding Guidelines for PEP 659 Optimization:

To maximize CPython 3.11+ execution speed, developers should design code to be **monomorphic**: 1. **Keep Types Consistent:** Avoid passing objects of different types to the same hot functions or attribute access sites. 2. **Favor Class Structure Consistency:** Avoid dynamically modifying instance attributes or class namespaces at runtime, as this increments `tp_version_tag` and invalidates caches globally.

Chapter 20

Python 3.12: Native Generics (PEP 695), Type statement, and Subinterpreters

20.0.1 18.1 PEP 695 Type Parameter Syntax

Python 3.12 introduced a clean, native syntax for generic classes, generic functions, and type aliases using type parameters enclosed in square brackets.

1. Transition from Legacy Declarations

Prior to Python 3.12, defining generic types required importing and manually instantiating `TypeVar` variables, along with explicitly defining variance constraints (covariant, contravariant, or invariant):

```
1 # Legacy Python 3.11-
2 from typing import TypeVar, Generic
3
4 T = TypeVar('T', covariant=True) # Manual variance specification
5
6 class Stack(Generic[T]):
7     pass
```

PEP 695 replaces this boiler-plate with a cleaner syntax:

```
1 # Modern Python 3.12+
2 class Stack[T]:
3     pass
```

Under this syntax: * CPython automatically creates the type variable `T` (instance of `typing.TypeVar`) and binds it to the class scope. * **Static Auto-Variance**: Static type checkers (like Mypy/Pyright) automatically infer the variance of `T` by analyzing how the variable is used inside the class definition, completely eliminating the need for manual `covariant=True` or `contravariant=True` flags.

2. Bound and Constraint Specifications

Type parameter bounds and constraints are defined directly inside the square brackets: * **Bound Syntax** (`T: bound_type`): Restricts the type variable to a specific subclass or type. `python def process[T: int](val: T) -> T: return val` * **Constraint Syntax** (`T: (type1, type2, ...)`): Restricts the type variable to one of a set of explicit types. `python def parse[T: (str, bytes)](data: T) -> T: return data`

20.0.2 18.2 Annotation Scopes and Lazy Evaluation

In Python 3.11 and below, type parameter bounds and type aliases were evaluated eagerly at module import time. This caused circular import errors (if a type variable referenced a class defined later in another module) and increased startup memory footprints. PEP 695 resolves this by introducing **Annotation Scopes** and lazy evaluation.

1. Annotation Scopes

When the compiler encounters type parameters or the `type` statement, it wraps their evaluation code inside a new lexical scope called an **Annotation Scope**: * An annotation scope is a nested compiler-generated scope (similar to a hidden function block). * Variables and bounds defined inside this scope are evaluated **lazily** only when they are explicitly queried at runtime.

2. Bytecode Disassembly of the `type` Statement

Let's analyze how CPython compiles a modern type alias:

```
1 type Vector3D[T] = tuple[T, T, T]
```

Compiled Bytecode:

1	1	0	LOAD_CONST	0	('Vector3D')
2		2	LOAD_CONST	1	('T')
3		4	LOAD_CONST	2	(<code object Vector3D at 0x...>)
4		6	MAKE_FUNCTION	0	
5		8	SET_FUNCTION_ATTRIBUTE	8	(closure/annotations helper)
6		10	CALL_FUNCTION	0	
7		12	BUILD_TYPEALIAS		
8		14	STORE_NAME	0	(Vector3D)

- **MAKE_FUNCTION**: Creates a lazy evaluation function from the nested code object representing the alias value `tuple[T, T, T]`.
- **BUILD_TYPEALIAS**: Pops the evaluation function, type parameters, and name, then constructs an instance of `typing.TypeAliasType`.
- **STORE_NAME**: Binds the type alias object to `Vector3D`.

3. C-Level Representation of `TypeAliasType`

At the C level, `TypeAliasType` is represented by a dedicated struct wrapping the properties of the type alias: * `__name__`: Name of the alias. * `__type_params__`: A tuple of type parameters. * `__value__`: The aliased type. It uses a C descriptor getter that evaluates the lazy function code object on the first access, caching the result to avoid redundant evaluations.

20.0.3 18.3 PEP 684: Per-Interpreter GIL and Subinterpreters

CPython has supported subinterpreters via its C-API since Python 1.5. However, because they all shared a single Global Interpreter Lock (GIL), only one interpreter could execute bytecode at a time, preventing true multi-core parallel execution. Python 3.12 introduced **PEP 684**, providing a dedicated, isolated GIL for each subinterpreter.

1. C-Level Structural Separation

CPython separates runtime state into two major structures: * **PyRuntimeState** (`pycore_runtime.h`): Process-wide, shared global resources, including system memory allocators, signal handlers, and GC arena lists. * **PyInterpreterState** (`pycore_pystate.h`): Interpreter-specific resources.

Under PEP 684, each **PyInterpreterState** is allocated its own dedicated GIL struct (`struct _gil_runtime_state`). This allows multiple interpreter instances to execute bytecode in parallel on separate threads, leveraging separate CPU cores.

```

1 CPython Process Layout with Per-Interpreter GIL:
2
3 ===== [ PyRuntimeState ] =====
4 /                                     \
5 [ PyInterpreterState A ]             [ PyInterpreterState B ]
6   - sys.modules                      - sys.modules
7   - Private Object Heap              - Private Object Heap
8   - Dedicated GIL A                  - Dedicated GIL B
9     |                                |
10    v                                v
11   ( OS Thread 1 )                   ( OS Thread 2 )
12 =====

```

2. Isolation Boundaries

To prevent concurrency race conditions, interpreters maintain strict boundaries: * **Private Module Cache**: Each subinterpreter has its own `sys.modules` dictionary. * **Heap Isolation**: Python heap objects (`PyObject` references) cannot be shared directly between interpreters. If interpreter A attempts to read or write a `PyObject` allocated in interpreter B: 1. Reference counting checks (`Py_INCREF/Py_DECREF`) will trigger race conditions. 2. Garbage collection passes running in interpreter A will attempt to collect memory belonging to interpreter B, causing memory corruption. * **Communication Channel**: Subinterpreters communicate exclusively through serialized message channels (such as `_xxsubinterpreters`). Data is serialized (e.g. converted to raw bytes or using shared memory structures via the C buffer protocol) and then re-materialized as fresh objects in the target interpreter's private heap.

3. Parallel Execution Code Example

We can spawn subinterpreters and execute code concurrently using the `_xxsubinterpreters` module:

```

1 import _xxsubinterpreters as interpreters
2 import threading
3
4 def run_in_subinterpreter(interp_id, script):
5     # Run script inside the isolated interpreter context
6     interpreters.run_string(interp_id, script)
7
8 # Create a new subinterpreter with an isolated GIL
9 interp_id = interpreters.create()
10
11 script = """
12 import time
13 # Executes concurrently on a separate CPU core
14 result = sum(i * i for i in range(10_000_000))
15 print("Subinterpreter result calculation complete:", result)
16 """
17
18 # Spawn a separate OS thread to run the subinterpreter in parallel

```

```
19 thread = threading.Thread(target=run_in_subinterpreter, args=(interp_id, script
   ))
20 thread.start()
21 thread.join()
22
23 # Destroy the subinterpreter and release resources
24 interpreters.destroy(interp_id)
```

Chapter 21

Python 3.11 to 3.12: Exception Groups (except*) and Traceback trees

21.0.1 19.1 Exception Groups (ExceptionGroup & BaseExceptionGroup / PEP 654)

Introduced in Python 3.11, `ExceptionGroup` and `BaseExceptionGroup` allow raising and handling multiple unrelated exceptions simultaneously. This is critical for concurrent frameworks (such as `asyncio TaskGroups` or `Trio`) where multiple background operations can crash concurrently.

1. Exception Group Class Layout and Hierarchy

The class hierarchy separates groups of standard exceptions from groups containing system-level exceptions:

```
1 class BaseExceptionGroup(BaseException):
2     def __init__(self, message: str, exceptions: Sequence[BaseException]) ->
      None:
3         self.message = message
4         self.exceptions = list(exceptions)
5
6 class ExceptionGroup(BaseExceptionGroup, Exception):
7     pass
```

- `BaseExceptionGroup` inherits directly from `BaseException`. It can wrap any exception instance, including system-terminating errors like `KeyboardInterrupt`, `SystemExit`, or `GeneratorExit`.
- `ExceptionGroup` inherits from both `BaseExceptionGroup` and `Exception`. It can only wrap exceptions that inherit from `Exception`. If a program attempts to instantiate an `ExceptionGroup` containing a `BaseException` that is not an `Exception` subclass, CPython raises a `TypeError` at runtime.

2. CPython C-Struct Representation

At the C level, exception groups are managed by the `PyExceptionGroupObject` structure, which extends the standard exception header to support the nested tree structure:

```
1 /* Include/cpython/exceptions.h */
2 typedef struct {
3     PyBaseExceptionObject base; /* Base exception header containing dict,
      traceback, context */
```



```

4     PyObject *msg;                /* Message describing the group (
   PyUnicodeObject) */
5     PyObject *excs;               /* Tuple containing the child exception objects
   (PyTupleObject) */
6 } PyExceptionGroupObject;

```

When an exception group is instantiated: 1. `base.args` is populated with a tuple containing the `message` and `exceptions` sequence. 2. `msg` holds the string identifier directly for fast attribute lookup. 3. `excs` holds the underlying exceptions wrapped inside a flat, immutable tuple.

21.0.2 19.2 The `except*` Clause and Exception Tree Filtering

To handle individual branches of an exception group, Python 3.11 introduced the `except*` statement. Unlike a traditional `except` statement, which catches a single exception instance, `except*` extracts a matching subset from the Exception Group hierarchy, letting unmatched exceptions propagate.

1. Compilation and Bytecode Mechanics

The compiler generates a specialized sequence of bytecodes for `except*` blocks to split the exception tree dynamically at runtime.

Consider the two patterns below:

```

1 # Standard try/except
2 try:
3     raise ValueError("Error")
4 except ValueError as e:
5     pass

```

Disassembly of standard `except`:

1	1	0	NOP	
2	2	SETUP_FINALLY		8 (to 12)
3				
4	2	4	LOAD_GLOBAL	1 (ValueError)
5		6	LOAD_CONST	1 ('Error')
6		8	CALL	1
7		10	RAISE_VARARGS	1
8				
9	3	>>	12 PUSH_EXC_INFO	
10			14 CHECK_EXCEPT	1 (ValueError)
11			16 POP_JUMP_IF_FALSE	14 (to 46)
12			18 STORE_FAST	0 (e)
13			...	

Now, consider the `except*` variant:

```

1 # Modern try/except*
2 try:
3     raise ExceptionGroup("Main Group", [ValueError("Error A"), TypeError("Error B")])
4 except* ValueError as eg:
5     print("Caught Val:", eg.exceptions)

```

Disassembly of `except*`:

1	1	0	NOP	
2	2	SETUP_FINALLY		12 (to 16)
3				
4	2	4	LOAD_GLOBAL	1 (ExceptionGroup)

```

5          6 LOAD_CONST          1 ('Main Group')
6          8 ... (building ValueError and TypeError objects)
7         10 BUILD_LIST          2
8         12 CALL                2
9         14 RAISE_VARARGS       1
10
11 3      >> 16 PUSH_EXC_INFO
12         18 CHECK_EXCEPT_STAR 1 (ValueError)
13         20 POP_JUMP_IF_FALSE 16 (to 52)
14         22 STORE_FAST         0 (eg)
15         ...

```

2. Bytecode Analysis and VM Stack Transitions

The crucial bytecode introduced for PEP 654 is `CHECK_EXCEPT_STAR`. Its execution flow proceeds as follows:

```

1 Stack State Before CHECK_EXCEPT_STAR:
2 +-----+
3 | ValueError (type tag) | <-- TOP OF STACK (TOS)
4 +-----+
5 | ExceptionGroup Object | <-- TOS1
6 +-----+
7
8 Execution details of CHECK_EXCEPT_STAR:
9 1. Pops the target exception type (TOS: ValueError).
10 2. Inspects the active ExceptionGroup (TOS1).
11 3. Executes a C-level filtering function:
12    - Splits the ExceptionGroup into a matched group containing all ValueError
13      instances.
14    - Pushes the matched group.
15    - Pushes the remaining unmatched exception group (containing TypeError).
16
17 Stack State After CHECK_EXCEPT_STAR (on match):
18 +-----+
19 | Matched ExceptionGroup | <-- TOS (Bound to 'eg' local variable)
20 +-----+
21 | Unmatched ExceptionGrp | <-- TOS1 (Propagated or checked by next handler)
22 +-----+

```

If the match group is empty (no `ValueError` instances were found), `CHECK_EXCEPT_STAR` pushes `None` to TOS, and `POP_JUMP_IF_FALSE` jumps directly to the next handler block.

21.0.3 19.3 Exception Tree Filtering Algorithm

The core mechanics of splitting exception groups are defined in CPython's exception runtime library. The algorithm must recursively inspect the tree of exceptions and cleanly separate them.

```

1          ExceptionGroup ("Main")
2          /                \
3      TypeError("B")      ExceptionGroup ("Sub")
4                          /                \
5          ValueError("A")      KeyError("C")

```

If we execute `except* ValueError`: 1. The runtime calls the internal C function `exception_group_filter` (`eg, match_value_error_func`). 2. It traverses the children of the group: `* TypeError("B")`: Does not

match. Added to the unmatched collection. * `ExceptionGroup("Sub")`: Recursively calls `exception_group_filter`. * Inside `ExceptionGroup("Sub")`: * `ValueError("A")`: Matches! Added to the sub-matched collection. * `KeyError("C")`: Does not match. Added to the sub-unmatched collection. * Since there were matches inside Sub, a new `ExceptionGroup("Sub")` is instantiated, containing only `ValueError("A")`. This is returned as the sub-matched tree. * A new `ExceptionGroup("Sub")` is instantiated containing only `KeyError("C")` and returned as the sub-unmatched tree. 3. The top-level execution collects the returned structures: * matched tree: `ExceptionGroup("Main", [ExceptionGroup("Sub", [ValueError("A")])])` * unmatched tree: `ExceptionGroup("Main", [TypeError("B"), ExceptionGroup("Sub", [KeyError("C")])])`

Here is the equivalent C-like pseudocode of the recursive filtering function:

```

1 PyObject* exception_group_filter(PyObject *eg, PyObject *match_type) {
2     PyObject *matched_list = PyList_New(0);
3     PyObject *unmatched_list = PyList_New(0);
4
5     PyObject *excs = ((PyExceptionGroupObject*)eg)->excs;
6     Py_ssize_t size = PyTuple_GET_SIZE(excs);
7
8     for (Py_ssize_t i = 0; i < size; i++) {
9         PyObject *exc = PyTuple_GET_ITEM(excs, i);
10        if (PyExceptionGroup_Check(exc)) {
11            // Recursive split
12            PyObject *sub_match = NULL, *sub_unmatch = NULL;
13            split_exception_group(exc, match_type, &sub_match, &sub_unmatch);
14            if (sub_match != NULL) {
15                PyList_Append(matched_list, sub_match);
16            }
17            if (sub_unmatch != NULL) {
18                PyList_Append(unmatched_list, sub_unmatch);
19            }
20        } else {
21            // Leaf node matching
22            if (PyErr_GivenExceptionMatches(exc, match_type)) {
23                PyList_Append(matched_list, exc);
24            } else {
25                PyList_Append(unmatched_list, exc);
26            }
27        }
28    }
29
30    PyObject *matched_group = NULL;
31    if (PyList_Size(matched_list) > 0) {
32        matched_group = create_exception_group_from_list(eg, matched_list);
33    }
34    PyObject *unmatched_group = NULL;
35    if (PyList_Size(unmatched_list) > 0) {
36        unmatched_group = create_exception_group_from_list(eg, unmatched_list);
37    }
38
39    return PyTuple_Pack(2, matched_group, unmatched_group);
40 }

```

21.0.4 19.4 Traceback Representation and `add_note()` Internals

1. Traceback Trees

Because exception groups represent hierarchical trees of errors, CPython's traceback generator is modified to format tracebacks recursively as tree diagrams.

When printed, the output represents the nesting level using visual indicators:

```

1 + Exception Group Traceback (most recent call last):
2 |   File "example.py", line 4, in <module>
3 |       raise ExceptionGroup("Main Group", [ValueError("Error A"), TypeError("
      Error B")])
4 | ExceptionGroup: Main Group (2 sub-exceptions)
5 +-+----- 1 -----
6 | ValueError: Error A
7 | +----- 2 -----
8 | TypeError: Error B
9 | +-----

```

If there are nested exception groups, the branch indicators prefix the output recursively (e.g., +-+----- 1.1 -----).

2. PEP 678 Exception Notes and C-Level Internals

PEP 678 introduces `BaseException.add_note(note)`, enabling users to attach custom text to exceptions without modifying their instantiation arguments. This is highly valuable for diagnostic tools, test frameworks (like `pytest`), and tracing asynchronous executions.

When called, `add_note()` stores the string inside a `__notes__` list attribute on the exception object. The C-level implementation handles memory allocations and type checks directly:

```

1 /* Objects/exceptions.c */
2 static PyObject *
3 BaseException_add_note(PyBaseExceptionObject *self, PyObject *note)
4 {
5     if (!PyUnicode_Check(note)) {
6         PyErr_SetString(PyExc_TypeError, "note must be a string");
7         return NULL;
8     }
9
10    PyObject *dict = self->dict;
11    if (dict == NULL) {
12        dict = PyDict_New();
13        if (dict == NULL) {
14            return NULL;
15        }
16        self->dict = dict;
17    }
18
19    PyObject *notes = PyDict_GetItemWithError(dict, &_amp;Py_ID(__notes__));
20    if (notes == NULL) {
21        if (PyErr_Occurred()) {
22            return NULL;
23        }
24        notes = PyList_New(0);
25        if (notes == NULL) {
26            return NULL;
27        }
28        if (PyDict_SetItem(dict, &_amp;Py_ID(__notes__), notes) < 0) {
29            Py_DECREF(notes);
30            return NULL;
31        }
32        Py_DECREF(notes);
33    }
34
35    if (PyList_Append(notes, note) < 0) {
36        return NULL;
37    }
38

```

```
39     Py_RETURN_NONE;  
40 }
```

At runtime, the traceback printing logic calls `PyObject_GetAttr(exc, &_Py_ID(__notes__))`. If the list is found, it iterates over each note string and prints it directly after the exception traceback and type representation.

Part VII

Volume VII: The GIL-less Future & JIT Compilers

Chapter 22

Python 3.13: Free-Threaded Build & GIL Removal Internals

22.0.1 20.1 The Free-Threaded CPython Paradigm Shift

Introduced provisionally in Python 3.13 (PEP 703), the free-threaded build of CPython (also known as the “no-GIL” build) allows executing bytecode in parallel across multiple OS threads without the constraint of the Global Interpreter Lock (GIL).

In standard CPython, the GIL serves as a single global mutex protecting the interpreter state and all Python objects from concurrent access. This guarantees that only one thread executes bytecode at a time, simplifying the VM design and preventing memory races. However, it severely limits Python’s ability to scale on multi-core processors.

To remove the GIL safely, the CPython runtime has been fundamentally re-engineered. Moving from single-threaded assumptions to a scalable parallel execution engine required re-designing three pillars of the runtime: 1. **Reference Counting**: Overcoming atomic bus contention on object refcount updates. 2. **Memory Allocation**: Ensuring lock-free, thread-local allocations for object instantiation. 3. **Collection Mutability**: Protecting built-in data structures (dictionaries, lists, sets) from race conditions.

22.0.2 20.2 Biased Reference Counting (BRC) and Object Headers

Reference counting is the primary mechanism for memory management in CPython. In a standard build, incrementing/decrementing reference counts is a simple, non-atomic operation: `op->ob_refcnt++`. In a thread-parallel environment, standard operations would cause race conditions. However, using standard atomic operations (`atomic_add` / `atomic_sub` or `LOCK XADD` instructions) globally would destroy performance due to cache coherency traffic (bus locks) across CPU cores.

To solve this, PEP 703 introduces **Biased Reference Counting (BRC)**. BRC divides an object’s reference count state based on thread ownership.

1. Biased Object Headers

An object is “biased” toward the thread that allocated it (the owner thread). The owner thread uses fast, non-atomic CPU instructions to modify the reference count. Non-owner threads (foreign threads) must use atomic operations on a separate “shared” reference count field.

Under no-GIL configurations, the standard `PyObject` header (`Include/object.h`) is redefined:

```
1 typedef struct _object {
2     uintptr_t ob_ref_local;          /* Local reference count and owner thread ID */
3     uintptr_t ob_ref_shared;        /* Shared reference count and object status
    flags */
```

```
4 PyObject *ob_type;          /* Pointer to object type descriptor */
5 } PyObject;
```

2. Bit Layout and Fields representation

The fields `ob_ref_local` and `ob_ref_shared` encode multiple pieces of metadata to optimize memory usage:

- `ob_ref_local`:
 - **Thread ID Bits (Upper Bits)**: Holds the memory address of the owner thread's thread-state structure (`PyThreadState*`).
 - **Local Refcount (Lower Bits)**: A small counter tracking reference updates made by the owner thread. Typically, the lower 16 or 32 bits are used for this count.
- `ob_ref_shared`:
 - **Shared Refcount (Upper Bits)**: A signed counter tracking reference additions/subtractions made by non-owner threads.
 - **Status Flags (Lower Bits)**: Bits reserved for encoding object states:
 - * `STATE_STATIC` (bit 0): Set if the object is static (e.g., statically allocated builtin types).
 - * `STATE_IMMORTAL` (bit 1): Set if the object is immortal (reference counts are never modified).
 - * `STATE_DEFERRED` (bit 2): Set if the object uses deferred reference counting (GC-managed objects).

3. BRC Reference Counting Algorithm

When modifying a reference (via `Py_INCREF` or `Py_DECREF`), the runtime performs owner-thread detection:

Owner Thread State Pointer = `ob_ref_local & THREAD_STATE_MASK`

```
1      [Reference Count Change Request]
2      |
3      Get Current ThreadState (Tstate)
4      |
5      Is Tstate == (ob_ref_local & MASK)?
6      /      \
7      Yes      No
8      /      \
9  [Non-Atomic Local Increment]  [Atomic Shared Increment]
10 (Fast path: CPU registers)    (Slow path: CAS / Bus Lock)
```

Increment Code Walkthrough

```
1 static inline void
2 _Py_INCREF_Specialized(PyObject *op, PyThreadState *tstate)
3 {
4     uintptr_t local = op->ob_ref_local;
5     // Fast path: current thread is the owner
6     if ((local & _Py_THREAD_ID_MASK) == (uintptr_t)tstate) {
7         op->ob_ref_local = local + _Py_REF_INCREMENT;
8     }
9     // Slow path: foreign thread
10    else {
```



```

11     _Py_Incref_Shared(op);
12 }
13 }

```

Decrement and Deallocation Walkthrough When a reference is decremented, if the local count reaches zero, the owner thread merges the shared reference count:

```

1 static inline void
2 _Py_DECREF_Specialized(PyObject *op, PyThreadState *tstate)
3 {
4     uintptr_t local = op->ob_ref_local;
5     if ((local & _Py_THREAD_ID_MASK) == (uintptr_t)tstate) {
6         uintptr_t new_local = local - _Py_REF_INCREMENT;
7         if ((new_local & _Py_REF_MASK) == 0) {
8             // Local count hit zero; merge shared reference counts to check for
9             deallocation
10             _Py_Merge_And_Dealloc(op);
11         } else {
12             op->ob_ref_local = new_local;
13         }
14     } else {
15         _Py-Decref_Shared(op);
16     }
17 }

```


During `_Py_Merge_And_Dealloc(op)`, the owner thread atomically reads and clears `ob_ref_shared`. If the sum of `ob_ref_local` (cleared of thread state bits) and `ob_ref_shared` is less than or equal to zero, it calls the type's `tp_dealloc` slot to free the object.

22.0.3 20.3 Thread-Safe Allocation via mimalloc

CPython's classic allocator (`pymalloc`) is optimized for small objects but relies on single-threaded assumptions. Under a free-threaded runtime, sharing a global allocator lock would lead to lock contention. Consequently, PEP 703 replaces `pymalloc` with **mimalloc**, a thread-safe, metadata-compact allocator developed by Microsoft.

1. Mimalloc Structural Hierarchy

Mimalloc organizes memory allocations into hierarchical structures to minimize thread synchronization:



- **Segments (typically 4MB):** Large contiguous memory blocks requested from the operating system.

- **Pages (typically 64KB - 512KB):** Subdivisions of segments containing blocks of a single size class (e.g., 32-byte blocks, 64-byte blocks).
- **Heaps:** Thread-local handles containing lists of active pages.

2. Thread-Local Allocations

Every OS thread maintains its own thread-local mimalloc heap (`mi_heap_t`). When a thread executes `PyObject_New()`, it attempts to allocate from its thread-local heap: 1. It locates the active page corresponding to the requested size class. 2. It pops a block from the page's thread-local **free list**. This list is accessed only by the owning thread, requiring no atomic locks or memory barriers.

3. Cross-Thread Deallocations and Thread-Safe Free Lists

When a foreign thread deallocates an object, writing directly back to the owner thread's thread-local free list would cause race conditions. To resolve this, mimalloc uses a secondary, atomic free list for each page:

```
1 [Foreign Thread deallocates block]
2 |
3 | v (Atomic CAS operation)
4 +-----+
5 | Atomic Free List | (Attached to Page)
6 +-----+
7 |
8 | v (During thread-local heap maintenance)
9 +-----+
10 | Local Free List | (Accessed non-atomically by Owner Thread)
11 +-----+
```

When a foreign thread frees memory: 1. It atomically prepends the freed block to the page's **atomic free list** (`thread_free`) using a lock-free Compare-And-Swap (CAS) loop. 2. When the owner thread exhausts its thread-local **free list**, it cleans up and merges the `thread_free` list into its local list. This deferred merging drastically reduces cache line bouncing across CPU cores.

22.0.4 20.4 Lock-free Collections & Thread Safety

In standard CPython, code executing modifications to dictionaries, lists, or sets does not need to worry about internal consistency, as the GIL serializes all operations. Without the GIL, concurrent reads/writes could corrupt collection structures, leading to segmentation faults or memory corruption.

1. Dictionary Locking System

CPython's compact dictionary layout was modified to support concurrent read and write operations. * **Read Paths (Lock-free):** Operations like dict lookups (`PyDict_GetItem`) read the hash index and table entries without acquiring locks. They rely on atomic pointer reads and memory barriers to ensure they read consistent states. * **Write Paths (Fine-grained Locks):** Modifying operations (insertion, deletion) acquire a localized lock associated with the dictionary instance. Instead of locking the global runtime, CPython locks only the specific dictionary being modified.

To avoid allocating a full OS mutex for every dictionary (which would consume excessive memory), CPython uses a pool of shared locks, indexing into them using a hash of the dictionary's memory address.

2. Lock Arrays and PyMutex

To support light, low-overhead synchronization, CPython implements a highly efficient locking primitive: **PyMutex**.

```

1 typedef struct {
2     uint8_t v; /* Lock state byte */
3 } PyMutex;

```

- **State Byte v:**
 - Bit 0: Indicates if the lock is held (1) or free (0).
 - Bit 1: Indicates if there are threads waiting in the queue (1) or not (0).
- **Fast Path:** A thread attempts to acquire the lock using a single atomic test-and-set instruction (`atomic_compare_exchange` on bit 0). If successful, it proceeds without sleeping or system calls.
- **Slow Path:** If the lock is held, the thread registers itself in a global wait queue associated with the lock's memory address and suspends execution, waiting to be woken up by the lock holder.

22.0.5 20.5 Generational Garbage Collection (GC) in a GIL-less World

CPython's cyclic garbage collector identifies self-referencing loops that reference counting alone cannot reclaim. In a standard build, the GC runs synchronously on the main thread during allocations. In a free-threaded environment, running the GC requires coordinated synchronization across all running threads to prevent mutator threads from modifying object references during the collection pass.

1. Stop-the-World (STW) Synchronization

To run the cyclic collection safely, CPython implements a Stop-the-World mechanism:

```

1 [GC Thread initiates GC] ---> Sets global GC status flag
2                               |
3   [All other threads reach safe-points / bytecode loop boundary]
4                               |
5                               Suspends all mutator threads
6                               |
7                               Executes GC Collection Pass
8                               |
9                               Resumes all mutator threads

```

1. A thread triggering a collection sets a global GC request flag.
2. All other running threads check this flag at defined bytecode boundaries (safe-points).
3. Upon detecting the flag, threads suspend their execution and block on an OS condition variable.
4. Once all threads are verified as stopped, the GC thread safely performs the collection pass, scanning object references and breaking cyclic loops.
5. After completion, the GC thread signals the condition variable, resuming all suspended threads.

2. Deferred Reference Counting (DRC)

For objects created during compilation (such as functions, classes, code objects, and constants), standard reference counting is completely bypassed. These objects are marked as `STATE_DEFERRED` or `STATE_IMMORTAL` inside `ob_ref_shared`. * **Immortal Objects:** Reference count updates are completely skipped. They are never deallocated during the lifetime of the process. * **Deferred Objects:** Reference count updates are ignored at runtime. The GC takes full responsibility for tracking these objects and determining when they are no longer reachable, freeing their memory during GC cycles. This bypasses atomic reference operations on highly accessed structural components.

Chapter 23

Python 3.13: Copy-and-Patch JIT Compiler Architecture

23.0.1 21.1 The Copy-and-Patch JIT Paradigm (PEP 744)

Python 3.13 introduces an experimental **Copy-and-Patch Just-In-Time (JIT) compiler** (PEP 744). Unlike classic template JITs (which compile bytecode to slow sequences of function calls) or heavy optimizing tracing/method JITs like PyPy or V8 (which compile code via complex graph building and register allocation at runtime), CPython’s copy-and-patch JIT achieves fast compilation times with near-zero runtime compilation overhead.

1. Optimization and Compilation Tradeoffs

Traditional JIT compilers (e.g., LLVM-based JITs) generate highly optimized assembly code at runtime, but the compilation process itself requires substantial CPU cycles and memory. This makes them unsuitable for interactive runtimes or workloads with short execution times, where the compilation cost outweighs the speedup.

Copy-and-patch JIT solves this by splitting the compilation workload: * **Compile-Time (Static Build Phase)**: High-performance compilers (Clang/LLVM) compile small C functions (micro-operations) into native object files (ELF, Mach-O, or COFF). These object files are parsed to extract the raw machine code bytes (stencils) and metadata describing the location of symbols and offsets (holes). * **Runtime (JIT Compilation Phase)**: The runtime JIT engine traverses a hot execution trace, copies the stencils into a memory buffer using simple memory copies (`memcpy`), and fills the holes (patches the relocations) with runtime values like frame offsets, constants, or jump destinations.

This architecture delivers machine-native code execution speeds with a compiler that executes in microseconds.

23.0.2 21.2 Stencil Relocation and C-Level Internals


1. Machine Code Stencils and Relocations

At the C level, the JIT engine represents bytecode instructions as stencils. A stencil is a sequence of native instructions containing placeholders (holes) that must be resolved with runtime context:

```
1 /* Python/jit.c (conceptual layout) */
2 typedef struct {
3     uint32_t offset;           /* Offset in bytes from start of stencil where
    hole begins */
```

```
4      uint8_t type;                /* Relocation type (e.g., absolute 64-bit, PC-
   relative 32-bit) */
5      uintptr_t value;            /* Raw value or identifier of symbol to patch into
   hole */
6  } JITRelocation;
7
8  typedef struct {
9      const unsigned char *code;    /* Array of raw machine code bytes */
10     size_t code_size;              /* Size of machine code array */
11     const JITRelocation *relocs;   /* Pointer to array of relocation entries */
12     size_t reloc_count;            /* Count of relocations for this stencil */
13 } JITStencil;
```

2. Visual Representation of Copy-and-Patch Stitching




During the build phase, CPython uses Clang to compile each bytecode case. A Python script (Tools/cases_generator/jit.py) parses the resulting object files, extracts the machine code, and outputs a C header file (jit_cases.c.h) containing arrays of static stencils.

23.0.3 21.3 Tier 2 Micro-Ops (uops) and the Optimization Pipeline

CPython 3.13 splits execution into a multi-tiered pipeline: 1. **Tier 1 (Specializing Interpreter)**: The standard interpreter loop (from Chapter 17) executes bytecode. If a block of bytecode is executed frequently, it triggers dynamic tracing. 2. **Tier 2 (Tracer / Optimizer)**: A tracing engine compiles the bytecode block into a flat linear sequence of simplified instructions called **Micro-Ops (uops)**. 3. **Tier 2 Optimization**: The uop trace undergoes basic optimizations: * **Dead Code Elimination**: Removing operations whose outputs are unused. * **Type Propagation**: Propagating specialized types through the trace to eliminate redundant type-checks. * **Abstract Evaluation**: Simulating stack heights to eliminate redundant push and pop operations. 4. **Tier 2 JIT Execution**: The optimized uop trace is compiled by the Copy-and-Patch JIT engine, producing native machine code.

1. Bytecode to Native Compilation Flow



8	Machine Code	(Stencils + Holes)	(uop linear trace)
9	+-----+	+-----+	+-----+

2. Tier 2 uop Compilation Loop

The JIT compiler processes optimized uops using a loop that matches uop IDs to static stencils:

```

1  /* Python/jit.c */
2  void* _PyJIT_CompileTrace(PyExecutorObject *executor, _PyUOpInstruction *trace,
   size_t trace_len) {
3      // 1. Allocate writable memory page
4      unsigned char *code_buffer = allocate_jit_memory();
5      unsigned char *cursor = code_buffer;
6
7      for (size_t i = 0; i < trace_len; i++) {
8          _PyUOpInstruction uop = trace[i];
9          // Retrieve static stencil generated at build time
10         JITStencil stencil = get_stencil_for_uop(uop.opcode);
11
12         // Copy machine code template
13         memcpy(cursor, stencil.code, stencil.code_size);
14
15         // Patch relocations inside the copied code
16         for (size_t j = 0; j < stencil.reloc_count; j++) {
17             JITRelocation reloc = stencil.relocs[j];
18             uintptr_t patch_value = resolve_uop_symbol(uop, reloc.value);
19             patch_hole(cursor + reloc.offset, reloc.type, patch_value);
20         }
21
22         cursor += stencil.code_size;
23     }
24
25     // 2. Transition memory permissions to Read/Execute (RX)
26     make_jit_memory_executable(code_buffer, cursor - code_buffer);
27     return code_buffer;
28 }

```

23.0.4 21.4 Relocation Hole Patching and Memory Security

Patching relocation holes requires writing machine-specific offsets directly into the copied stencils. The JIT engine implements several relocation type handlers matching the host architecture:

1. Relocation Implementations

- **x86_64 Relocations:**

- R_X86_64_64: Direct absolute 64-bit address patch.
- R_X86_64_PC32: 32-bit PC-relative address patch (used for relative jumps between stencils).

- **AArch64 (ARM64) Relocations:**

- R_AARCH64_MOVW_UABS_G0 & R_AARCH64_MOVW_UABS_G1: Patches lower/upper bits of immediate load operations (movz/movk).
- R_AARCH64_CONDBR19: Patches conditional branch instruction target offsets.

2. Memory Security: W^X Enforcement

Modern operating systems enforce **Write-Once-Read-Execute (W^X)** permissions to prevent security vulnerabilities. A memory page cannot be simultaneously writable and executable.

To adhere to this, CPython manages execution memory in two steps: 1. **Allocation Phase:** The JIT allocator requests pages from the OS with Read-Write (RW) permissions (using `mmap` with `PROT_READ | PROT_WRITE` on POSIX, or `VirtualAlloc` with `PAGE_READWRITE` on Windows). 2. **Compilation & Patching:** The compiler copies stencils and writes relocation values into the RW pages. 3. **Permissions Transition:** Once patching is complete, the JIT locks down the memory, requesting the OS to transition permissions to Read-Execute (RX) (using `mprotect` with `PROT_READ | PROT_EXEC`, or `VirtualProtect` with `PAGE_EXECUTE_READ`). 4. **Execution:** The executor executes the native machine code block by calling it via a C function pointer.

Chapter 24


Python 3.12 to 3.13: Subinterpreters & Per-Interpreter GIL Parallelism

24.0.1 22.1 Isolation of Interpreter State (PEP 684)

Before Python 3.12, multi-threaded concurrency in CPython was constrained by a single process-wide Global Interpreter Lock (GIL). Although subinterpreters could be spawned via the C-API, they all shared this single GIL and process-wide global structures, preventing them from running concurrently across multiple CPU cores.

PEP 684 introduced support for **isolated subinterpreters**, each running with its own per-interpreter GIL. This configuration allows a single OS process to run multiple Python interpreters concurrently, achieving true multi-core parallel execution.

1. Process Memory Layout: Shared vs. Isolated States



2. The PyInterpreterState Structure and Refactoring

To achieve GIL isolation, CPython moved all state variables out of C-static/global variables into fields nested inside the `PyInterpreterState` structure. This structure is defined in CPython's internal headers (`Include/internal/pycore_interp.h`):

```

1 struct _is {
2     struct _is *next;                /* Linkage for global list of
3     struct _ceval_state ceval;        /* Per-interpreter evaluation loop
4     configuration and GIL */
5     struct _gc_state gc;              /* Per-interpreter garbage collector
6     generations */
7
8     /* Private Heaps and Object Storage */
9     PyObject *modules;                /* Isolated dictionary containing loaded
10    modules (sys.modules) */
11    PyObject *sysdict;                 /* Isolated sys module context variables
12    */
13    PyObject *builtins;                /* Builtins dictionary context */
14
15    /* Execution Contexts */
16    struct _dict_state dict_state;     /* Per-interpreter dictionary structures
17    and caching */
18    struct _types_state types;         /* Per-interpreter type caches and static
19    type tables */
20
21    int64_t id;                        /* Unique identifier for the
22    subinterpreter */
23    int gil_status;                    /* Configuration representing if this
24    interpreter owns a GIL */
25 };
26 typedef struct _is PyInterpreterState;

```

- `struct _ceval_state ceval`: Wraps the evaluation loop structures, including the interpreter's private Global Interpreter Lock (`ceval.gil`). This lock is completely isolated from other interpreters.
- `struct _gc_state gc`: Isolates the generational lists and flags of the cyclic Garbage Collector. A collection run in one subinterpreter does not block or stop other running interpreters.
- `PyObject *modules`: Isolates module namespaces. When subinterpreter 1 executes `import sys`, it resolves to a different dictionary instance than the `sys` module in subinterpreter 2.

24.0.2 22.2 The C-API Subinterpreter Initialization Configuration

From a C extension, subinterpreters can be configured and created using the Python C-API. In Python 3.12+, the runtime exposes structure configurations that specify whether the newly created interpreter should share the GIL or initialize its own:

```

1 /* C-API Example: Creating an Isolated Subinterpreter with its own GIL */
2 #include <Python.h>
3
4 void execute_in_subinterpreter() {
5     // 1. Define configuration for isolated subinterpreter
6     PyInterpreterConfig config = {
7         .use_main_obmalloc = 0,        /* Use separate mimalloc heap allocations
8         */
9         .allow_fork = 0,                /* Disallow fork operations for safety */
10        .allow_exec = 0,                 /* Disallow exec operations */
11        .allow_threads = 1,              /* Enable threading within subinterpreter
12        */
13        .allow_daemon_threads = 0,
14    };
15 }

```

```


12     .check_multi_interp_extensions = 1, /* Enforce strict multi-interpreter
    extension isolation */
13     .gil = PyInterpreterConfig_OWN_GIL /* REQUEST PRIVATE PER-INTERPRETER
    GIL */
14 };
15
16 PyThreadState *sub_tstate = NULL;
17 PyStatus status = Py_NewInterpreterFromConfig(&sub_tstate, &config);
18
19 if (PyStatus_Exception(status)) {
20     fprintf(stderr, "Failed to initialize subinterpreter.\n");
21     return;
22 }
23
24 // The current OS thread now holds the GIL for the new subinterpreter.
25 // We can execute code inside this isolated context:
26 PyRun_SimpleString("import sys; print('Subinterpreter ID:', sys.getsizeof(
    sys.modules))");
27
28 // 2. Shut down the subinterpreter and release the private GIL
29 Py_EndInterpreter(sub_tstate);
30 }

```

24.0.3 22.3 Cross-Interpreter Communication Channels

Because subinterpreters run on separate heaps with private reference counting states, passing raw object pointers (`PyObject*`) between them is strictly forbidden. Sharing pointers would lead to race conditions when multiple threads increment or decrement reference counts concurrently.

To pass data, the runtime uses a **shared-nothing** message passing architecture, serializing objects across isolation boundaries.



1. **Serialization:** The sender interpreter extracts the data. It uses serialization protocols (such as `pickle` or a specialized C-level marshal mechanism) to copy the object's value into a raw, C-native memory buffer.
2. **Transmission:** The buffer is queued in a thread-safe global memory channel that belongs to the process-global space (external to any specific interpreter heap).
3. **Deserialization:** The receiving interpreter pulls the buffer, decodes the values, and allocates new objects representing the data on its own heap.

Parallel Computation Code Example

In Python 3.12+, you can interact with subinterpreters using the experimental `_xxsubinterpreters` module. Below is a complete script demonstrating true multi-core parallel computation by distributing work across subinterpreters:

```

1 import threading
2 import time
3 import _xxsubinterpreters as interpreters
4
5 # Worker script to execute inside the subinterpreter
6 worker_script = """
7 import time
8 import _xxsubinterpreters as interpreters
9
10 # Get the channel to send results back
11 channel_id = interpreters.get_current_channel()
12 # Perform some CPU-heavy computation
13 result = sum(i * i for i in range(10_000_000))
14
15 # Send the calculation result back to the main interpreter
16 interpreters.channel_send(channel_id, result)
17 """
18
19 def run_worker(interp_id, channel_id):
20     # Bind the current thread to the subinterpreter and run the computation
21     interpreters.run_string(interp_id, worker_script, shared={'channel_id':
22         channel_id})
23
24 def main():
25     # 1. Create a channel for communication
26     channel_id = interpreters.channel_create()
27
28     # 2. Create an isolated subinterpreter (with its own GIL by default in
29     # 3.12+)
30     interp_id = interpreters.create()
31
32     # 3. Spawn a separate OS thread to execute the subinterpreter concurrently
33     thread = threading.Thread(target=run_worker, args=(interp_id, channel_id))
34
35     print("Spawning worker thread...")
36     start_time = time.perf_counter()
37     thread.start()
38
39     # The main interpreter runs in parallel on the primary thread
40     main_result = sum(i * i for i in range(5_000_000))
41     print("Main interpreter calculation complete:", main_result)
42
43     # 4. Wait for the subinterpreter to send its result and join the thread
44     sub_result = interpreters.channel_recv(channel_id)
45     thread.join()
46
47     duration = time.perf_counter() - start_time
48     print(f"Subinterpreter calculation complete: {sub_result}")
49     print(f"Total parallel execution duration: {duration:.4f} seconds")
50
51     # 5. Destroy the subinterpreter and clean up channels
52     interpreters.destroy(interp_id)
53     interpreters.channel_destroy(channel_id)
54
55 if __name__ == '__main__':
56     main()

```


Part VIII

Volume VIII: Runtime Internals & C Extensions

Chapter 25

CPython Memory Allocator (PyMalloc) & Generational Garbage Collection

25.0.1 23.1 CPython's Memory Allocation Engine (PyMalloc)

For large allocations (greater than 512 bytes), CPython forwards the request directly to the system's standard C library allocator (`malloc()`). However, for small objects (≤ 512 bytes) which represent the vast majority of Python allocations, standard operating system allocators introduce high fragmentation and locking overhead. To resolve this, CPython implements a custom small-object allocator called **PyMalloc**.

1. PyMalloc Memory Hierarchy

PyMalloc structures memory into three distinct layers to minimize operating system allocation calls:

- * **Arenas (256 KB)**: Contiguous memory blocks allocated from the operating system, aligned to 256 KB boundaries. Arenas manage memory at the virtual memory level and contain exactly 64 pools.
- * **Pools (4 KB)**: Subdivisions of arenas that match the operating system's virtual page size. Each pool is dedicated to a single **size class** (e.g., all blocks in a pool are 32 bytes, or all are 64 bytes).
- * **Blocks**: The actual chunks of memory returned to the interpreter. Blocks range from 8 bytes to 512 bytes, aligned to 8-byte steps (giving 64 distinct size classes).

2. C-level Struct Definitions in `obmalloc.c`

The structures for pools and arenas are defined in CPython's memory management source file (`Objects/obmalloc.c`).

Pool Header Struct Every 4 KB pool begins with a header that tracks block allocations and linked lists of free pools:

```
1  /* Objects/obmalloc.c */
2  typedef uint8_t block;
3
4  struct pool_header {
5      union {
6          block *as_block;
7          uint32_t as_uint;
8      } nextfree;          /* Pointer to the next available free block in the
9                          pool */
10     union {
11         block *as_block;
12         uint32_t as_uint;
```

```

12     } firstfree;           /* Pointer to the first free block in the pool */
13     struct pool_header *nextpool; /* Link to the next pool in the active list
14     */
15     struct pool_header *prevpool; /* Link to the previous pool in the active
16     list */
17     uint32_t refcount;       /* Number of allocated blocks currently in use */
18     uint32_t szidx;         /* Size class index of this pool */
19     uint32_t freeblocks;    /* Number of free blocks remaining in the pool */
20 };
21 typedef struct pool_header *poolp;

```

- **nextfree**: Points to the next block in a singly linked list of free blocks inside the pool. PyMalloc uses **in-place pointers** inside the free blocks themselves to implement this list without consuming extra memory.
- **firstfree**: Tracks the boundary of allocated blocks versus unallocated space in the pool.
- **szidx**: Specifies which size class this pool belongs to.

Arena Object Struct CPython tracks arenas using an array of `arena_object` descriptors:

```

1 /* Objects/obmalloc.c */
2 struct arena_object {
3     uintptr_t address;           /* Base virtual memory address of the 256 KB
4     block */
5     block* pool_address;        /* Pointer to the first pool inside this arena
6     */
7     uint32_t nfreepools;        /* Number of pools currently free in this arena
8     */
9     uint32_t ntotalpools;       /* Total number of pools inside this arena (
10     usually 64) */
11     struct arena_object* nextarena; /* Link to the next arena */
12     struct arena_object* prevarena; /* Link to the previous arena */
13 };

```

3. Size Class Formula

The size class for an allocation request of size S is calculated by:

$$\text{Class Index} = \left\lceil \frac{S}{8} \right\rceil - 1$$

This mapped alignment restricts memory fragmentation, ensuring that small object requests are answered in $\mathcal{O}(1)$ time by locating the active pool array matching the calculated class index.

25.0.2 23.2 Generational Garbage Collection and Cycle Detection

While reference counting manages the vast majority of object lifecycles, it cannot identify self-referencing cyclic loops (e.g., `x = []`; `x.append(x)`). To reclaim cyclic memory leaks, CPython runs a cyclic Garbage Collector (GC) as a background system.

1. GC Memory Prefix and Header Layout

Every object tracked by the GC has an extra header prefix in memory before its normal `PyObject` structure. When `PyObject_GC_New()` is called, it allocates memory block of size:

$$\text{Allocation Size} = \text{sizeof(PyGC_Head)} + \text{sizeof(PyObject)}$$


```

1 +-----+
2 |                PyGC_Head                |
3 | - gc_next: Linkage to other GC tracked objects |
4 | - gc_prev: Linkage to other GC tracked objects |
5 | - gc_refs: Temporary reference state          |
6 +-----+
7 |                PyObject                | <--- Object Pointer
8 | Returned to User                        |
9 | - ob_refcnt: Standard reference counter      |
10 | - ob_type: Type pointer                    |
11 +-----+

```

The pointer returned to the user points directly to the PyObject start. The VM accesses the GC header by subtracting the header size:

$$\text{PyGC_Head Pointer} = (\text{PyGC_Head*})\text{op} - 1$$

The C-level layout of PyGC_Head is defined in Include/internal/pycore_gc.h:

```

1 typedef union _gc_head {
2     struct {
3         uintptr_t gc_next; /* Pointer to next object in generation list */
4         uintptr_t gc_prev; /* Pointer to previous object in generation list */
5         /*
6          * Py_ssize_t gc_refs; /* Reference count copy used during cycle checks */
7         */
8         } gc;
9         double dummy; /* Forces double-word alignment adjustments */
10 } PyGC_Head;

```

2. The Three Generations

The GC divides tracked objects into three generations based on survival age: * **Generation 0 (Gen 0)**: Where new objects are registered. Collected frequently. * **Generation 1 (Gen 1)**: Objects that survived a Gen 0 collection pass. * **Generation 2 (Gen 2)**: Long-lived objects. Collected least frequently.

Collection is triggered when the number of allocations minus deallocations in a generation exceeds a configured threshold:

$$\text{Allocations} - \text{Deallocations} > \text{Threshold}$$

Default thresholds are (700, 10, 10). You can view or change them via `gc.get_threshold()` and `gc.set_threshold()`.

25.0.3 23.3 The Cycle Detection Algorithm

To find and break reference cycles, the GC executes the following steps during a collection pass:

1. Initializing gc_refs

The GC copy-assigns the reference count of every object in the collection generation to its `gc_refs` field:

$$\text{gc_refs} = \text{ob_refcnt}$$

2. Traversing Internal References


For each object in the generation list, the GC calls its `tp_traverse` slot (a C function that lists all objects referenced by the object). If a referenced object is also in the collection generation, the GC decrements that object's `gc_refs`:

$$\text{gc_refs} = \text{gc_refs} - 1$$

After traversing all objects, any object that still has `gc_refs > 0` must be reachable from outside the generation list (e.g., from local variables, global parameters, or a higher generation).

3. Resolving and Splitting Collections

The GC splits the candidate list into two sets: `reachable` and `unreachable`.



1. If an object has `gc_refs > 0`, it is moved back to the `reachable` list.
2. All objects reachable *from* that object (transitively traversed via `tp_traverse`) are also marked as `reachable` and moved out of the `unreachable` list, even if their `gc_refs` was zero.
3. Any remaining objects in the `unreachable` set are confirmed to be part of a cyclic garbage loop and are marked for deallocation.

4. Clearing Cycles and Deallocation

The GC breaks the reference cycle by calling the `tp_clear` slot on each object in the `unreachable` set. This sets internal pointers (like dictionary entries or list items) to `None` or clears them, which decrements the reference counts of the target objects and allows the standard reference counter to deallocate the memory.

Chapter 26

C Extensions & Python C-API Interoperability

26.0.1 24.1 Writing a Pure C Extension Module

Python allows writing modules directly in C or C++ to access low-level OS APIs, optimize performance-critical paths, or link with native libraries. A C extension is compiled into a shared library (a `.so` file on Unix, or a `.pyd` file on Windows) that Python can load dynamically at runtime using `import`.

1. Detailed C-Extension with Error Handling and Keywords

A robust C extension must handle keyword arguments, perform error checks, raise Python exceptions on failure, and manage resources.

Below is an implementation of a custom C module exposing a safe division function:

```
1 #define PY_SSIZE_T_CLEAN
2 #include <Python.h>
3
4 /* 1. Core C Function Implementation with Keywords & Exception Handling */
5 static PyObject*
6 custom_safe_divide(PyObject* self, PyObject* args, PyObject* keywds) {
7     double numerator;
8     double denominator;
9
10    /* Argument keyword names array */
11    static char *kwlist[] = {"numerator", "denominator", NULL};
12
13    /* Parse Python argument tuple & dict into native double values */
14    if (!PyArg_ParseTupleAndKeywords(args, keywds, "dd", kwlist, &numerator, &denominator)) {
15        return NULL; /* Returns NULL to signal an exception occurred during parsing */
16    }
17
18    /* Check for division by zero and raise a Python ZeroDivisionError */
19    if (denominator == 0.0) {
20        PyErr_SetString(PyExc_ZeroDivisionError, "Denominator cannot be zero.");
21        return NULL; /* Return NULL to propagate the raised exception */
22    }
23
24    double result = numerator / denominator;
25
26    /* Convert native C double back to a Python PyFloatObject */
27    return PyFloat_FromDouble(result);
```

```

28 }
29
30 /* 2. Methods Table Registration */
31 static PyMethodDef CustomMethods[] = {
32     {
33         "safe_divide",
34         (PyCFunction)(void (*)(void))custom_safe_divide,
35         METH_VARARGS | METH_KEYWORDS,
36         "Divides numerator by denominator, returning float safely."
37     },
38     {NULL, NULL, 0, NULL} /* Sentinel marker to end array */
39 };
40
41 /* 3. Module Definition Structure */
42 static struct PyModuleDef custommodule = {
43     PyModuleDef_HEAD_INIT,
44     "custom_c_math", /* Module Name */
45     "A custom high-performance C extension module.", /* Docstring */
46     -1, /* Module state size (-1 = global state) */
47     CustomMethods
48 };
49
50 /* 4. Initialization Hook called upon 'import custom_c_math' */
51 PyMODINIT_FUNC
52 PyInit_custom_c_math(void) {
53     return PyModule_Create(&custommodule);
54 }

```

- **METH_VARARGS | METH_KEYWORDS:** Tells the interpreter that this function accepts both positional arguments (parsed into `args` as a tuple) and keyword arguments (parsed into `kwargs` as a dictionary).
- **PyArg_ParseTupleAndKeywords:** Resolves both types of inputs, mapping them into the C double variables based on the format string `"dd"` (double, double) and the keyword list `kwlist`.
- **Exception Signaling:** Returning a `NULL` pointer tells CPython's execution frame that an exception has been set inside the thread state. The interpreter pauses bytecode execution and executes its exception handler path.

26.0.2 24.2 Reference Counting and Ownership Rules in C-API

When interacting with the Python C-API, developers must manually manage reference counting. Standard Python code delegates reference updates to the compiler/interpreter, but in C, a single missing `decref` causes a memory leak, and an extra `decref` causes a crash (segmentation fault).

1. Reference Ownership Categories

Every pointer returning from a C-API function belongs to one of three reference categories:

New References The C function receives absolute ownership of the reference. It must decrement the reference count when finished, or pass ownership back to CPython (e.g., by returning the object from the function).

- *Examples:* `PyLong_FromLong()`, `PyList_New()`, `PyObject_Call()`.

Borrowed References The C function receives a pointer without ownership. It does not own the reference and must not decrement it unless it explicitly calls `Py_INCREF()` to claim ownership. If the owner of the object frees it, a borrowed pointer becomes dangling.

- *Examples:* `PyTuple_GetItem()`, `PyList_GetItem()`.

Stolen References Some C-API functions take ownership of references passed to them. The caller no longer owns the reference and must not decrement it, as the receiver will decrement it during its own deallocation pass.

- *Examples:* `PyTuple_SetItem()`, `PyList_SetItem()`.

```

1 void add_item_to_tuple_example(void) {
2     PyObject* my_tuple = PyTuple_New(1);           /* Returns New Reference
3     PyObject* my_val = PyLong_FromLong(42);        /* Returns New Reference
4     /*
5     /* PyTuple_SetItem steals the reference to 'my_val' */
6     PyTuple_SetItem(my_tuple, 0, my_val);
7
8     /*
9     Do NOT call Py_DECREF(my_val). 'my_tuple' now owns it.
10    If we decref'd my_val, freeing my_tuple would trigger a double-free
11    crash.
12    */
13    Py_DECREF(my_tuple); /* Safely deallocates tuple and my_val */
14 }
```

2. Safe Reference Macros

CPython provides macros to handle reference updates safely: `* Py_INCREF(op) / Py_DECREF(op)`: Non-null safe updates. Passing a NULL pointer triggers a crash. `* Py_XINCREF(op) / Py_XDECREF(op)`: Null-safe updates. They check if `op` is NULL and do nothing if it is, protecting code in error-handling block cleanups:

```

1 void error_handling_cleanup_example(PyObject *a, PyObject *b) {
2     // If a or b failed allocation and were set to NULL, XDECREF handles it
3     safely
4     Py_XDECREF(a);
5     Py_XDECREF(b);
6 }
```

26.0.3 24.3 C-API Stable ABI (PEP 384)

Standard C extensions are compiled against a specific Python version's headers (e.g., Python 3.10). They depend on concrete structure offsets (such as `sizeof(PyObject)` or offsets of `ob_type`). When Python updates its internal layouts (e.g., changing the type struct in 3.11), these compiled binaries break, requiring recompilation for the new Python version.

To resolve this compilation dependency, PEP 384 introduced the **Stable ABI** (Application Binary Interface) / Limited API.

```

1 +-----+
2 |           Standard C Extension           |
3 | Directly accesses structures (e.g.      | ---> Fast, but must compile
4 | op->ob_refcnt, op->ob_type)             | for every Python version.
5 +-----+
6
7 +-----+
8 |           Stable ABI Extension           |
9 | Uses opaque pointers and accessors      | ---> Compile once. Runs on
10 | (e.g., Py_REFCNT(op), PyType_GetSlot) | Python 3.5, 3.6, ... 3.13+
11 +-----+

```

1. Opaque Pointers and Opaque Structures

When compiling under the Stable ABI, structure definitions (like `PyObject` or `PyTypeObject`) are treated as **opaque structures**. You cannot compile code that accesses their internal fields directly:

```

1 /* Standard API (will not compile under Py_LIMITED_API) */
2 PyTypeObject* type = op->ob_type;
3
4 /* Stable ABI equivalent */
5 PyTypeObject* type = Py_TYPE(op);

```

2. Defining the Limited API

To restrict compilation to the Stable ABI, define `Py_LIMITED_API` before importing `<Python.h>` or set it as a compiler flag:

```

1 #define Py_LIMITED_API 0x030A0000 /* Targets Python 3.10 and later */
2 #define PY_SSIZE_T_CLEAN
3 #include <Python.h>

```

This guarantees that the resulting compiled binary links only with stable symbols exported by the dynamic library (`libpython3.so` or `python3.dll`), allowing the extension to load and run on any subsequent Python version without recompilation.

Chapter 27

Metaclasses, Descriptor Protocol, and type Slots

27.0.1 25.1 The Metaclass Class Creation Pipeline

In Python, classes are objects themselves. A **metaclass** is the class of a class; it defines how classes are constructed. The default metaclass for all objects is `type`.

When CPython executes a class definition block:

```
1 class MyClass(BaseClass, metaclass=MyMeta):  
2     x = 1
```

The VM resolves class creation by running these sequential steps:

1. Namespace Preparation (`__prepare__`)

Before executing the class body code, CPython checks if the metaclass has a `__prepare__` method. If present, it calls:

```
1 namespace = MyMeta.__prepare__('MyClass', (BaseClass,))
```

This returns a mapping object (usually a standard empty dictionary, but it can be a custom namespace like a `collections.OrderedDict` or a custom dictionary subclass). The interpreter then executes the class body code within this namespace object, populating it with attributes (methods, class variables, annotations).

2. Class Object Allocation (`__new__`)

After executing the class body, CPython invokes the metaclass `__new__` method to allocate the class object:

```
1 cls = MyMeta.__new__(MyMeta, 'MyClass', (BaseClass,), namespace)
```

At the C level, this redirects to the type object's allocation slot `PyType_Type.tp_new` (implemented in `Objects/typeobject.c` as `type_new`). This allocates a new `PyTypeObject` struct on the heap, setting up its class fields, base classes, and MRO (Method Resolution Order).

3. Class Initialization (`__init__`)

Once the class object is allocated, CPython initializes it by calling:

```
1 MyMeta.__init__(cls, 'MyClass', (BaseClass,), namespace)
```

This allows the metaclass to inspect or modify the newly created class object before returning it.

27.0.2 25.2 The Descriptor Protocol and Attribute Lookup Precedence

A descriptor is an object that customizes attribute access behavior by implementing methods in the descriptor protocol: `* __get__(self, instance, owner):` Customizes attribute reads. `* __set__(self, instance, value):` Customizes attribute writes. `* __delete__(self, instance):` Customizes attribute deletions.

1. Data Descriptors vs. Non-Data Descriptors

- **Data Descriptor:** Implements both `__get__` and `__set__` (or `__delete__`).
- **Non-Data Descriptor:** Only implements `__get__` (typically used for methods).

This separation is critical because it dictates how the attribute lookup algorithm prioritizes the descriptor over instance dictionaries.

2. Attribute Lookup Precedence Hierarchy

When retrieving an attribute `obj.name` (where `obj` is an instance of class `c`), CPython resolves the lookup path in this strict order:

1. **Class MRO Search for Data Descriptor:** Search the Method Resolution Order (MRO) of class `c` for an attribute named `name`. If found, and it is a **Data Descriptor**, call its `__get__` method and return the result:

```
result = DataDescriptor.__get__(desc, obj, C)
```

2. **Instance Dictionary Search:** Search the instance dictionary (`obj.__dict__`). If `name` is present, return the value directly, bypassing non-data descriptors:

```
result = obj.__dict__['name']
```

3. **Class MRO Search for Non-Data Descriptor:** Search the MRO of class `c` for an attribute named `name`. If found, and it is a **Non-Data Descriptor**, call its `__get__` method:

```
result = NonDataDescriptor.__get__(desc, obj, C)
```

4. **Class Attributes:** Search the MRO of class `c` for a standard class attribute. If found, return the value directly.
5. **Fallback to `__getattr__`:** If the attribute is not found, and the class defines `__getattr__`, call it:

```
result = obj.__getattr__('name')
```

6. **Raise `AttributeError`:** If all steps fail, raise `AttributeError`.


```

1      [Lookup obj.name]
2      |
3      Does "name" exist in class MRO as a Data Descriptor?
4      / \
5      Yes No
6      / \
7      [Call descriptor __get__] Does "name" exist in instance __dict__?
8      / \
9      Yes No
10     / \
11     [Return from __dict__] Does "name" exist in class MRO
12                             as a Non-Data Descriptor?
13                             / \
14                             Yes No
15                             / \
16                             [Call descriptor] Is there a class
attribute?
17                                     / \
18                                     Yes No
19                                     / \
20                             [Return value] Raise AttributeError

```

27.0.3 25.3 CPython Type Slots

To optimize attribute access and function calls at the C level, CPython uses **Type Slots**. Instead of looking up methods like `__repr__` or `__add__` in the class dictionary at runtime, CPython populates static C function pointers directly inside the `PyTypeObject` struct definition:

```

1 typedef struct _typeobject {
2     PyObject_VAR_HEAD
3     const char *tp_name;           /* For printing, in format <module>.<
name> */
4     Py_ssize_t tp_basicsize, tp_itemsize; /* For allocation */
5
6     /* Type Slots (Static Function Pointers) */
7     destructor tp_dealloc;         /* Deallocation handler */
8     reprfunc tp_repr;             /* Representation builder */
9     getattrofunc tp_getattr;      /* Attribute lookup hook */
10    setattrfunc tp_setattr;        /* Attribute assignment hook */
11
12    /* Protocol Table Slots */
13    PyNumberMethods *tp_as_number;  /* Math function pointers (e.g. nb_add
) */
14    PySequenceMethods *tp_as_sequence; /* Indexing function pointers (e.g.
sq_item) */
15    PyMappingMethods *tp_as_mapping; /* Map lookup function pointers (e.g.
mp_subscript) */
16 } PyTypeObject;

```

1. Protocol Sub-Tables

CPython groups protocol-specific slots into secondary struct tables pointed to by fields inside `PyTypeObject`:

- **PyNumberMethods**: Contains pointers for numeric operators (e.g., `nb_add` for `+`, `nb_multiply` for `*`):

```

1 typedef struct {
2     binaryfunc nb_add;
3     binaryfunc nb_subtract;
4     binaryfunc nb_multiply;
5     /* ... */
6 } PyNumberMethods;

```

- **PySequenceMethods:** Contains pointers for sequence operations (e.g., `sq_item` for indexing, `sq_length` for length):

```

1 typedef struct {
2     lenfunc sq_length;
3     binaryfunc sq_concat;
4     ssizeargfunc sq_item;
5     /* ... */
6 } PySequenceMethods;

```

- **PyMappingMethods:** Contains pointers for mapping lookups (e.g., `mp_subscript` for key lookups, `mp_ass_subscript` for key updates):

```

1 typedef struct {
2     lenfunc mp_length;
3     binaryfunc mp_subscript;
4     obobjargproc mp_ass_subscript;
5 } PyMappingMethods;

```

2. Slot Inheritance

When a class is created: 1. CPython copy-inherits these function pointer slots from the base classes defined in MRO. 2. If the subclass overrides a method (e.g., defines `def __repr__(self):`), CPython replaces the slot pointer (`tp_repr`) with a wrapper function (`slot_tp_repr`) that calls the Python method. 3. This inheritance ensures that C-level calls (like `a + b` executing `tp_as_number->nb_add(a, b)`) execute without dictionary searches, optimizing runtime speed.

Part IX

Volume IX: High Performance & Low Latency Concurrency

Chapter 28

LOW-LEVEL MEMORY OPTIMIZATION TECHNIQUES

28.0.1 26.1 `__slots__` Internals and Memory Mapping

By default, CPython allocates a dynamic dictionary (`__dict__`) for every instance of a user-defined class. This dictionary allows users to set arbitrary attributes at runtime, but introduces significant memory overhead. For small objects, storing a hash table (requiring a minimum of 8 entries, table pointers, and tracking states) can consume several hundred bytes per instance.

To optimize memory usage, developers can declare `__slots__` inside the class definition:

```
1 class OptimisedPoint:
2     __slots__ = ('x', 'y')
3     def __init__(self, x: float, y: float) -> None:
4         self.x = x
5         self.y = y
```

1. Compilation and Memory Layout Comparison

When a class defines `__slots__`, CPython alters the structure allocation of its instances: * **Without slots:** The class instances reserve a pointer to a dynamic `__dict__` struct and a `__weakref__` pointer, occupying 16 bytes (on 64-bit systems) just for metadata pointers, in addition to the hash table memory allocated when attributes are assigned. * **With slots:** CPython allocates a fixed array of pointers directly inside the instance struct itself. The class namespace defines **Member Descriptors** at fixed offset indices matching the slot attributes.

```
1 Without __slots__:
2 +-----+
3 |           PyObject           |
4 | - ob_refcnt                   |
5 | - ob_type                     |
6 +-----+
7 | - __dict__ pointer  -----> [ PyDictObject (Hash Table) ]
8 +-----+
9 | - __weakref__ pointer         |
10 +-----+
11
12 With __slots__:
13 +-----+
14 |           PyObject           |
15 | - ob_refcnt                   |
16 | - ob_type                     |
17 +-----+
18 | - slot pointer 0 ('x' float) | <-- Access via fixed offset (e.g., +16)
19 +-----+
```

```

20 | - slot pointer 1 ('y' float)      | <-- Access via fixed offset (e.g., +24)
21 +-----+

```

2. Bytecode Impact of slots

Accessing attributes on instances with slots bypasses dictionary lookups. The bytecode instructions execute via optimized offset reads:

```

1 # Without slots
2 10 LOAD_FAST          0 (self)
3 12 LOAD_ATTR          1 (x) ; Performs MRO lookup and falls back to __dict__
4
5 # With slots
6 10 LOAD_FAST          0 (self)
7 12 LOAD_ATTR          1 (x) ; Detects Member Descriptor; reads offset
   directly

```

Because the attribute offset is known at compile time, CPython skips the hash-calculation and hash-bucket collision checks, yielding a significant performance speedup in addition to memory reduction.

28.0.2 26.2 Buffer Protocol and memoryview Zero-Copy Slicing

In high-performance I/O or numerical applications, copying large arrays of binary data introduces CPU and memory bottlenecks. CPython's **Buffer Protocol** defines a C-level interface that allows objects to expose their raw, internal memory buffers directly to other Python components without allocating secondary copies.

1. C-level Py_buffer Struct Definition

An object exposes its memory layout by implementing the buffer protocol and populating a `Py_buffer` struct:

```

1 /* Include/object.h */
2 typedef struct {
3     void *buf;                /* Pointer to the start of the memory block */
4     PyObject *obj;            /* Reference to the exporting object (to
   prevent GC reclamation) */
5     Py_ssize_t len;           /* Total length of the buffer in bytes */
6     Py_ssize_t itemsize;      /* Size of a single element in bytes */
7     int readonly;             /* Set to 1 if the buffer is read-only */
8     const char *format;       /* Format string describing the elements (
   struct style, e.g. "f" for float) */
9     int ndim;                 /* Number of dimensions */
10    Py_ssize_t *shape;         /* Array of sizes for each dimension */
11    Py_ssize_t *strides;       /* Array of step offsets (strides) for each
   dimension */
12    Py_ssize_t *suboffsets;    /* Suboffsets array (used for nested arrays) */
13    void *internal;           /* Private data for allocator tracking */
14 } Py_buffer;

```

- `buf`: The raw C pointer to the data.
- `strides`: Dictates the distance in bytes to jump to reach the next element in a given dimension.

2. memoryview Wrapper and Zero-Copy Slicing

A `memoryview` is a Python-level wrapper around the C-level buffer protocol. It allows Python code to perform slicing operations that do not copy data. Instead, slicing creates a new `memoryview` object sharing the same `buf` pointer with adjusted `shape`, `strides`, and `len` parameters:

```

1 # Zero-copy slicing demo
2 raw_data = bytearray(b"Mastering CPython Memory Allocation")
3 mv = memoryview(raw_data)
4
5 # Slice shares the memory buffer of raw_data; no new allocations occur
6 mv_slice = mv[10:17]
7 print(mv_slice.tobytes()) # Output: b"CPython"
8
9 # Modifying the slice modifies the source object directly
10 mv_slice[0] = ord('c')
11 print(raw_data)           # Output: bytearray(b"Mastering cPython Memory
                             Allocation")

```

28.0.3 26.3 Compact Serialization: array vs. struct

Standard Python lists are arrays of pointers pointing to heap-allocated `PyObject` instances scattered across virtual memory. This introduces cache-locality issues (cache misses) and substantial reference counting overhead.

1. array.array

The `array` module provides a compact, contiguous data structure that stores homogeneous primitive types (integers, floats) directly in a contiguous memory block, mimicking C arrays:

```

1 import array
2 # Allocates a contiguous block of memory containing 1 million native 32-bit
   floats
3 float_array = array.array('f', [1.0] * 1_000_000)

```

Compared to a list of floats (which requires 8 bytes for the pointer, plus 24 bytes for each float object), `array.array` consumes exactly 4 bytes per element, yielding up to an 8x memory reduction and improved cache locality during sequential iterations.

2. Struct Packing

The `struct` module enables packing and unpacking Python objects to and from binary representation buffers matching C structures, which is ideal for network serialization or binary file I/O:

```

1 import struct
2
3 # Layout: 32-bit Integer (i), 32-bit float (f), and 8-byte string (8s)
4 # Total size: 4 + 4 + 8 = 16 bytes of contiguous binary data
5 binary_data = struct.pack('if8s', 42, 3.14159, b'data_str')
6
7 # Unpack back to Python primitives
8 unpacked_vals = struct.unpack('if8s', binary_data)
9 print(unpacked_vals) # Output: (42, 3.1415927410125732, b'data_str')

```

Chapter 29

CPU & I/O BOUND SYSTEM CONCURRENCY

29.0.1 27.1 Concurrency Paradigms: Threads vs. Processes vs. Coroutines (Asyncio)

Python supports three core concurrency paradigms, each targeting different system limitations. The table below outlines their architectural differences:

Paradigm	Execution Type	GIL Limitation	Primary Use Case	Overhead	Communication Method
Multithreading (threading)	Preemptive multitasking (managed by OS).	Yes (standard CPython blocks CPU scaling).	I/O-bound tasks (waiting for sockets, files).	Medium (thread stacks, context switches).	Shared memory (requires locks, mutexes).
Multiprocessing (multiprocessing)	Preemptive parallel execution (separate OS processes).	No (each process has its own GIL/heap).	CPU-bound computations (math, data parsing).	High (process fork/spawn cost, private heaps).	IPC (Pipes, queues, shared memory, pickles).
Asynchronous (asyncio / coroutines)	Cooperative multitasking (single thread, explicit yields).	Yes (single-threaded).	High-concurrency network services (web servers).	Low (coroutine objects are lightweight heap structures).	Direct variable access (no locks required).

29.0.2 27.2 Asyncio Event Loops and uvloop Internals

Python's standard `asyncio` event loop uses select-based system calls (`selectors` module, e.g. `epoll` on Linux, `kqueue` on macOS) to track file descriptors. While functional, the default implementation is written in pure Python and introduces overhead when wrapping callback schedules.

1. uvloop Optimization

`uvloop` is a drop-in replacement for the default `asyncio` event loop. Written in Cython and built on top of `libuv` (the high-performance asynchronous I/O engine powering Node.js), it replaces

CPython's loop implementation with C-level structures.

```

1 CPython Default asyncio Loop:
2 [Asyncio Code] ---> [Pure Python Event Loop] ---> [selectors (epoll/kqueue)]
3
4 uvloop Event Loop:
5 [Asyncio Code] ---> [Cython Wrapper] ---> [libuv (C-native epoll/kqueue)]

```

Libuv optimizes execution paths by: * **System Call Reduction**: Batching read/write operations to minimize transitions between user space and kernel space. * **Direct Memory Buffers**: Allocating internal buffers directly in C, avoiding intermediate Python byte allocations. * **Zero-Overhead Timers**: Implementing timer queues using binary heaps at the C level.

2. Configuring uvloop in Python

To use uvloop, register it as the default event loop policy at the entry point of your application:

```

1 import asyncio
2 import sys
3
4 # Register uvloop policy on supported platforms
5 if sys.platform != 'win32':
6     import uvloop
7     asyncio.set_event_loop_policy(uvloop.EventLoopPolicy())
8
9 async def main():
10     print("Running on optimized uvloop engine.")
11
12 if __name__ == '__main__':
13     asyncio.run(main())

```

With uvloop, network throughput can increase by 2x to 4x, reaching speeds comparable to implementations in Go or Node.js.

29.0.3 27.3 Multiprocessing Shared Memory IPC

Standard multiprocessing in Python relies on **Pipes** and **Queues** for Inter-Process Communication (IPC). When a process sends an object to another process: 1. The sender **pickles** (serializes) the object into bytes. 2. The bytes are written to an OS socket or pipe. 3. The receiver reads the bytes and **unpickles** (deserializes) them to allocate new objects on its private heap.

For large data structures (like lists of floats or images), this serialization pipeline degrades performance.

1. Shared Memory Architecture

Python 3.8 introduced the `multiprocessing.shared_memory` module, which maps a block of virtual memory across the address spaces of multiple OS processes. Both processes can read and write to the same physical RAM block directly, avoiding serialization overhead.

<pre> 1 Process 1 (Address Space) 2 +-----+ 3 Virtual Memory 4 [Mapped Segment] -+-----+ 5 +-----+ 6 </pre>	<pre> 1 Process 2 (Address Space) 2 +-----+ 3 Virtual Memory 4 +-----+ [Mapped Segment] 5 +-----+ 6 </pre>
---	--


```

7          +-----+
8          |   Physical RAM   |
9          | [Shared Block]  |
10         +-----+

```

2. Robust Shared Memory Code Example

Below is a complete script demonstrating parent and child processes communicating using shared memory:

```

1 import time
2 from multiprocessing import Process
3 from multiprocessing.shared_memory import SharedMemory
4
5 def child_process_worker(shm_name):
6     # 1. Attach to existing shared memory block using unique name
7     existing_shm = SharedMemory(name=shm_name)
8
9     # 2. Access buffer as a memoryview array slice
10    buffer = existing_shm.buf
11
12    print(f"[Child] Connected to shared memory. Current contents: {bytes(buffer[:10])}")
13
14    # Modify data in place (zero-copy modification)
15    for i in range(10):
16        buffer[i] = ord('A') + i
17
18    print(f"[Child] Modified buffer contents in place.")
19
20    # 3. Clean up the shared memory handle
21    existing_shm.close()
22
23 def main():
24     # 1. Allocate a shared memory block of 1024 bytes
25     shm = SharedMemory(create=True, size=1024)
26     shm_name = shm.name
27     print(f"[Parent] Allocated Shared Memory block named: {shm_name}")
28
29     # Initialize buffer contents
30     shm.buf[:10] = b"0123456789"
31     print(f"[Parent] Initial buffer contents: {bytes(shm.buf[:10])}")
32
33     # 2. Spawn a child process, passing the shared memory block name
34     p = Process(target=child_process_worker, args=(shm_name,))
35     p.start()
36     p.join() # Wait for child process to finish writing
37
38     # 3. Parent reads the updated data immediately without IPC serialization
39     print(f"[Parent] Updated buffer contents read by Parent: {bytes(shm.buf[:10])}")
40
41     # 4. Release and destroy shared memory block
42     shm.close()
43     shm.unlink() # Instructs OS to free the shared memory resource
44
45 if __name__ == '__main__':
46     main()

```

Chapter 30

NUMERICAL HIGH-PERFORMANCE DATA STRUCTURES

30.0.1 28.1 NumPy Contiguous Layouts and Strides

NumPy arrays are designed to speed up mathematical operations on homogeneous datasets by utilizing contiguous blocks of memory. Unlike standard Python lists, which contain pointers to scattered object instances, a NumPy array stores the raw values directly in a single, contiguous array block, which maximizes cache hit rates.

1. C-Contiguous vs. Fortran-Contiguous Layouts

The mapping of multi-dimensional arrays to linear virtual memory is defined by the array's layout: * **C-Contiguous (Row-Major)**: Elements along rows are stored adjacent to each other in memory. The last index changes fastest when traversing the memory sequentially. This is the default layout in NumPy. * **Fortran-Contiguous (Column-Major)**: Elements along columns are stored adjacent to each other in memory. The first index changes fastest when traversing memory sequentially.

```
1 2D Array Layout (3x3):
2  [[1, 2, 3],
3   [4, 5, 6],
4   [7, 8, 9]]
5
6 C-Contiguous (Row-Major):
7  [ 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 ]
8
9 Fortran-Contiguous (Column-Major):
10 [ 1 | 4 | 7 | 2 | 5 | 8 | 3 | 6 | 9 ]
```

2. Stride Calculations

NumPy locates any element in an N-dimensional array using its **strides** vector. A stride defines the number of bytes that must be skipped in memory to move to the next index in a given dimension.

For a 2D array of shape (R, C) and data type size D (in bytes), the stride vectors are:

$$\text{Strides}_C = (C \times D, D) \quad (\text{Row-Major} / \text{C-Contiguous})$$

$$\text{Strides}_F = (D, R \times D) \quad (\text{Column-Major} / \text{Fortran-Contiguous})$$

$$\text{Memory Address} = \text{Base Address} + \sum_{i=0}^{N-1} (\text{Index}_i \times \text{Stride}_i)$$

For example, given a 3x3 array of 64-bit floats ($D = 8$ bytes) in C-contiguous layout: * **strides** = (24, 8). To move to the next row, jump 24 bytes (3 elements). To move to the next column, jump 8 bytes (1 element).

Accessing elements along strides that do not match the contiguous memory direction (e.g. slicing rows from a column-major array) breaks CPU cache locality, as the processor must retrieve values from scattered memory addresses (cache misses).

3. Vectorization and SIMD Pipelines

Contiguous arrays enable compiler vectorization. Instead of executing operations on individual elements sequentially (Scalar execution), modern CPU architectures support **Single Instruction, Multiple Data (SIMD)** instructions (e.g. AVX-512, ARM NEON).

```

1 Scalar Addition (Iterative loops):
2   Step 1: A[0] + B[0] -> C[0]
3   Step 2: A[1] + B[1] -> C[1]
4   ...
5
6 SIMD Vector Addition (Single clock cycle):
7   Vector Register A: [ A[0] | A[1] | A[2] | A[3] ]
8                     +
9   Vector Register B: [ B[0] | B[1] | B[2] | B[3] ]
10                    =
11   Vector Register C: [ C[0] | C[1] | C[2] | C[3] ]

```

By storing numbers contiguously, CPython's underlying math libraries (like BLAS or LAPACK) load entire memory blocks into hardware vector registers, computing operations across multiple elements in a single CPU clock cycle.

30.0.2 28.2 PyArrow Columnar Format and Zero-Copy Sharing

For large-scale data analysis or microservice architectures, exchanging dataframes between distinct processes (like C++ loaders, Python analytics scripts, and JVM processing nodes) introduces performance bottlenecks.

1. Apache Arrow Columnar Format

PyArrow implements the Apache Arrow format, which defines a standardized, language-independent **columnar memory layout**.

```

1 Record-Oriented Layout (Standard CSV / Database):
2 [Row 0: ID, Name, Salary] | [Row 1: ID, Name, Salary] | ...
3
4 Columnar-Oriented Layout (Apache Arrow):
5 [All IDs (contiguous)] | [All Names (contiguous)] | [All Salaries (contiguous)]

```

Columnar storage optimizes performance in three ways: * **Vectorized Scan Operations**: The CPU can scan a single attribute column (e.g. calculating average salary) without loading unrelated columns (like Name or ID) into cache memory. * **Contiguity**: Column values are stored contiguously, allowing direct SIMD instruction execution. * **Null Bitmaps**: Null values are tracked in a separate bit mask, avoiding sentinel values inside the data array.

2. Zero-Copy Sharing Protocols

Because Arrow memory layouts are identical across languages, PyArrow can map memory buffers across process boundaries using shared memory or plasma stores without performing copying or serialization.

Below is an example mapping data between PyArrow tables and Python memory layouts with zero-copy overhead:

```
1 import pyarrow as pa
2 import numpy as np
3
4 # Create a contiguous NumPy array
5 np_data = np.arange(1_000_000, dtype=np.int64)
6
7 # Wrap the NumPy array directly in a PyArrow Buffer
8 # PyArrow shares the memory address pointer of the NumPy array (zero-copy)
9 arrow_array = pa.Array.from_numpy_to_arrow(np_data)
10
11 # Create an Arrow Table
12 table = pa.Table.from_arrays([arrow_array], names=['metrics'])
13
14 # Verify that the underlying memory address remains identical
15 np_address = np_data.__array_interface__['data'][0]
16 arrow_address = table.column(0).chunks[0].buffers()[1].address
17
18 print(f"NumPy Memory Address: {hex(np_address)}")
19 print(f"Arrow Memory Address: {hex(arrow_address)}")
20 print(f"Are addresses identical (Zero-Copy)? {np_address == arrow_address}")
```

— Papermill, pandas, and machine-learning frameworks (like PyTorch or TensorFlow) utilize PyArrow columns as direct input sources, bypassing serialization bottlenecks and memory duplication.

Chapter 31

PROFILING, BENCHMARKING, AND DIAGNOSTICS

31.0.1 29.1 Deterministic Profiling (cProfile) vs. Sampling Profiling (py-spy)

Identifying execution bottlenecks is the first step toward optimization. Python developers must choose between two primary profiling methodologies depending on the target environment.

1. Deterministic Profiling via cProfile

cProfile is a built-in module that provides **deterministic profiling**. It monitors every function call, function return, and exception raise event in the program.

Under the hood, cProfile registers a profile hook at the C level using CPython's execution frame hooks (similar to `sys.setprofile`):

```
1 [CPython VM Evaluator]
2   |
3   +----> Triggers Hook ----> [cProfile Event handler] (Logs timestamp)
4   |
5   [Execute Bytecode]
6   |
7   +----> Triggers Hook ----> [cProfile Event handler] (Calculates
   duration)
```

Advantages

- **Exact Counts:** Provides exact call counts for every function in the execution tree.
- **Granular Statistics:** Tracks exact cumulative time spent inside a function vs. time spent in its child calls.

Disadvantages

- **High Overhead:** Hooking into every call/return introduces significant execution overhead (typically slowing programs down by 2x to 10x).
- **Measurement Distortion:** For fast helper functions or recursion loops, the execution cost of the profiler hook itself can be larger than the function being profiled, distorting the final statistics.

2. Sampling Profiling via py-spy

py-spy is an out-of-process sampling profiler written in Rust. Instead of intercepting the executing bytecode, it queries the operating system to read the virtual memory pages of the target


```

9     global leak_accumulator
10    if 'leak_accumulator' not in globals():
11        leak_accumulator = []
12
13    # Allocate a large list of string objects
14    leaked_data = [f"Leaked String Index {i}" for i in range(10_000)]
15    leak_accumulator.append(leaked_data)
16
17 def main():
18     # 1. Take initial baseline snapshot
19     print("Capturing baseline memory snapshot...")
20     snapshot_baseline = tracemalloc.take_snapshot()
21
22     # Run operations that allocate memory and leak it
23     print("Running leaky operations...")
24     for _ in range(5):
25         simulate_memory_leak()
26         time.sleep(0.1)
27
28     # 2. Take secondary snapshot after operations
29     print("Capturing comparison snapshot...")
30     snapshot_current = tracemalloc.take_snapshot()
31
32     # 3. Compare snapshots, grouping allocations by file line number
33     stats = snapshot_current.compare_to(snapshot_baseline, 'lineno')
34
35     print("\n=== TOP MEMORY ALLOCATION DIFFERENCES ===")
36     for stat in stats[:3]:
37         print(stat)
38         # Print the exact traceback line that allocated the leaked memory
39         print("Traceback:")
40         for frame in stat.traceback:
41             print(f"    File {frame.filename}, line {frame.lineno}: {frame.line}")
42         print("-" * 50)
43
44 if __name__ == '__main__':
45     main()

```

31.0.3 29.3 Crash Debugging via `faulthandler`

When a Python program crashes due to a low-level C error (such as a segmentation fault in a C extension, stack overflow, or memory corruption), the OS immediately terminates the process with a core dump. Standard Python try-except blocks cannot catch these crashes, and Python stack traces are lost, leaving developers with only a generic `Segmentation Fault` message.

The `faulthandler` module registers signal handlers at the operating system level for critical signals:

Signal	Description	Trigger Example
<code>SIGSEGV</code>	Segmentation fault.	Dereferencing a invalid or <code>NULL</code> pointer in C.
<code>SIGFPE</code>	Floating-point exception.	Division by zero or overflow at C level.
<code>SIGBUS</code>	Bus error.	Unaligned memory access or hardware fault.
<code>SIGILL</code>	Illegal instruction.	Execution of corrupted machine code stencils.

When one of these signals is caught: 1. The OS suspends normal process execution and invokes the handler registered by `faulthandler`. 2. `faulthandler` safely queries CPython's current

thread state (`PyThreadState_Get()`). 3. It prints the Python call stack of all running threads directly to the standard error output (`sys.stderr`), using only async-signal-safe system calls. 4. The process then terminates as usual.

To enable crash traceback logging at the start of your application:

```
1 import faulthandler
2 import sys
3
4 # Enable crash reporting, directing output to stderr
5 faulthandler.enable(file=sys.stderr, all_threads=True)
```

Alternatively, you can enable it from the terminal without modifying code by setting the environment variable:


```
1 export PYTHONFAULTHANDLER=1
```

Chapter 32

CAPSTONE PROJECT: HIGH-FREQUENCY ORDER BOOK & TRADING ENGINE

32.0.1 30.1 High-Frequency Trading Engine Architecture

A High-Frequency Trading (HFT) matching engine requires deterministic execution latencies, minimal memory allocation to avoid garbage collection pauses, and concurrent network I/O. The diagram below illustrates the modular architecture of the trading engine:



32.0.2 30.2 Optimized Order Book Implementation

To achieve sub-millisecond execution times, the engine uses: 1. `__slots__`: Avoids class dictionary allocation overhead for incoming order objects. 2. `collections.deque`: Provides $\mathcal{O}(1)$ insertion and pop times for ordering arrays at individual price points. 3. **Dictionary Cache**: Tracks active order locations by ID to achieve $\mathcal{O}(1)$ cancellations.

Below is a complete implementation of the order book engine:

```

1 import asyncio
2 from collections import deque
3 import sys
4
5 class Order:
6     __slots__ = ('order_id', 'side', 'price', 'quantity', 'timestamp')
7
8     def __init__(self, order_id: str, side: str, price: float, quantity: int)
9     -> None:
10         self.order_id = order_id
11         self.side = side          # 'B' for Bid, 'S' for Sell (Ask)
12         self.price = price
13         self.quantity = quantity
14         self.timestamp = asyncio.get_event_loop().time()
15
16     def __repr__(self) -> str:
17         return f"Order({self.order_id}, {self.side}, {self.price}, {self.
18         quantity})"
19
20 class OrderBook:
21     def __init__(self) -> None:
22         # Bids (buys): sorted in descending order (highest price first)
23         self.bids = {}
24         # Asks (sells): sorted in ascending order (lowest price first)
25         self.asks = {}
26         # Fast lookup map: order_id -> Order object
27         self.order_map = {}
28
29     def add_limit_order(self, order: Order) -> list:
30         trades = []
31         if order.side == 'B':
32             # 1. Match against Asks (sells)
33             while order.quantity > 0 and self.asks:
34                 best_ask_price = min(self.asks.keys())
35                 if order.price < best_ask_price:
36                     break # No cross; order cannot be matched immediately
37
38                 ask_queue = self.asks[best_ask_price]
39                 while order.quantity > 0 and ask_queue:
40                     matching_order = ask_queue[0]
41                     trade_qty = min(order.quantity, matching_order.quantity)
42
43                     # Execute trade
44                     order.quantity -= trade_qty
45                     matching_order.quantity -= trade_qty
46                     trades.append((order.order_id, matching_order.order_id,
47                     best_ask_price, trade_qty))
48
49                     if matching_order.quantity == 0:
50                         ask_queue.popleft()
51                         del self.order_map[matching_order.order_id]
52
53                     if not ask_queue:
54                         del self.asks[best_ask_price]
55
56             # 2. Add remaining quantity to Bids book
57             if order.quantity > 0:
58                 if order.price not in self.bids:
59                     self.bids[order.price] = deque()
60                 self.bids[order.price].append(order)
61                 self.order_map[order.order_id] = order

```

```

60     elif order.side == 'S':
61         # 1. Match against Bids (buys)
62         while order.quantity > 0 and self.bids:
63             best_bid_price = max(self.bids.keys())
64             if order.price > best_bid_price:
65                 break # No cross
66
67             bid_queue = self.bids[best_bid_price]
68             while order.quantity > 0 and bid_queue:
69                 matching_order = bid_queue[0]
70                 trade_qty = min(order.quantity, matching_order.quantity)
71
72                 # Execute trade
73                 order.quantity -= trade_qty
74                 matching_order.quantity -= trade_qty
75                 trades.append((order.order_id, matching_order.order_id,
best_bid_price, trade_qty))
76
77                 if matching_order.quantity == 0:
78                     bid_queue.popleft()
79                     del self.order_map[matching_order.order_id]
80
81                 if not bid_queue:
82                     del self.bids[best_bid_price]
83
84         # 2. Add remaining quantity to Asks book
85         if order.quantity > 0:
86             if order.price not in self.asks:
87                 self.asks[order.price] = deque()
88                 self.asks[order.price].append(order)
89                 self.order_map[order.order_id] = order
90
91         return trades
92
93     def cancel_order(self, order_id: str) -> bool:
94         if order_id not in self.order_map:
95             return False
96
97         order = self.order_map[order_id]
98         price = order.price
99         side = order.side
100
101         if side == 'B' and price in self.bids:
102             self.bids[price].remove(order)
103             if not self.bids[price]:
104                 del self.bids[price]
105         elif side == 'S' and price in self.asks:
106             self.asks[price].remove(order)
107             if not self.asks[price]:
108                 del self.asks[price]
109
110         del self.order_map[order_id]
111         return True
112
113     def get_top_of_book(self) -> tuple:
114         best_bid = max(self.bids.keys()) if self.bids else None
115         best_ask = min(self.asks.keys()) if self.asks else None
116         return best_bid, best_ask

```

32.0.3 30.3 High-Throughput Asyncio Server

To receive external market orders, we deploy an optimized asyncio TCP socket server. The server reads packet strings asynchronously, parsing and applying them directly to the order book instance.

```

1 class TradingEngineServer:
2     def __init__(self, host: str, port: int) -> None:
3         self.host = host
4         self.port = port
5         self.order_book = OrderBook()
6         self.order_counter = 0
7
8     async def handle_client(self, reader: asyncio.StreamReader, writer: asyncio
9         .StreamWriter) -> None:
10         print("[Engine] Client connection established.")
11         try:
12             while True:
13                 data = await reader.readline()
14                 if not data:
15                     break
16
17                 # Parse instruction string: "ADD B 100.50 500" or "CANCEL id"
18                 line = data.decode().strip()
19                 if not line:
20                     continue
21
22                 parts = line.split()
23                 cmd = parts[0]
24
25                 if cmd == "ADD":
26                     side, price_str, qty_str = parts[1], parts[2], parts[3]
27                     self.order_counter += 1
28                     order_id = f"ORD_{self.order_counter:06d}"
29
30                     order = Order(order_id, side, float(price_str), int(qty_str
31 ))
32                     trades = self.order_book.add_limit_order(order)
33
34                     # Send response back to broker client
35                     writer.write(f"ACK {order_id}\n".encode())
36                     for trade in trades:
37                         writer.write(f"TRADE {trade[0]} <-> {trade[1]} Price: {
38 trade[2]} Qty: {trade[3]}\n".encode())
39                     await writer.drain()
40
41                 elif cmd == "CANCEL":
42                     order_id = parts[1]
43                     success = self.order_book.cancel_order(order_id)
44                     status = "SUCCESS" if success else "NOT_FOUND"
45                     writer.write(f"CANCELED {order_id} Status: {status}\n".
46 encode())
47                     await writer.drain()
48
49                 # Periodic output of best bid/ask spreads
50                 bid, ask = self.order_book.get_top_of_book()
51                 print(f"[OrderBook] Best Bid: {bid} | Best Ask: {ask}")
52
53         except Exception as e:
54             print(f"[Engine] Exception encountered: {e}", file=sys.stderr)
55         finally:
56             writer.close()
57             await writer.wait_closed()

```

```
54         print("[Engine] Client connection terminated.")
55
56     async def run(self) -> None:
57         server = await asyncio.start_server(self.handle_client, self.host, self
58         .port)
59         addr = server.sockets[0].getsockname()
60         print(f"[Engine] Listening on socket: {addr}")
61         async with server:
62             await server.serve_forever()
63
64 if __name__ == '__main__':
65     # Runs the server loop locally on port 9999
66     # Run: telnet localhost 9999
67     # Input: ADD B 100.50 1000
68     # Input: ADD S 100.45 500
69     engine = TradingEngineServer('127.0.0.1', 9999)
70     try:
71         asyncio.run(engine.run())
72     except KeyboardInterrupt:
73         print("\n[Engine] Server stopped.")
```
